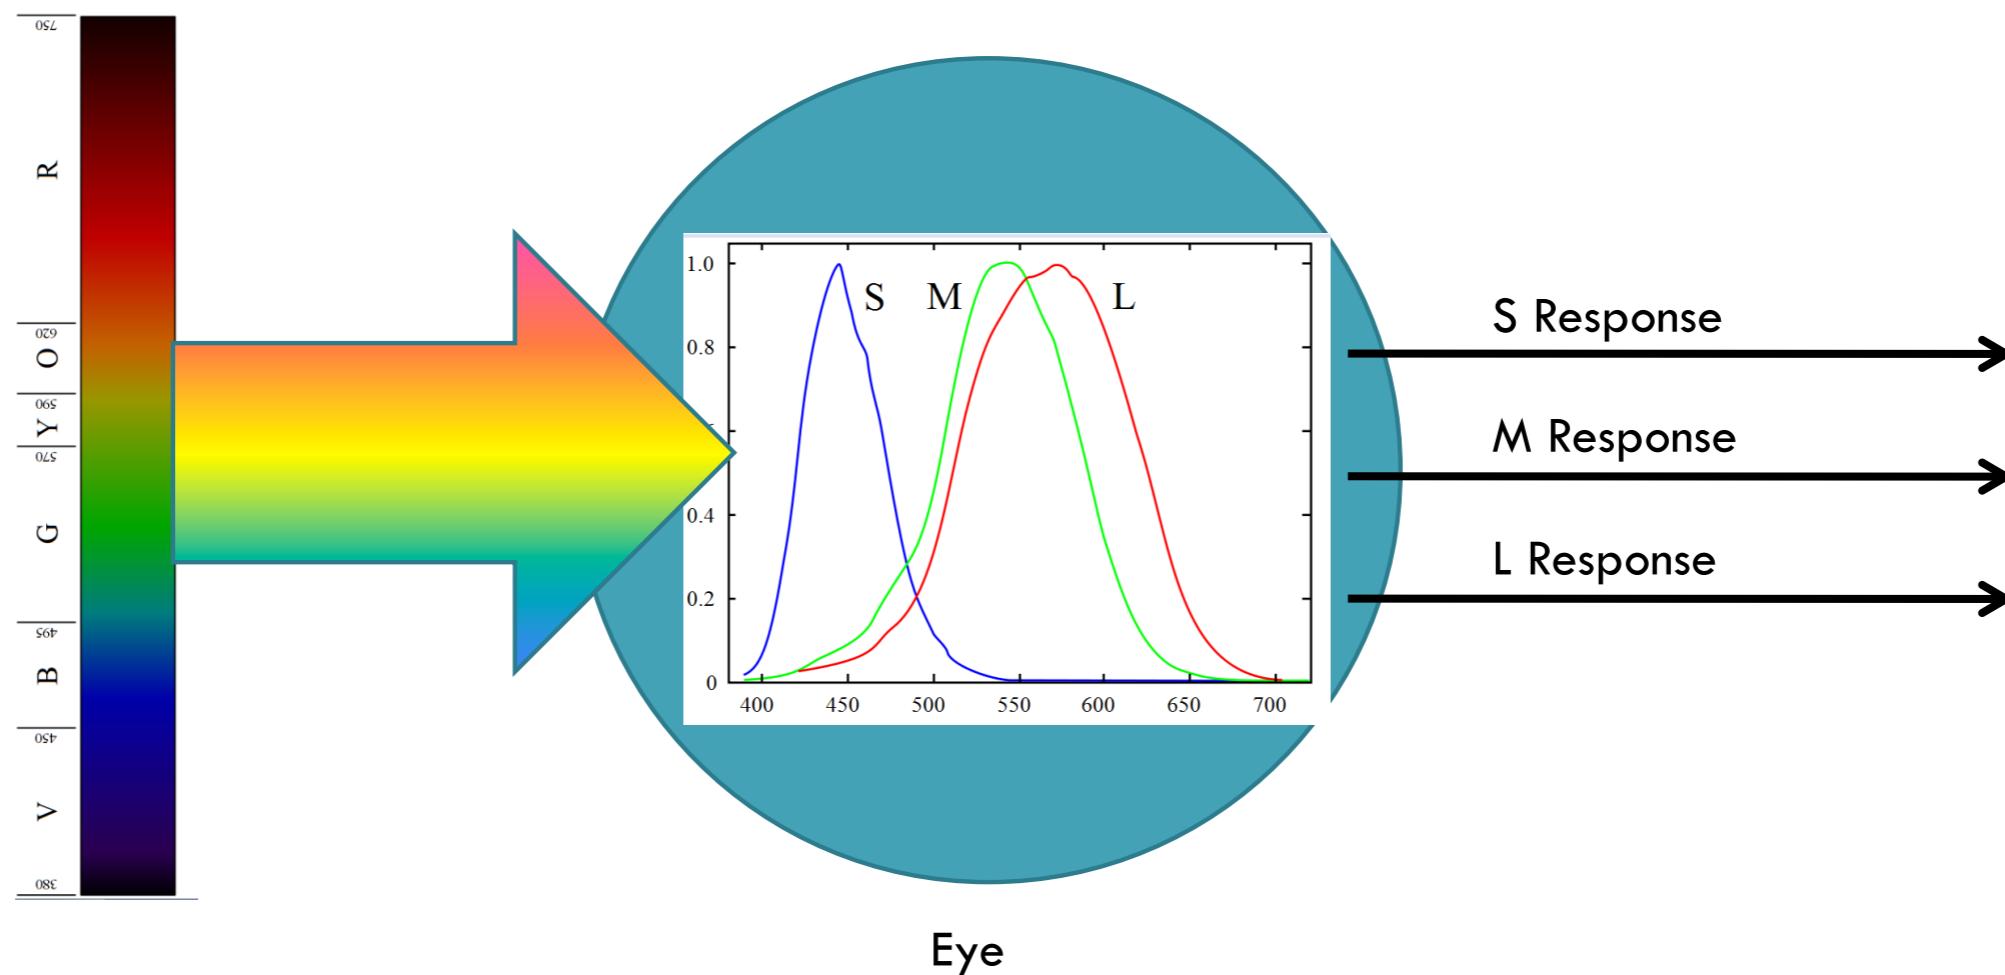
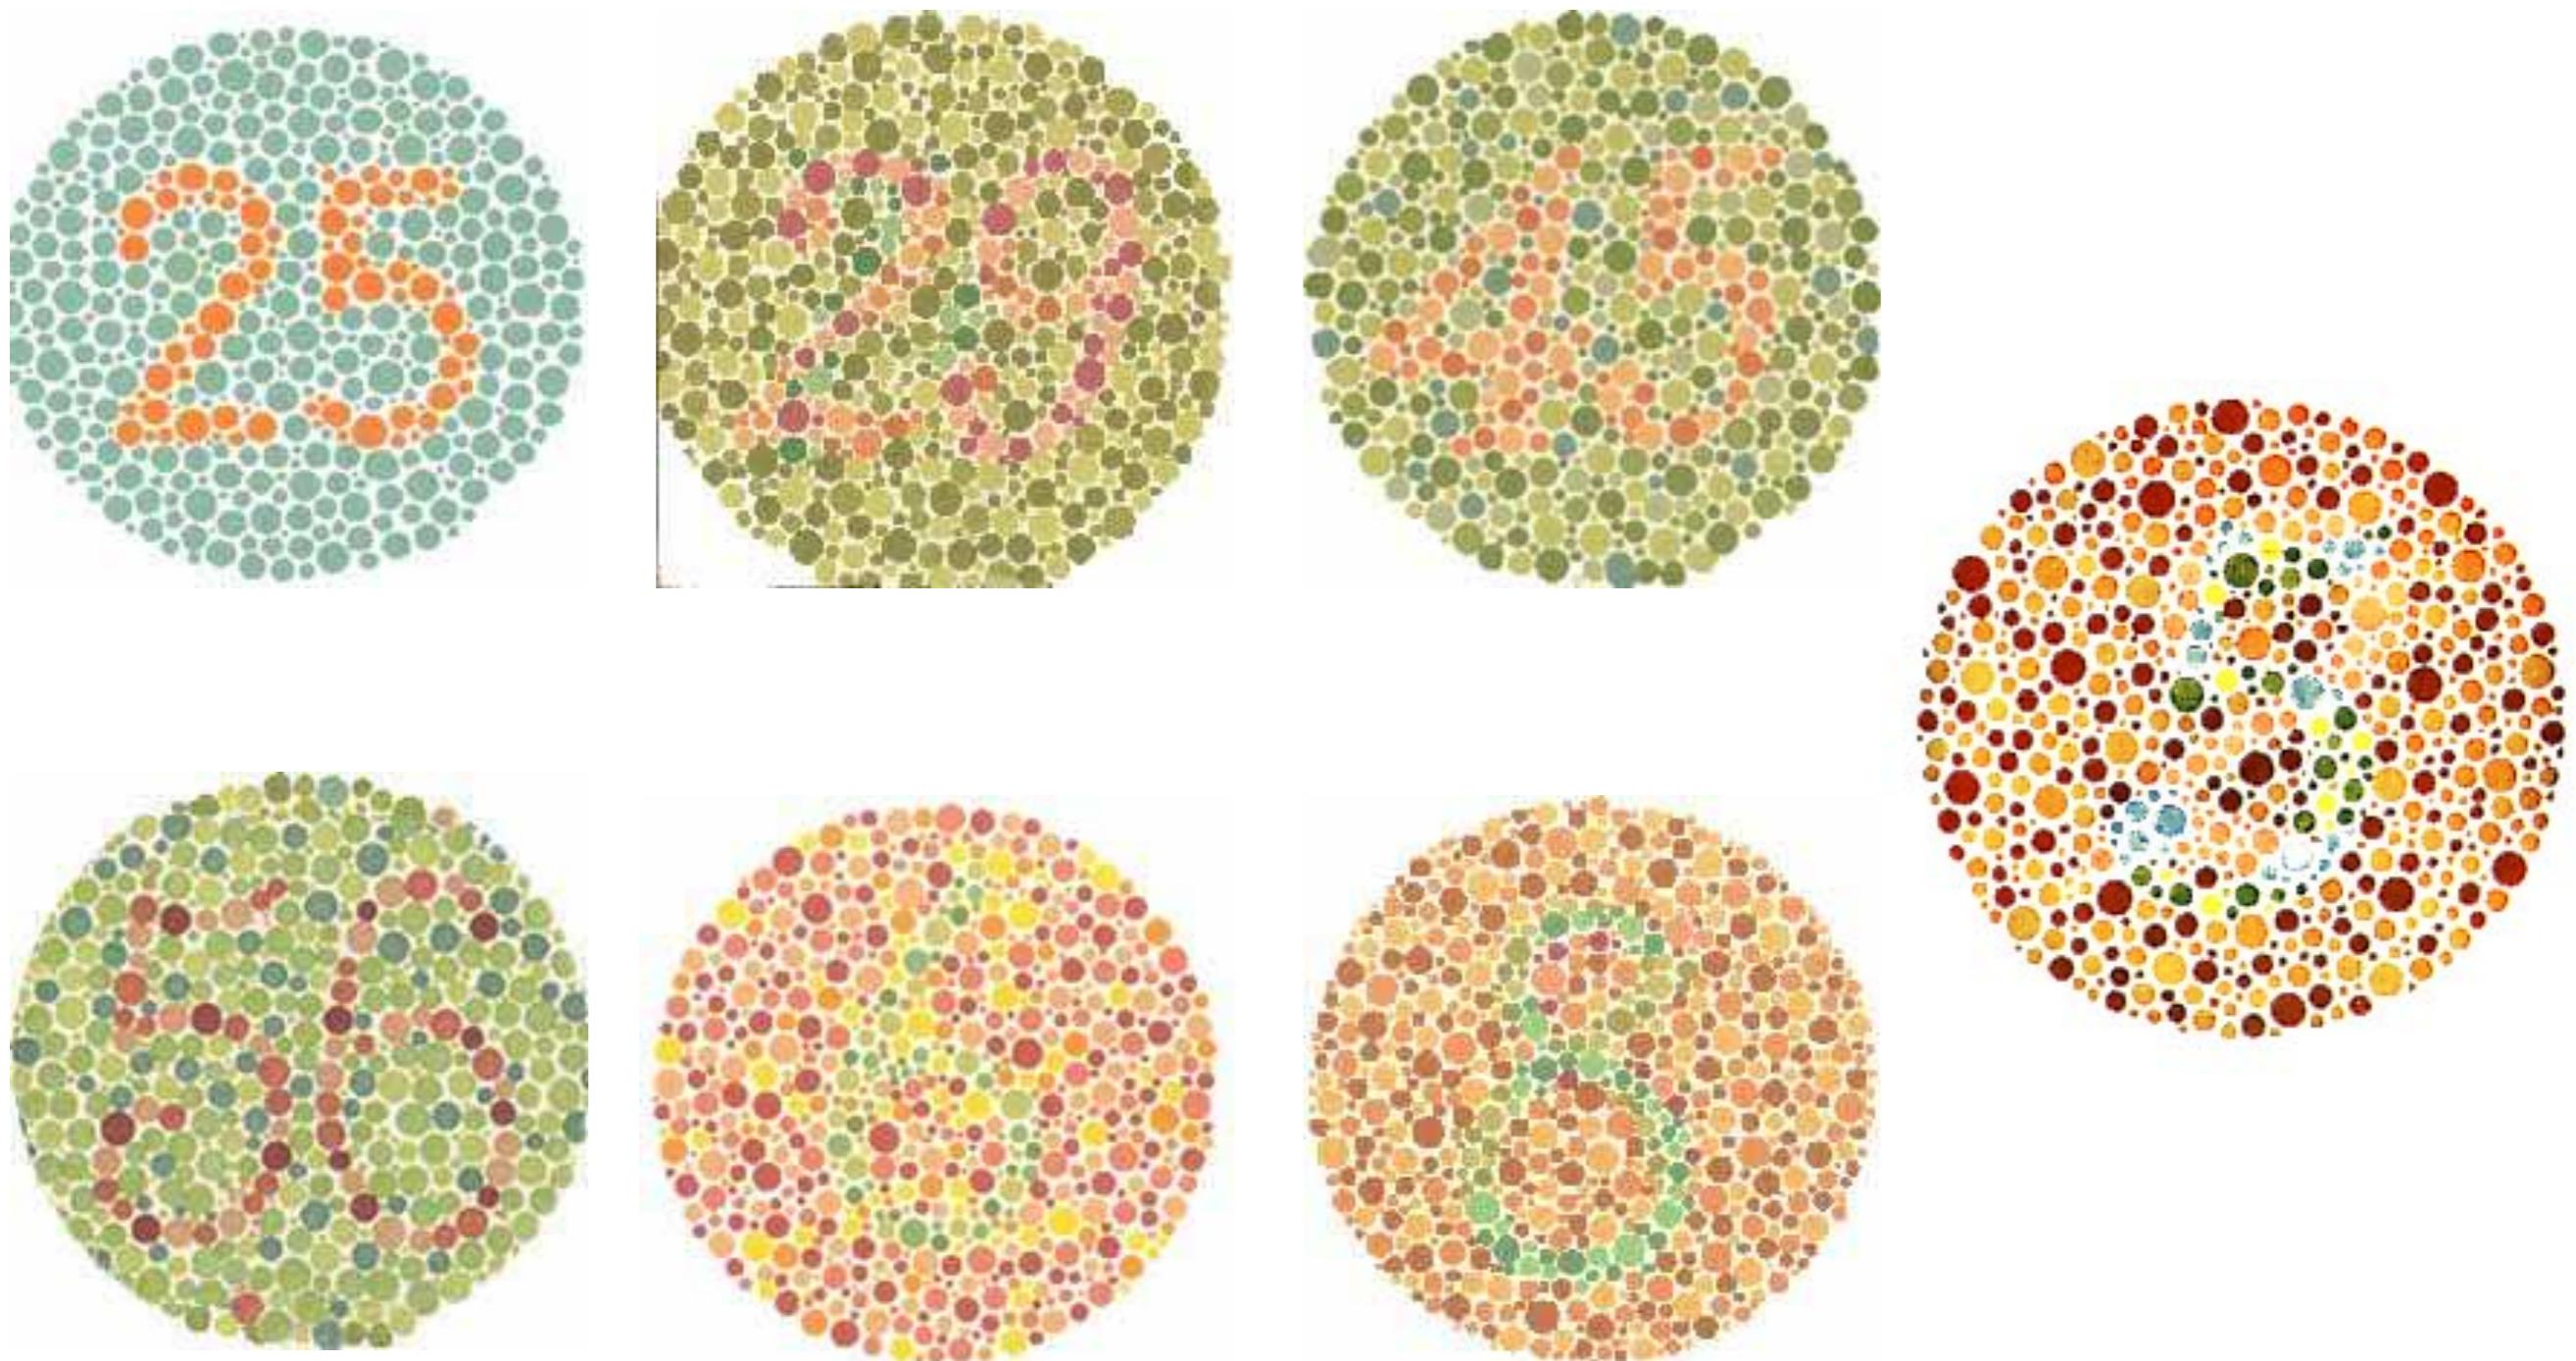


# Color Cameras

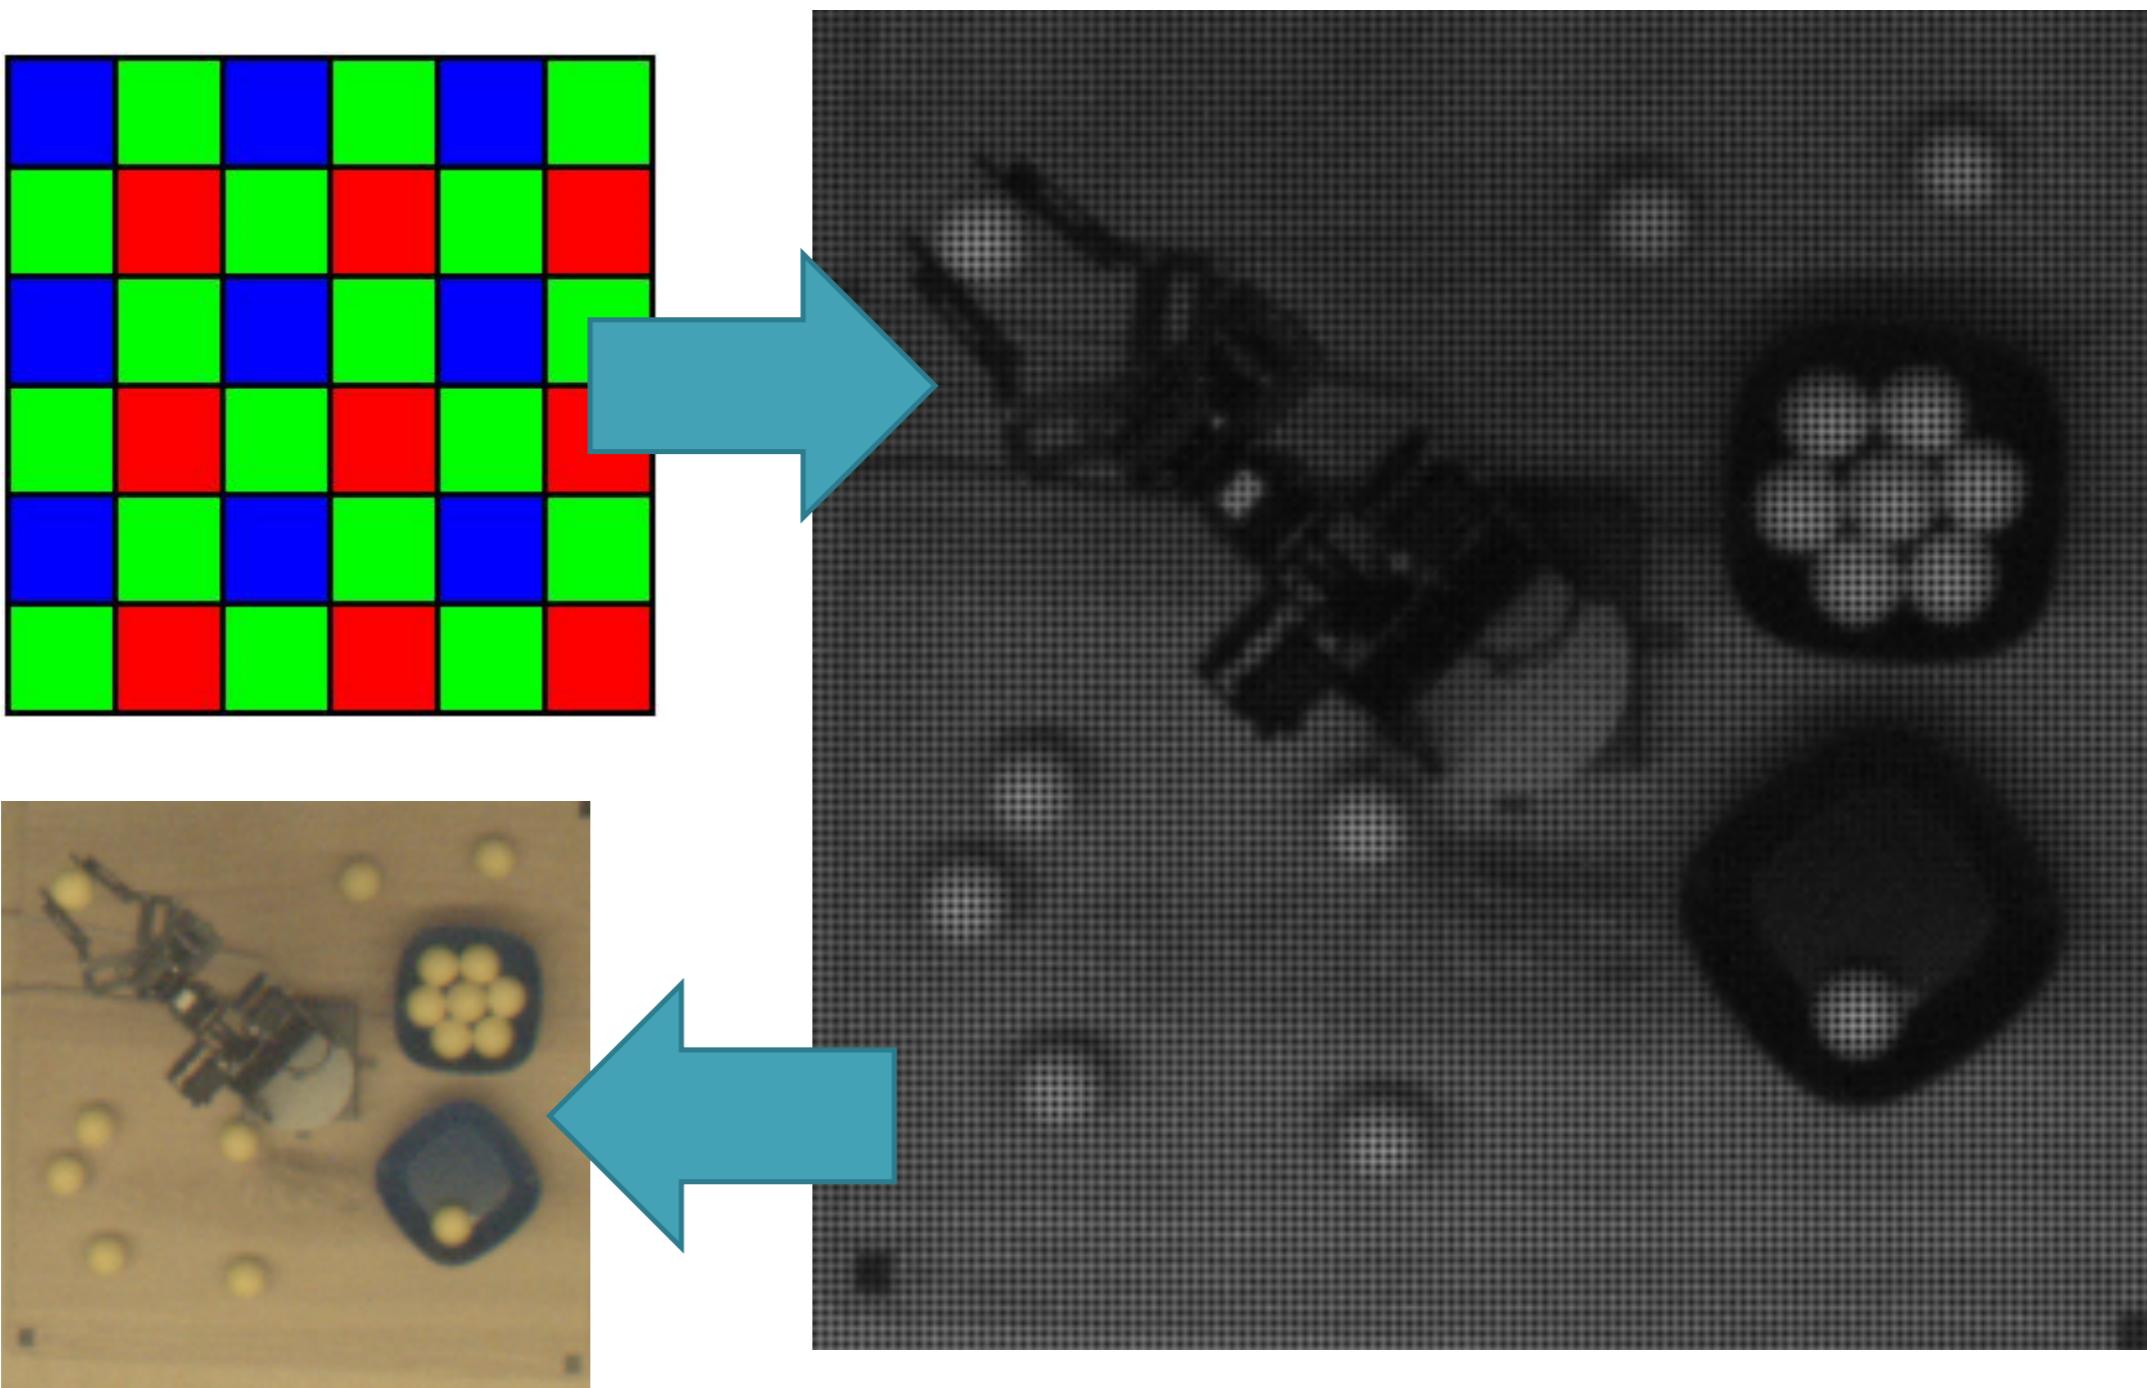
- Incoming light is described in terms of a *power spectral density*
- “Color” isn’t a physical property of light
  - ▶ It’s made up by our eyes and brain!
  - ▶ Different types of incoming light can have the same “color”



# Just for fun...

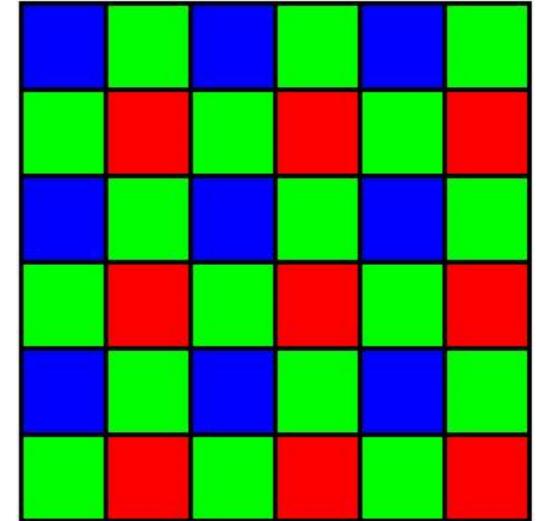


# Bayer Patterns



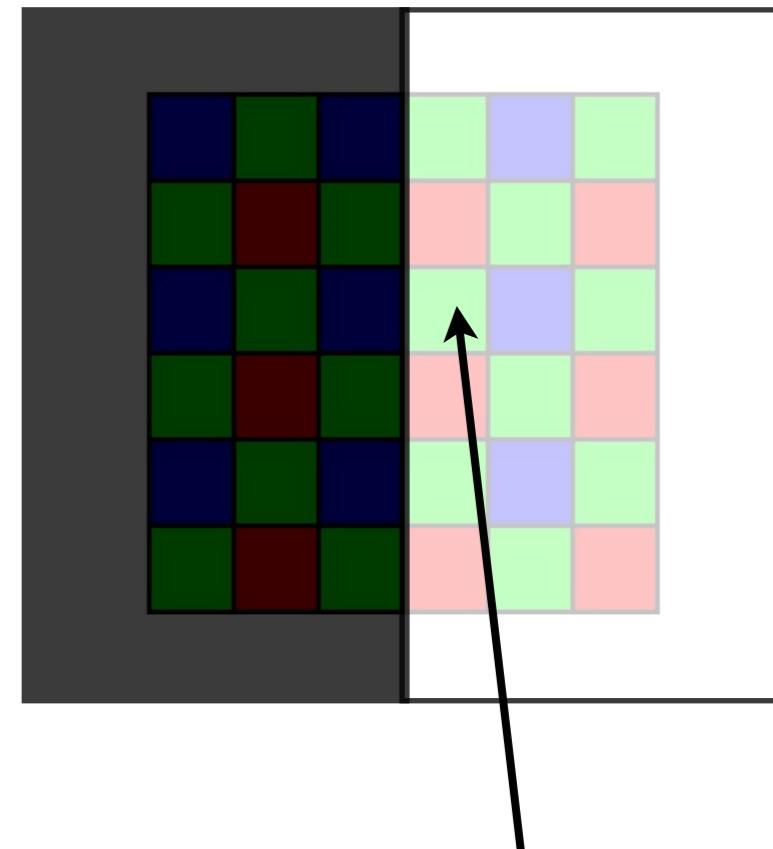
# Bayer Patterns

- Why does this matter?
  - ▶ At each pixel, two color channels are interpolated based on nearby pixels
- Thus, a color camera is more blurry than a monochrome camera.
  - ▶ e.g., Monochrome cameras give slightly better results for AprilTags



# Bayer Pattern Artifacts

- When the color of an area is uniform, Bayer patterns work well.
- What happens when there is a rapid change in color?
  - ▶ R, G, and B sub-pixels may observe different PSDs
  - ▶ Interpolated colors may not exist anywhere!



Average of nearby red pixels = red... so there will be a red output pixel even though the incoming light is either white or black.

# Why extract features from camera images?

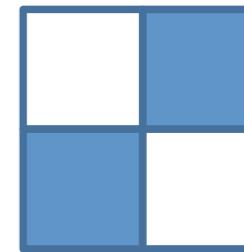
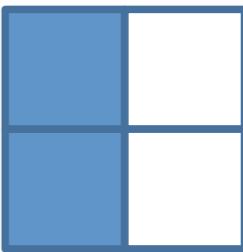
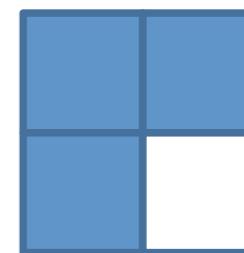
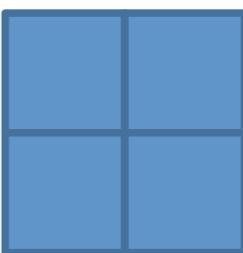
- Motivation: understanding images is really hard!
  - ▶ Lots of data
  - ▶ Some parts of the image are “boring”
- Idea: extract “good” features
  - ▶ From 1M pixels to 100s of features
  - ▶ Can make features robust



what do we mean by robust?  
detect the same objects over a range of lighting, viewpoint, rotation, scale changes

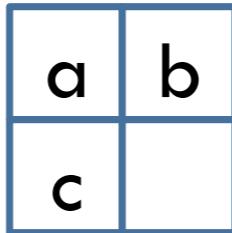
# Corner Detectors

- Intuitively, corners are a good feature.
  - ▶ Relatively easy to find
  - ▶ Trackable
- But what is a corner?
  - ▶ We're processing from bottom-up
  - ▶ No idea (yet) about objects
    - a corner  $\neq$  object corner
- What isn't a corner?
  - ▶ Uniform areas
  - ▶ Edges/lines

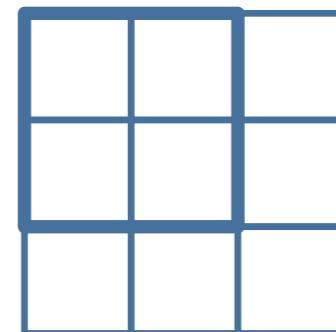


# Image Gradients

- Idea: let's look at gradients of a patch of pixels
  - ▶ Gradient at pixel a is  $(b-a, c-a)$

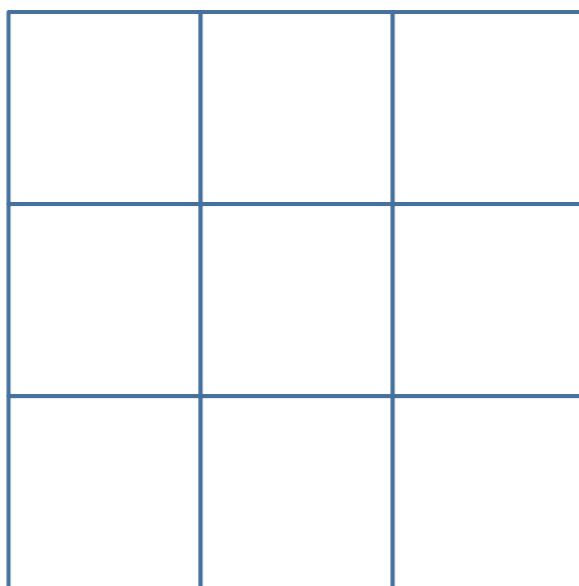


- Compute gradients for 2x2 area
  - ▶ We need 3x3 input...



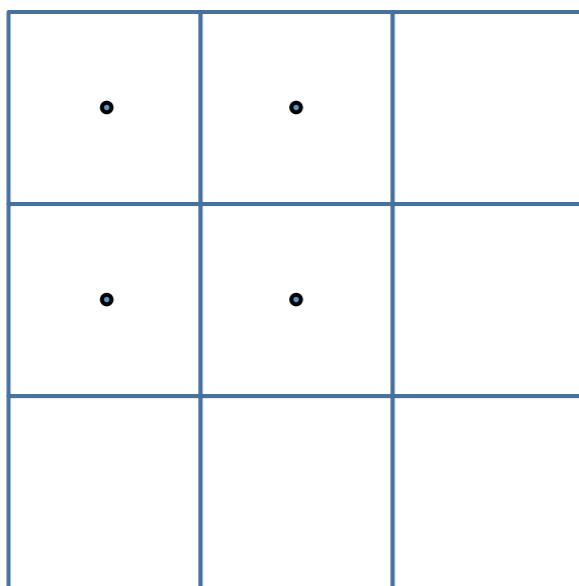
# Image Gradients

- Are these good corners?
  - ▶ (What are the gradients?)



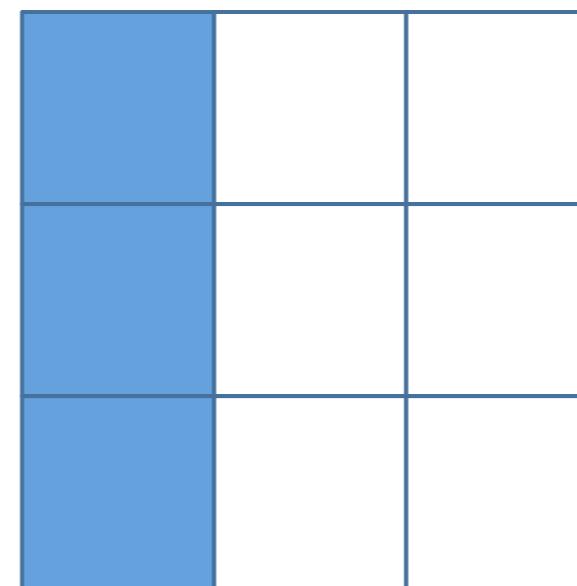
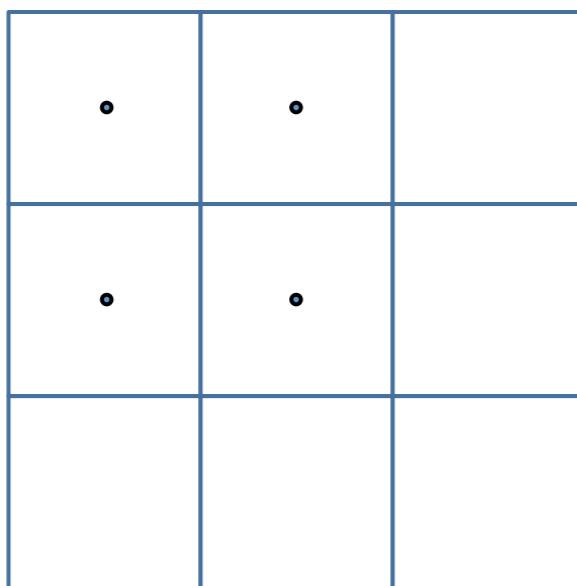
# Image Gradients

- Are these good corners?
  - ▶ (What are the gradients?)



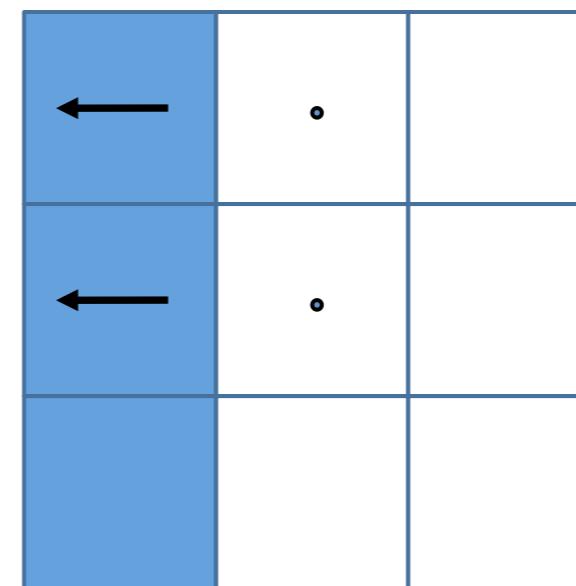
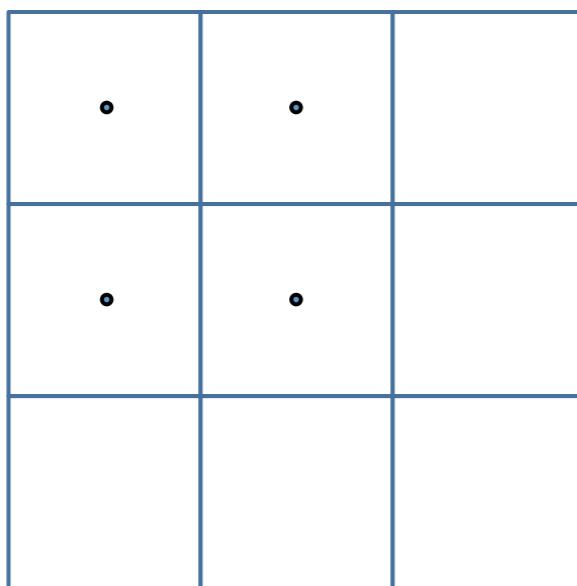
# Image Gradients

- Are these good corners?
  - ▶ (What are the gradients?)



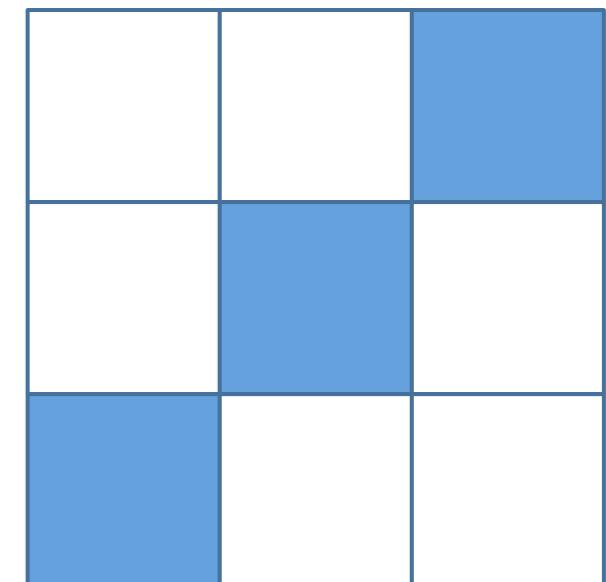
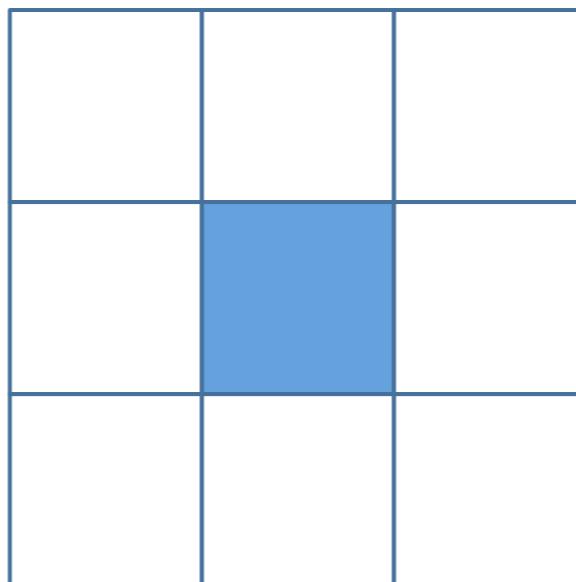
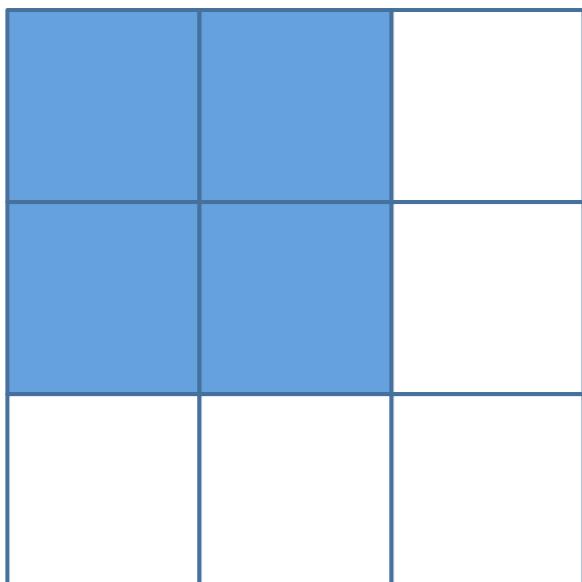
# Image Gradients

- Are these good corners?
  - ▶ (What are the gradients?)



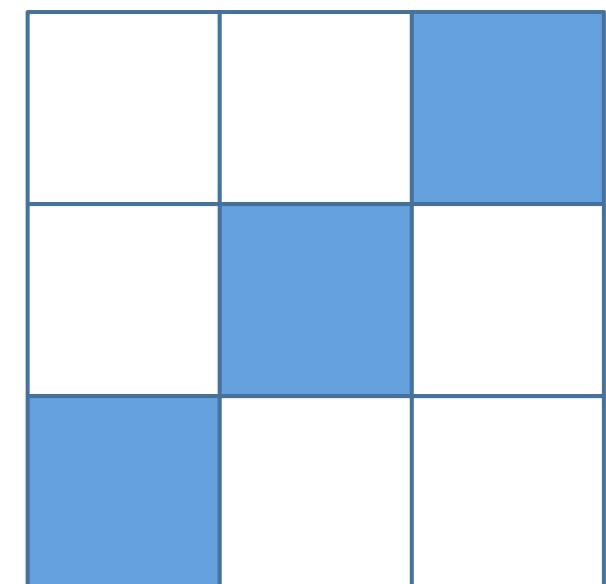
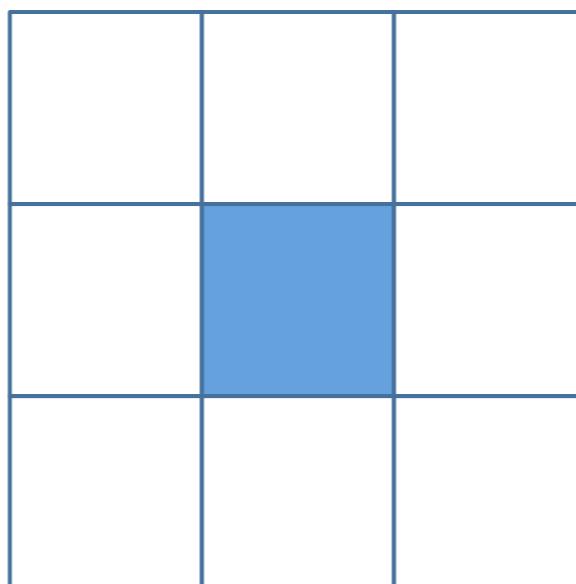
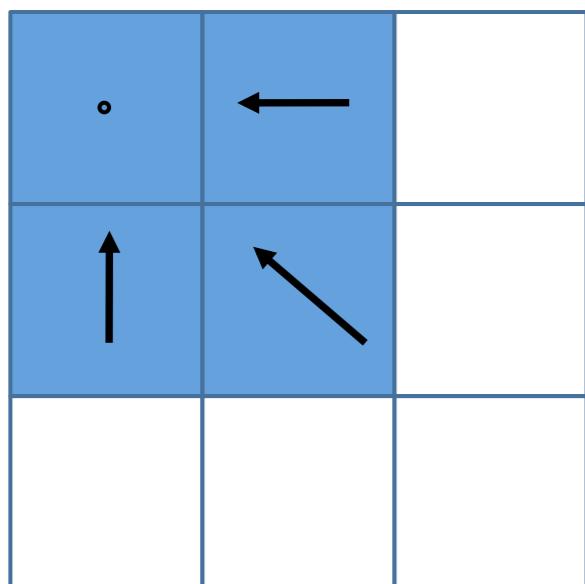
# Image Gradients

- Are these good corners?
  - ▶ (What are the gradients?)



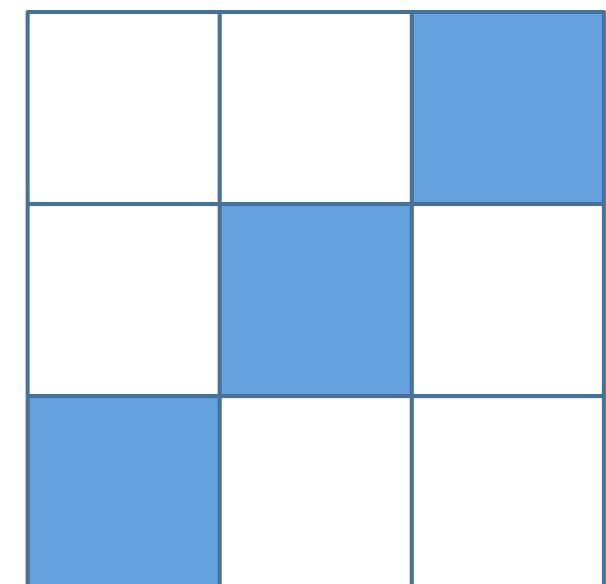
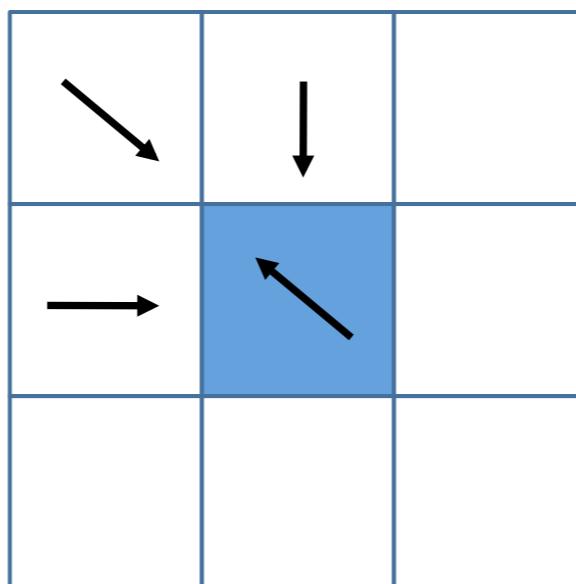
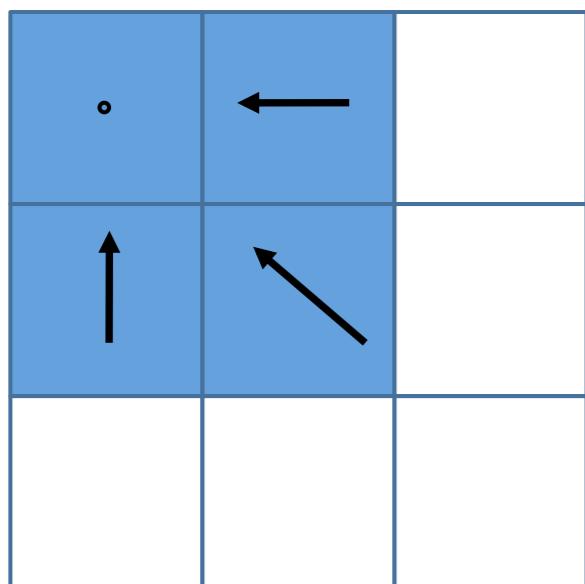
# Image Gradients

- Are these good corners?
  - ▶ (What are the gradients?)



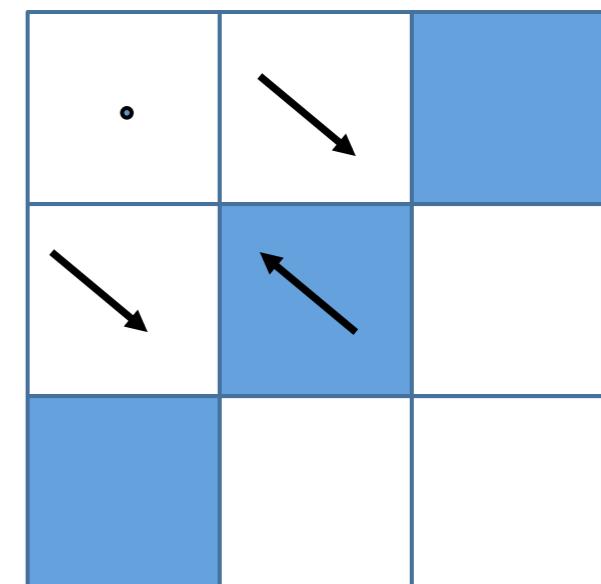
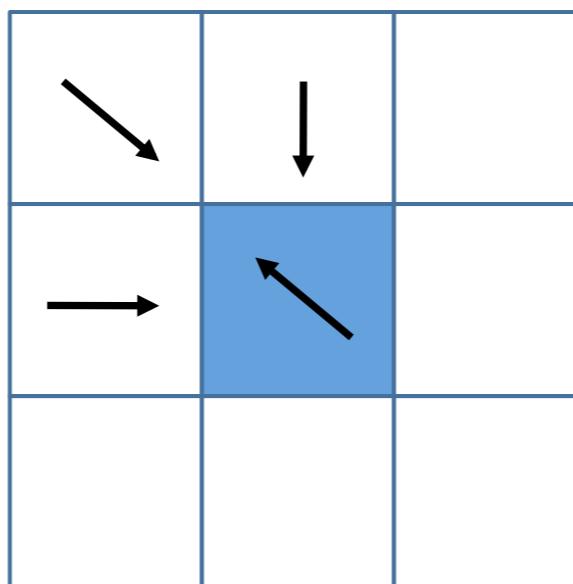
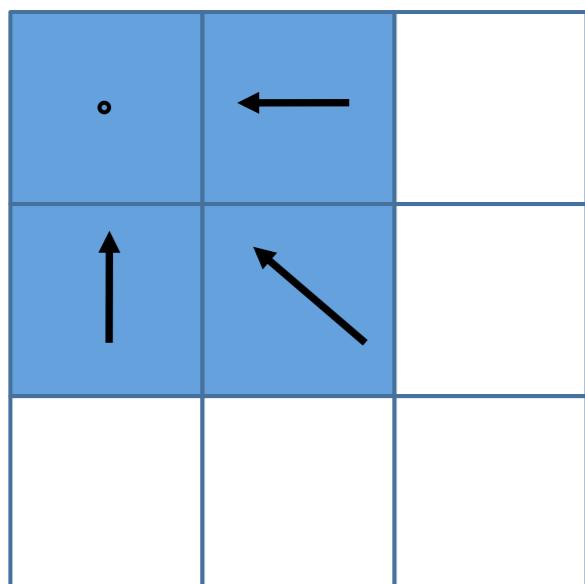
# Image Gradients

- Are these good corners?
  - ▶ (What are the gradients?)



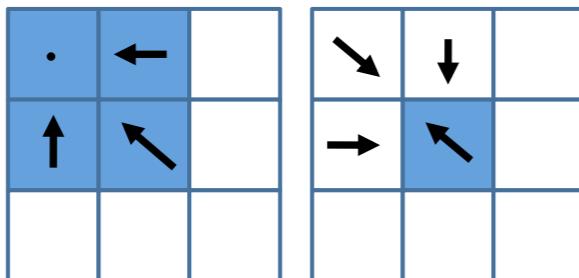
# Image Gradients

- Are these good corners?
  - ▶ (What are the gradients?)

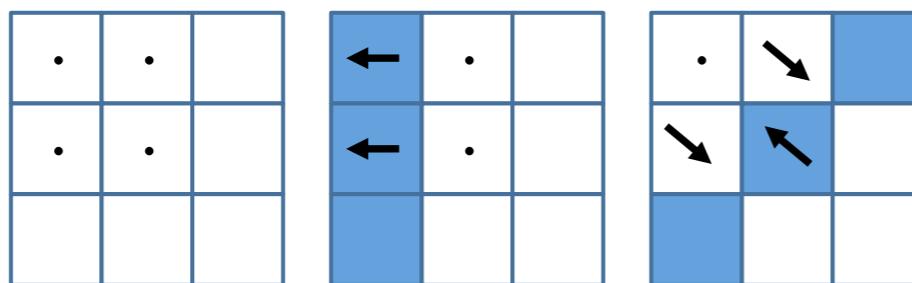


# Good and Bad Corners

- Good Corners

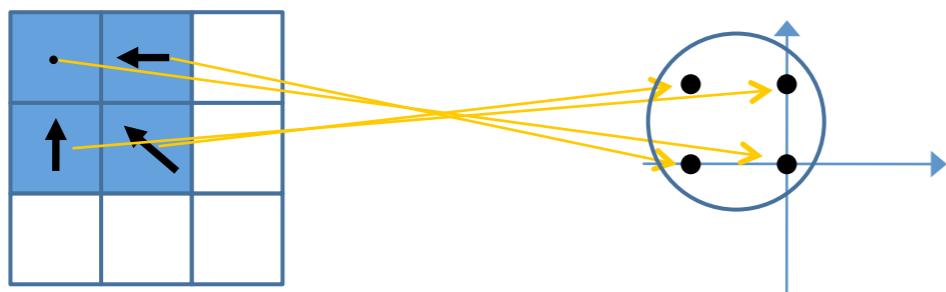


- Bad Corners

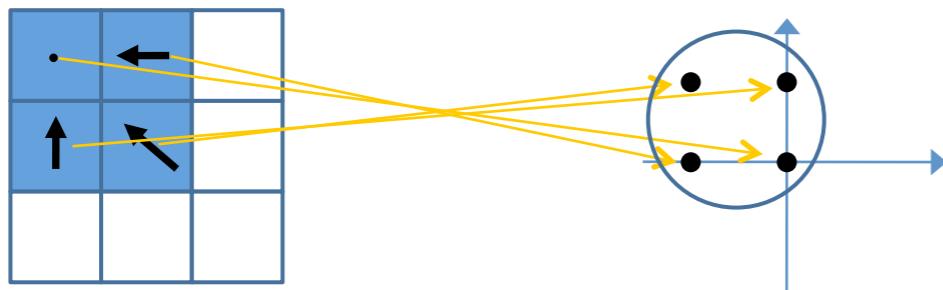


- What do good/bad corners have in common?

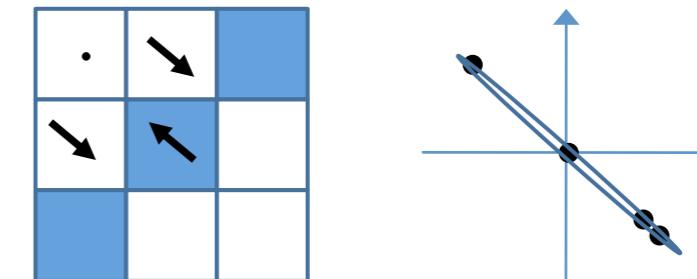
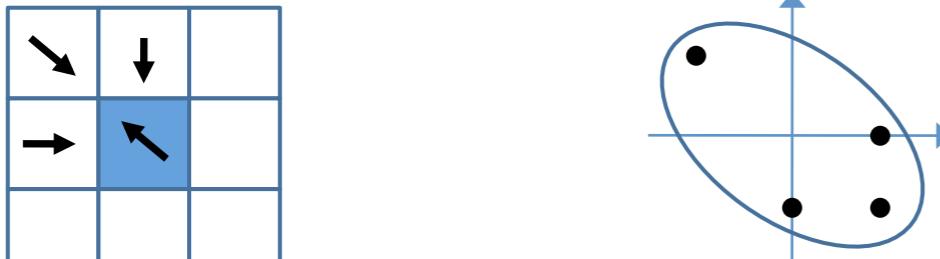
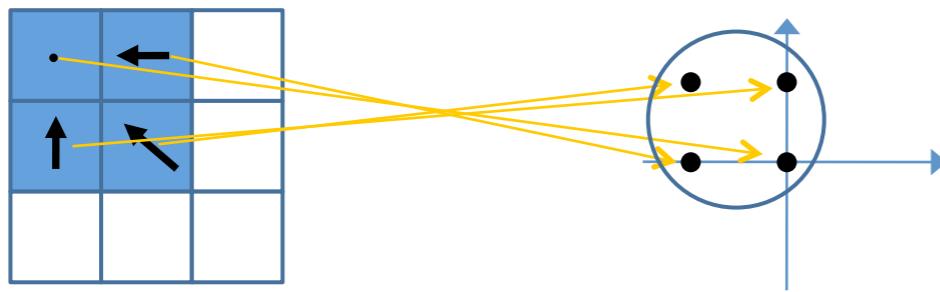
# Corners in gradient space



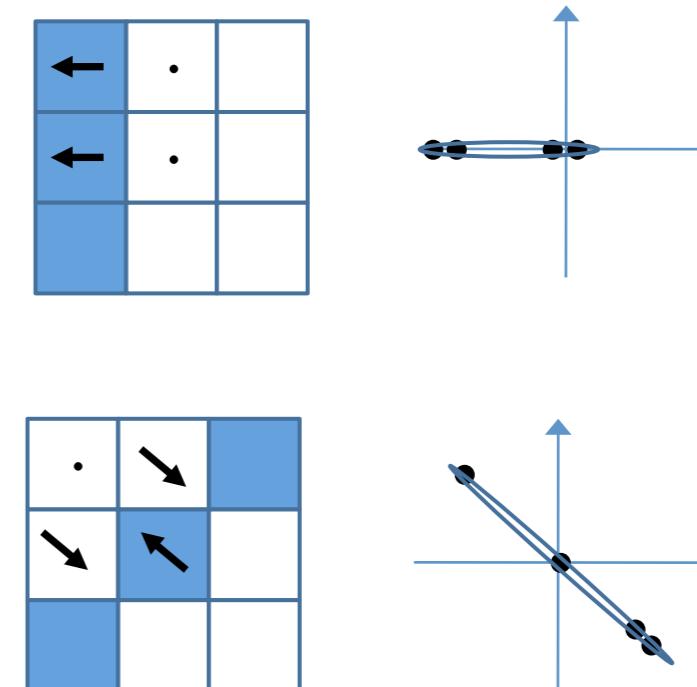
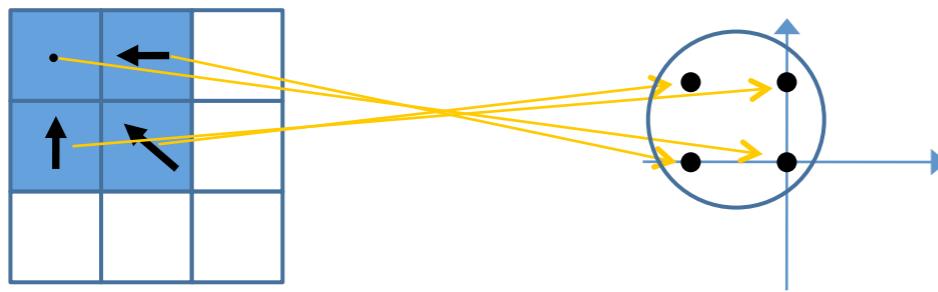
# Corners in gradient space



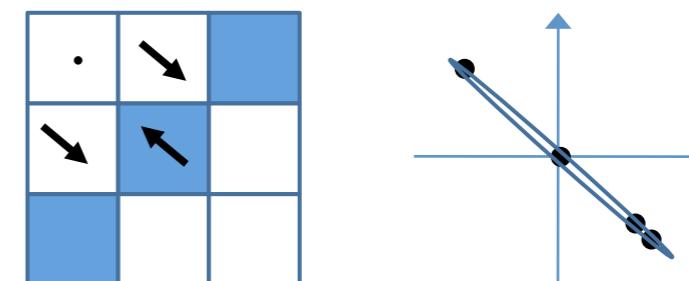
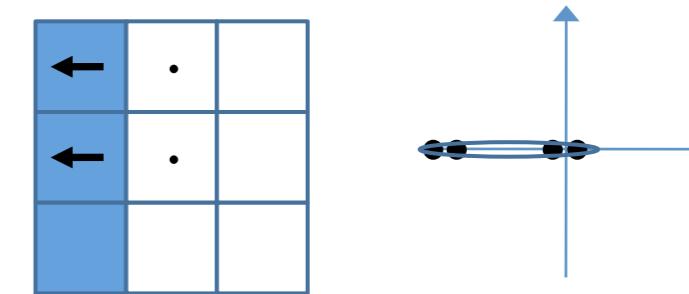
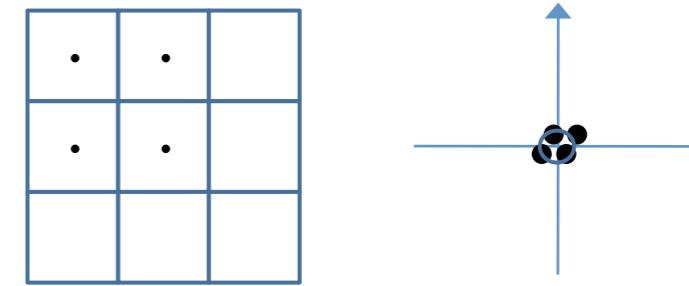
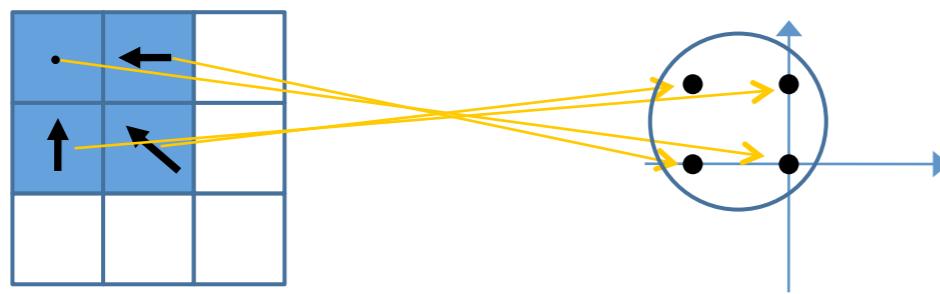
# Corners in gradient space



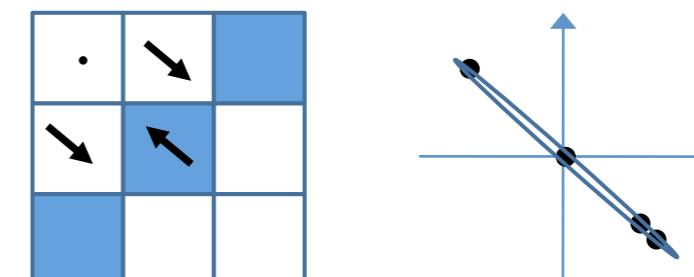
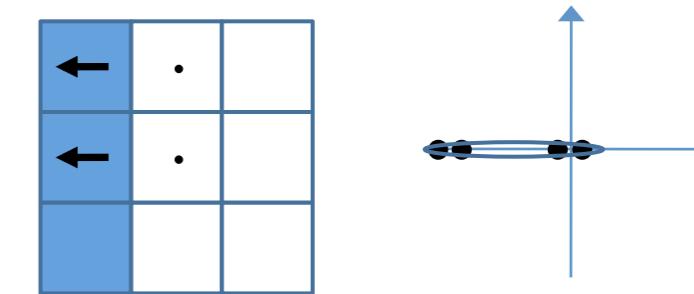
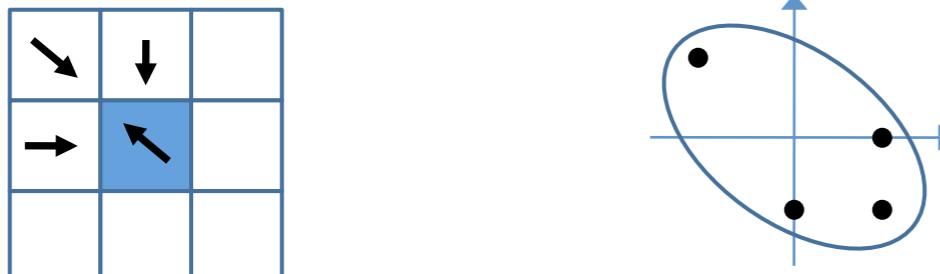
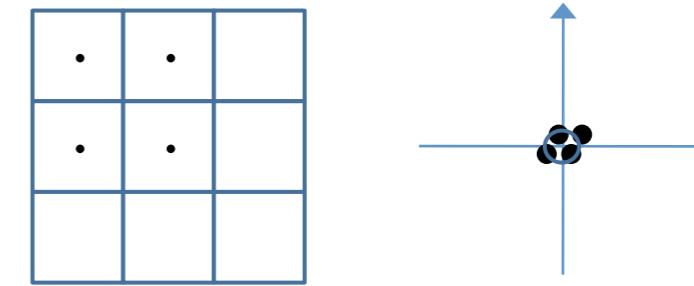
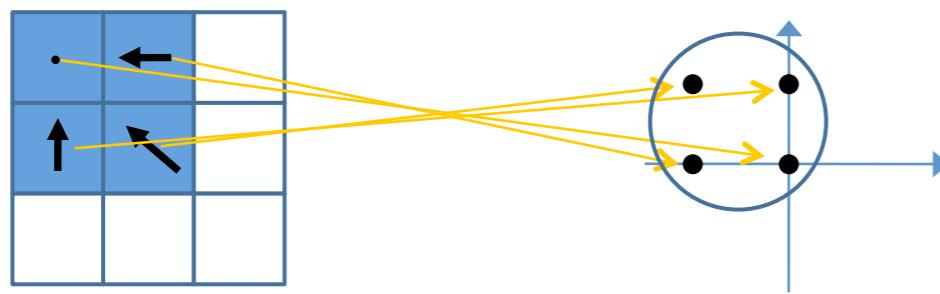
# Corners in gradient space



# Corners in gradient space

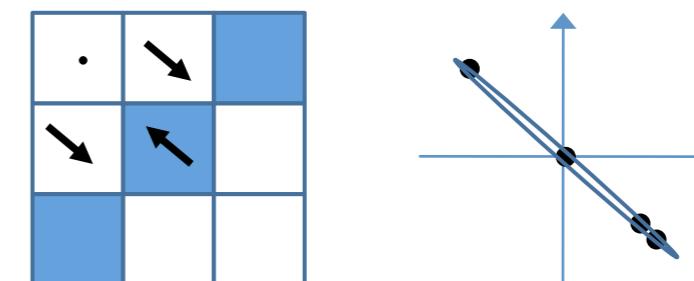
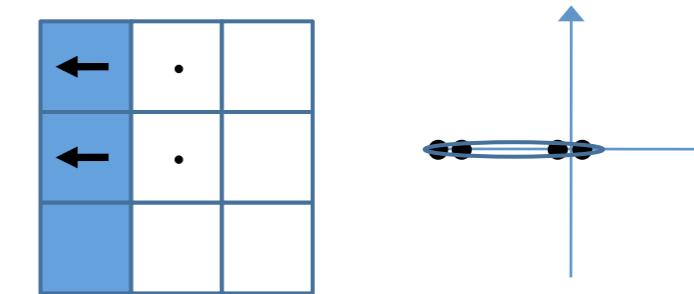
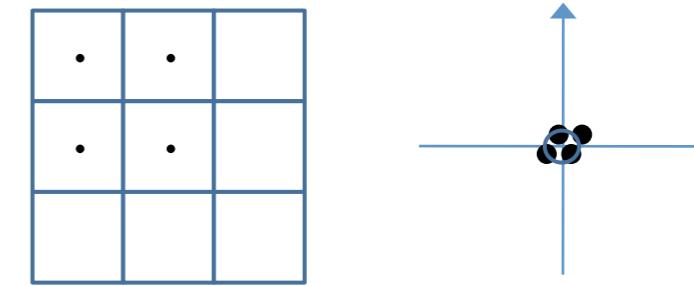
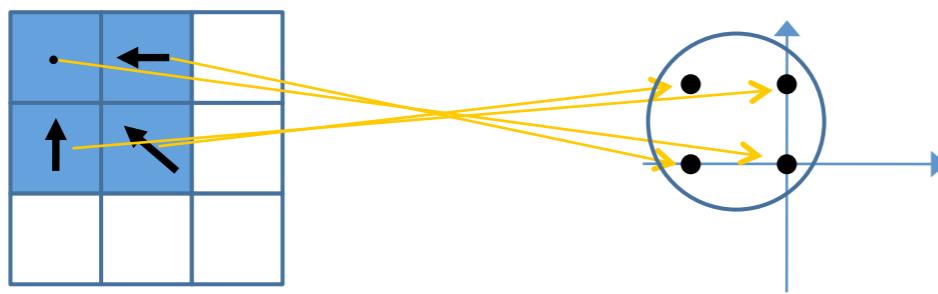


# Corners in gradient space



Good Corners

# Corners in gradient space

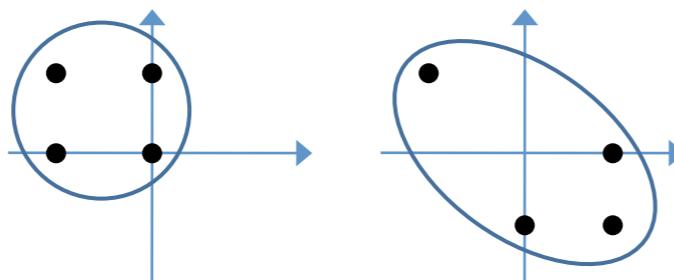


Good Corners

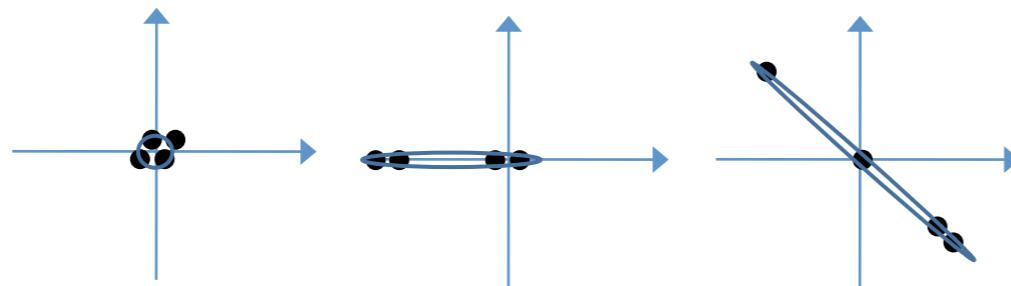
Bad Corners

# Harris Corner Detector

- Good Corners



- Bad Corners



- Idea: The “fatness” of the covariance ellipse of the gradient directions is a measure of “cornerness”
  - ▶ How do we compute this?

# Computing corner response

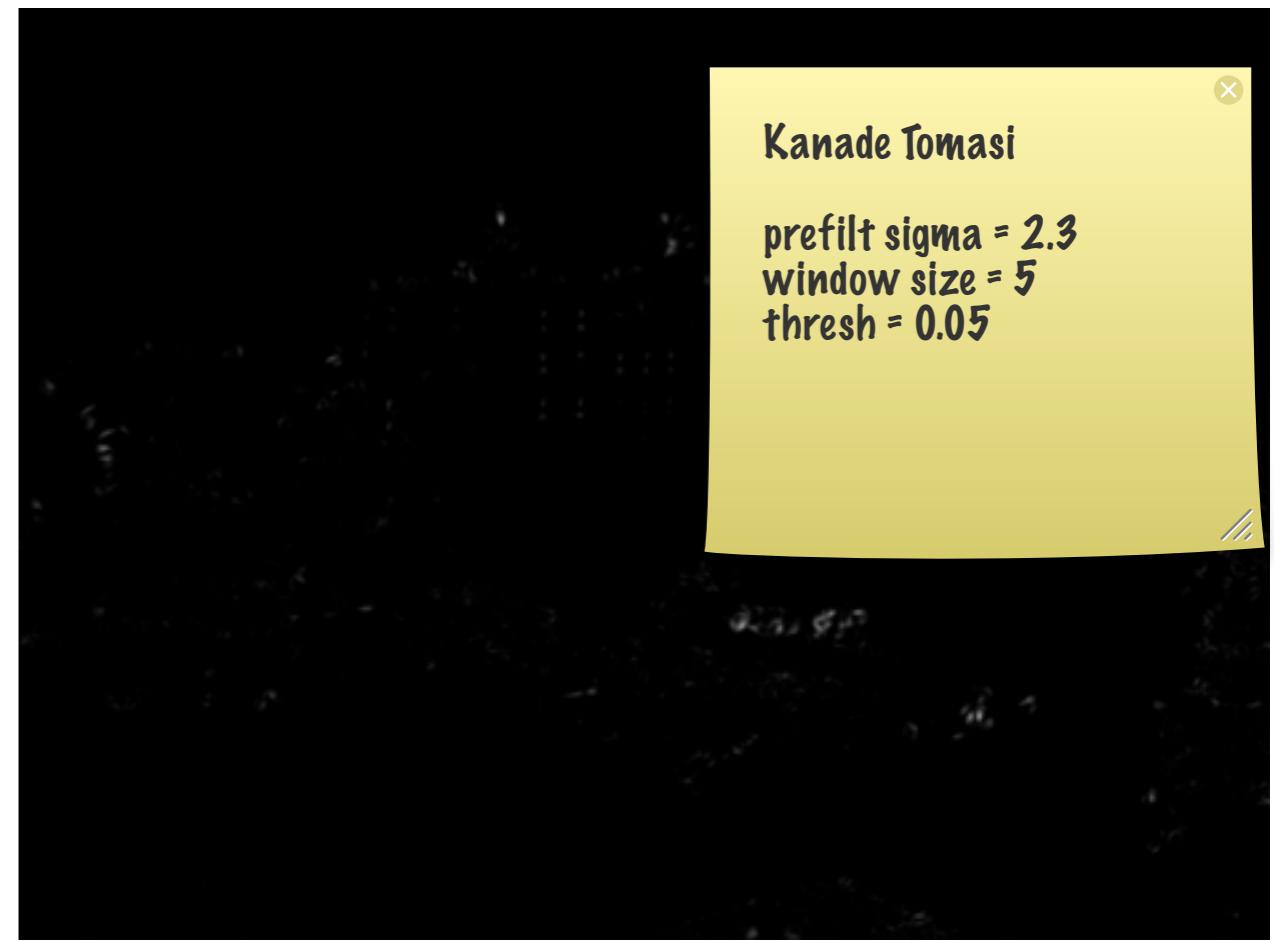
- Compute S matrix
  - ▶ Covariance of the gradients

$$S = \sum_i g_i g_i^T$$

- Compute eigen-values
  - ▶ Corner response = *smallest* eigen value
  - ▶ How bad (computationally) is this?
- Identify pixels with “good” corner responses
  - ▶ Thresholding

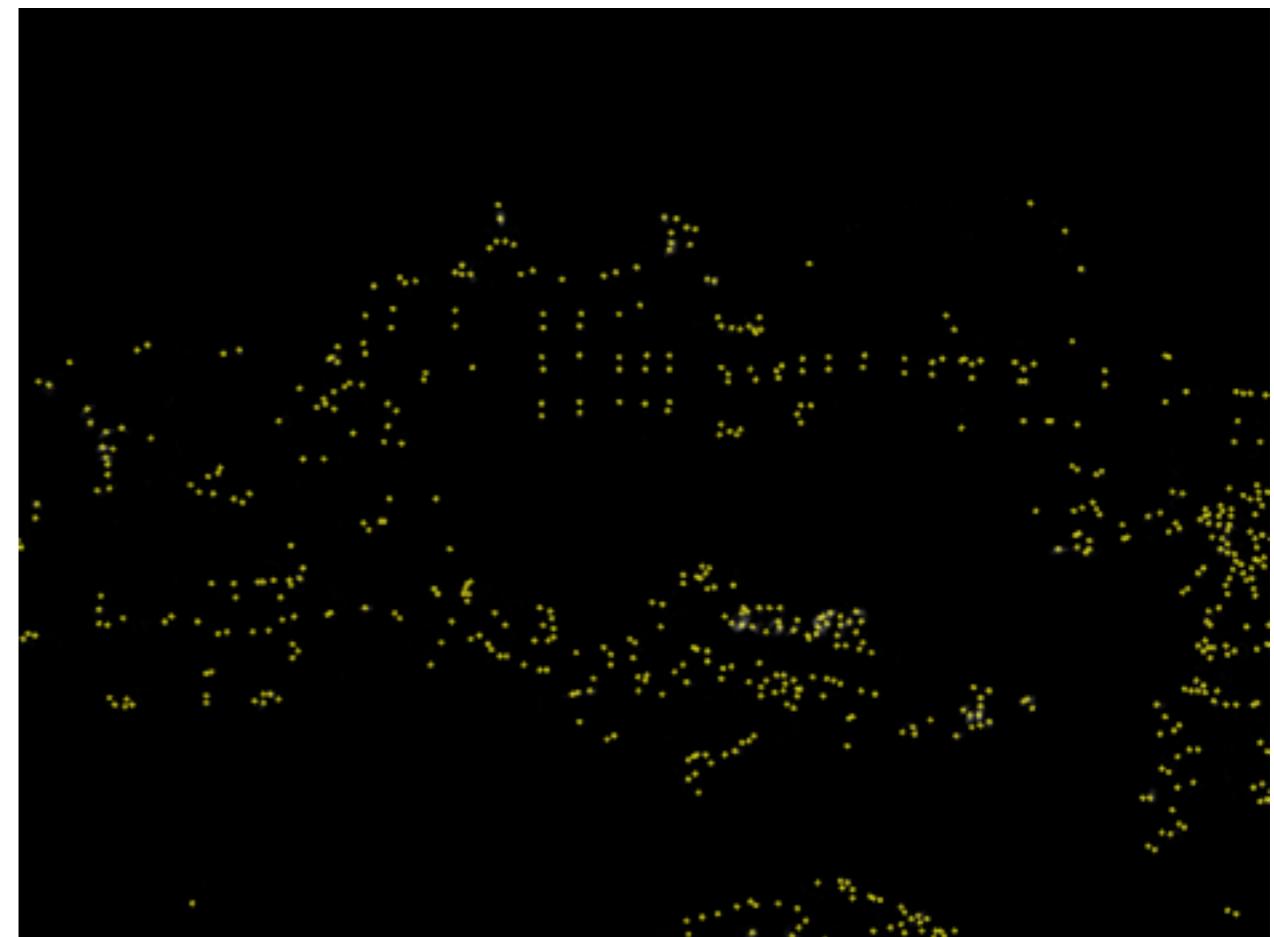
# Computing corner response

- Compute S matrix
  - ▶ Covariance of the gradients
$$S = \sum_i g_i g_i^T$$
- Compute eigen-values
  - ▶ Corner response = *smallest* eigen value
  - ▶ How bad (computationally) is this?
- Identify pixels with “good” corner responses
  - ▶ Thresholding
- A handful of practical issues!
  - Noise in original image
    - ▶ Creates false positives
    - ▶ Apply low-pass filter *first*
    - ▶ Also improves isotropicity of response
  - ▶ Local maximum suppression
- How big a patch to compute gradients over?



Kanade Tomasi

prefilt sigma = 2.3  
window size = 5  
thresh = 0.05



# Actually, I lied.

- There are two very closely related corner detectors
  - ▶ Kanade-Tomasi
    - what we described (uses eigenvalues)
  - ▶ Harris
    - identical, except uses (bad) approximation of eigenvalue.

$$\text{trace}(S) = \lambda_1 + \lambda_2$$

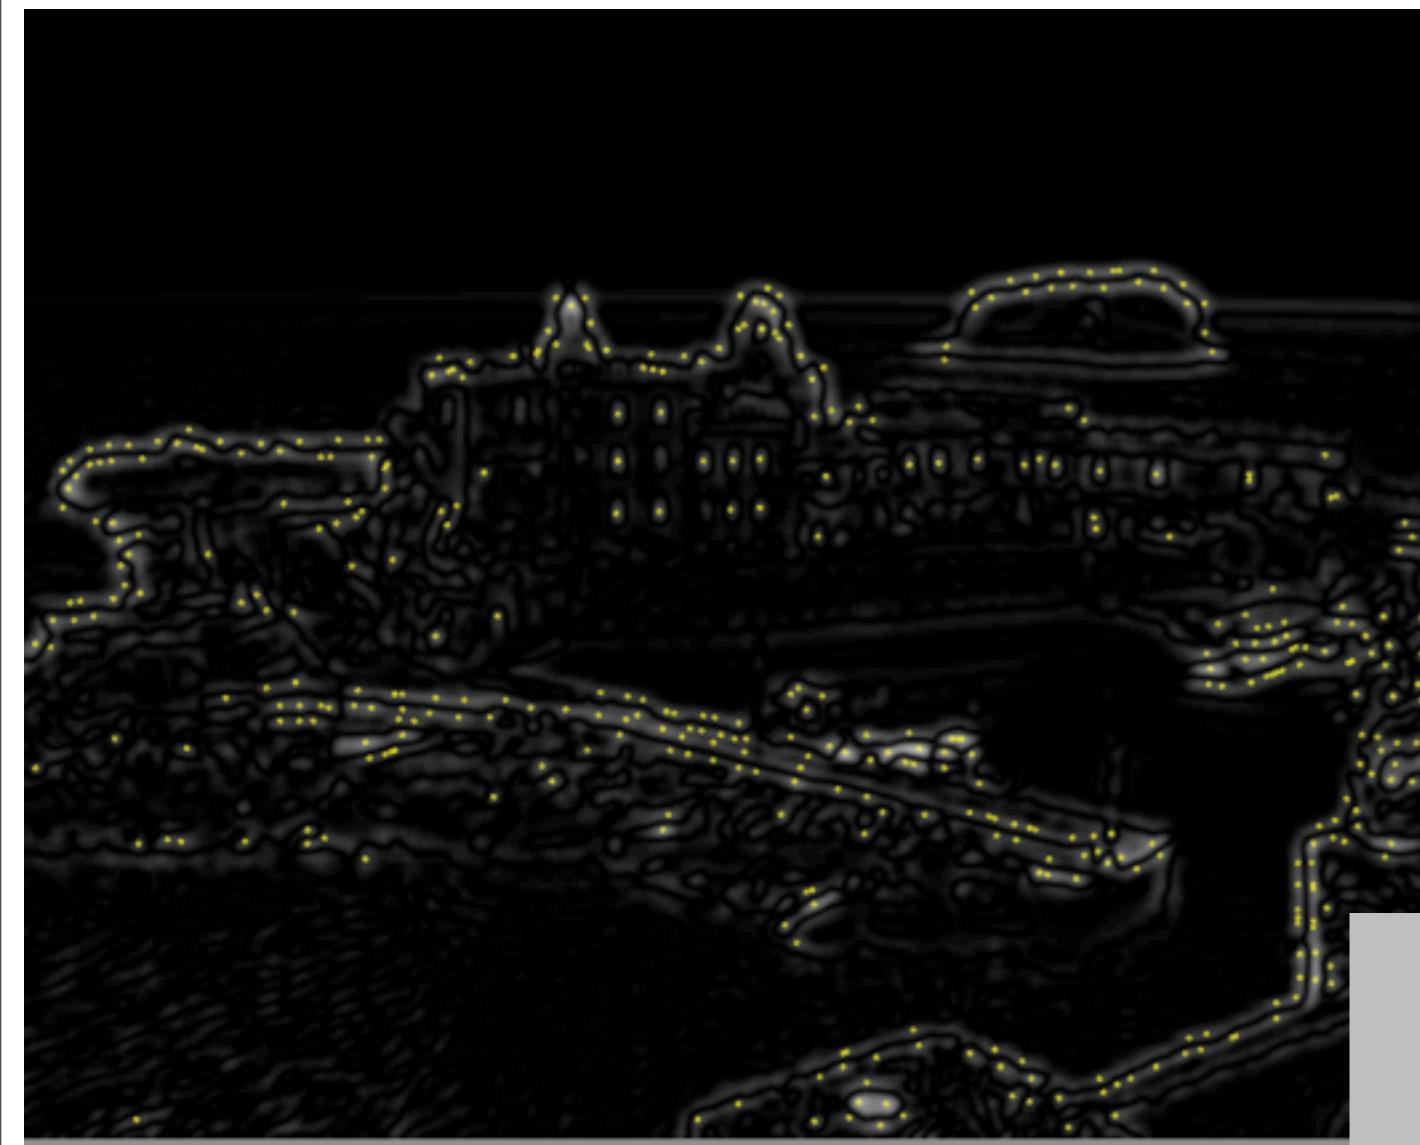
$$\det(S) = \lambda_1 \lambda_2$$

$$M = \det(S) - \kappa \text{trace}(S)^2$$

“Area of ellipse, minus a penalty for those that are highly eccentric.”

# Difference of Gaussians

- Another way of looking at corner detectors
  - ▶ Look for areas with high frequency in both directions
- What frequencies to look for?
  - ▶ We want a band pass
    - not too high (it's noise!)
    - not too low (it's not a corner!)
- Filter an image with two different Gaussians
  - ▶ Each corresponds to a low-pass filter
  - ▶ Difference corresponds to a band-pass filter



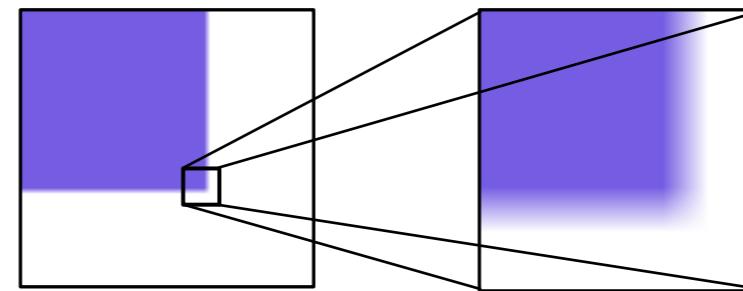
DoG, sigma1 = 5, sigma2 =  
7

Of course, we could  
evaluate this for  
different band passes.



# Multiple Scales

- Corners of high-resolution images are often blurry or noisy at fine levels of detail

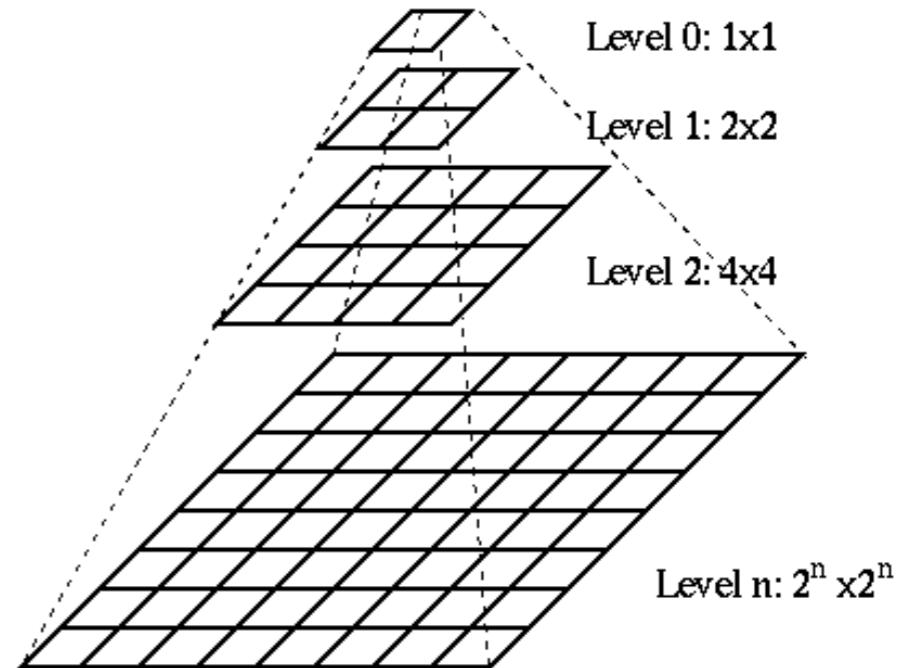


Corner indistinct at high resolution: no corner extracted

- Idea: run Harris corner detector on down-sampled versions of the image
  - ▶ Extract corners, blur, decimate, repeat.
- Idea: repeatedly compute DoG, increasing both sigma1 and sigma2
  - ▶ Look for successively lower frequency corners
  - ▶ Better yet, once we've band-limited "enough", we can decimate the image!

# Image Pyramids

- Look for features on multiple scales
  - ▶ Just repeat image processing algorithm on successively lower-resolution images
  - ▶ Must produce lower-resolution images
- Avoiding aliasing requires low-pass filters
  - ▶ Ideal low-pass filter?
  - ▶ Don't create new features when filtering
    - Avoid ringing!
    - Want a monotonic filter



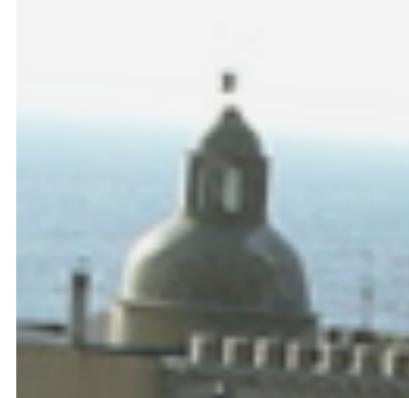
Can be shown that the only admissible filter is a Gaussian low-pass filter. However, can't achieve perfect frequency response... some tradeoffs necessary when building filter

# Feature Tracking

- We often want to track (or match) features across two frames.
  - ▶ Which corners in image A match those in image B?
  - ▶ i.e., data association
- Can we use *more* information?
  - ▶ Why not use the local appearance?

# Image patch patching

- Consider the pixel patch around a feature
  - ▶ Sum of absolute/squared (SAD/SSE) differences/errors
- How robust is this to small alignment errors/rotations/changes in viewpoint/etc.?



These will probably match



These probably won't

# Invariances

- Our goal: detect distinctive features, maximizing repeatability
  - ▶ Transform pixel patch into a space where a simple comparison (SAD/SSE) *is effective*.
- Scale invariance
  - ▶ Robust to changes in distance
- Rotation invariance
  - ▶ Robust to rotations of camera
- Affine invariance
  - ▶ Robust to tilting of camera
- Brightness invariance
  - ▶ Robust to minor changes in illumination

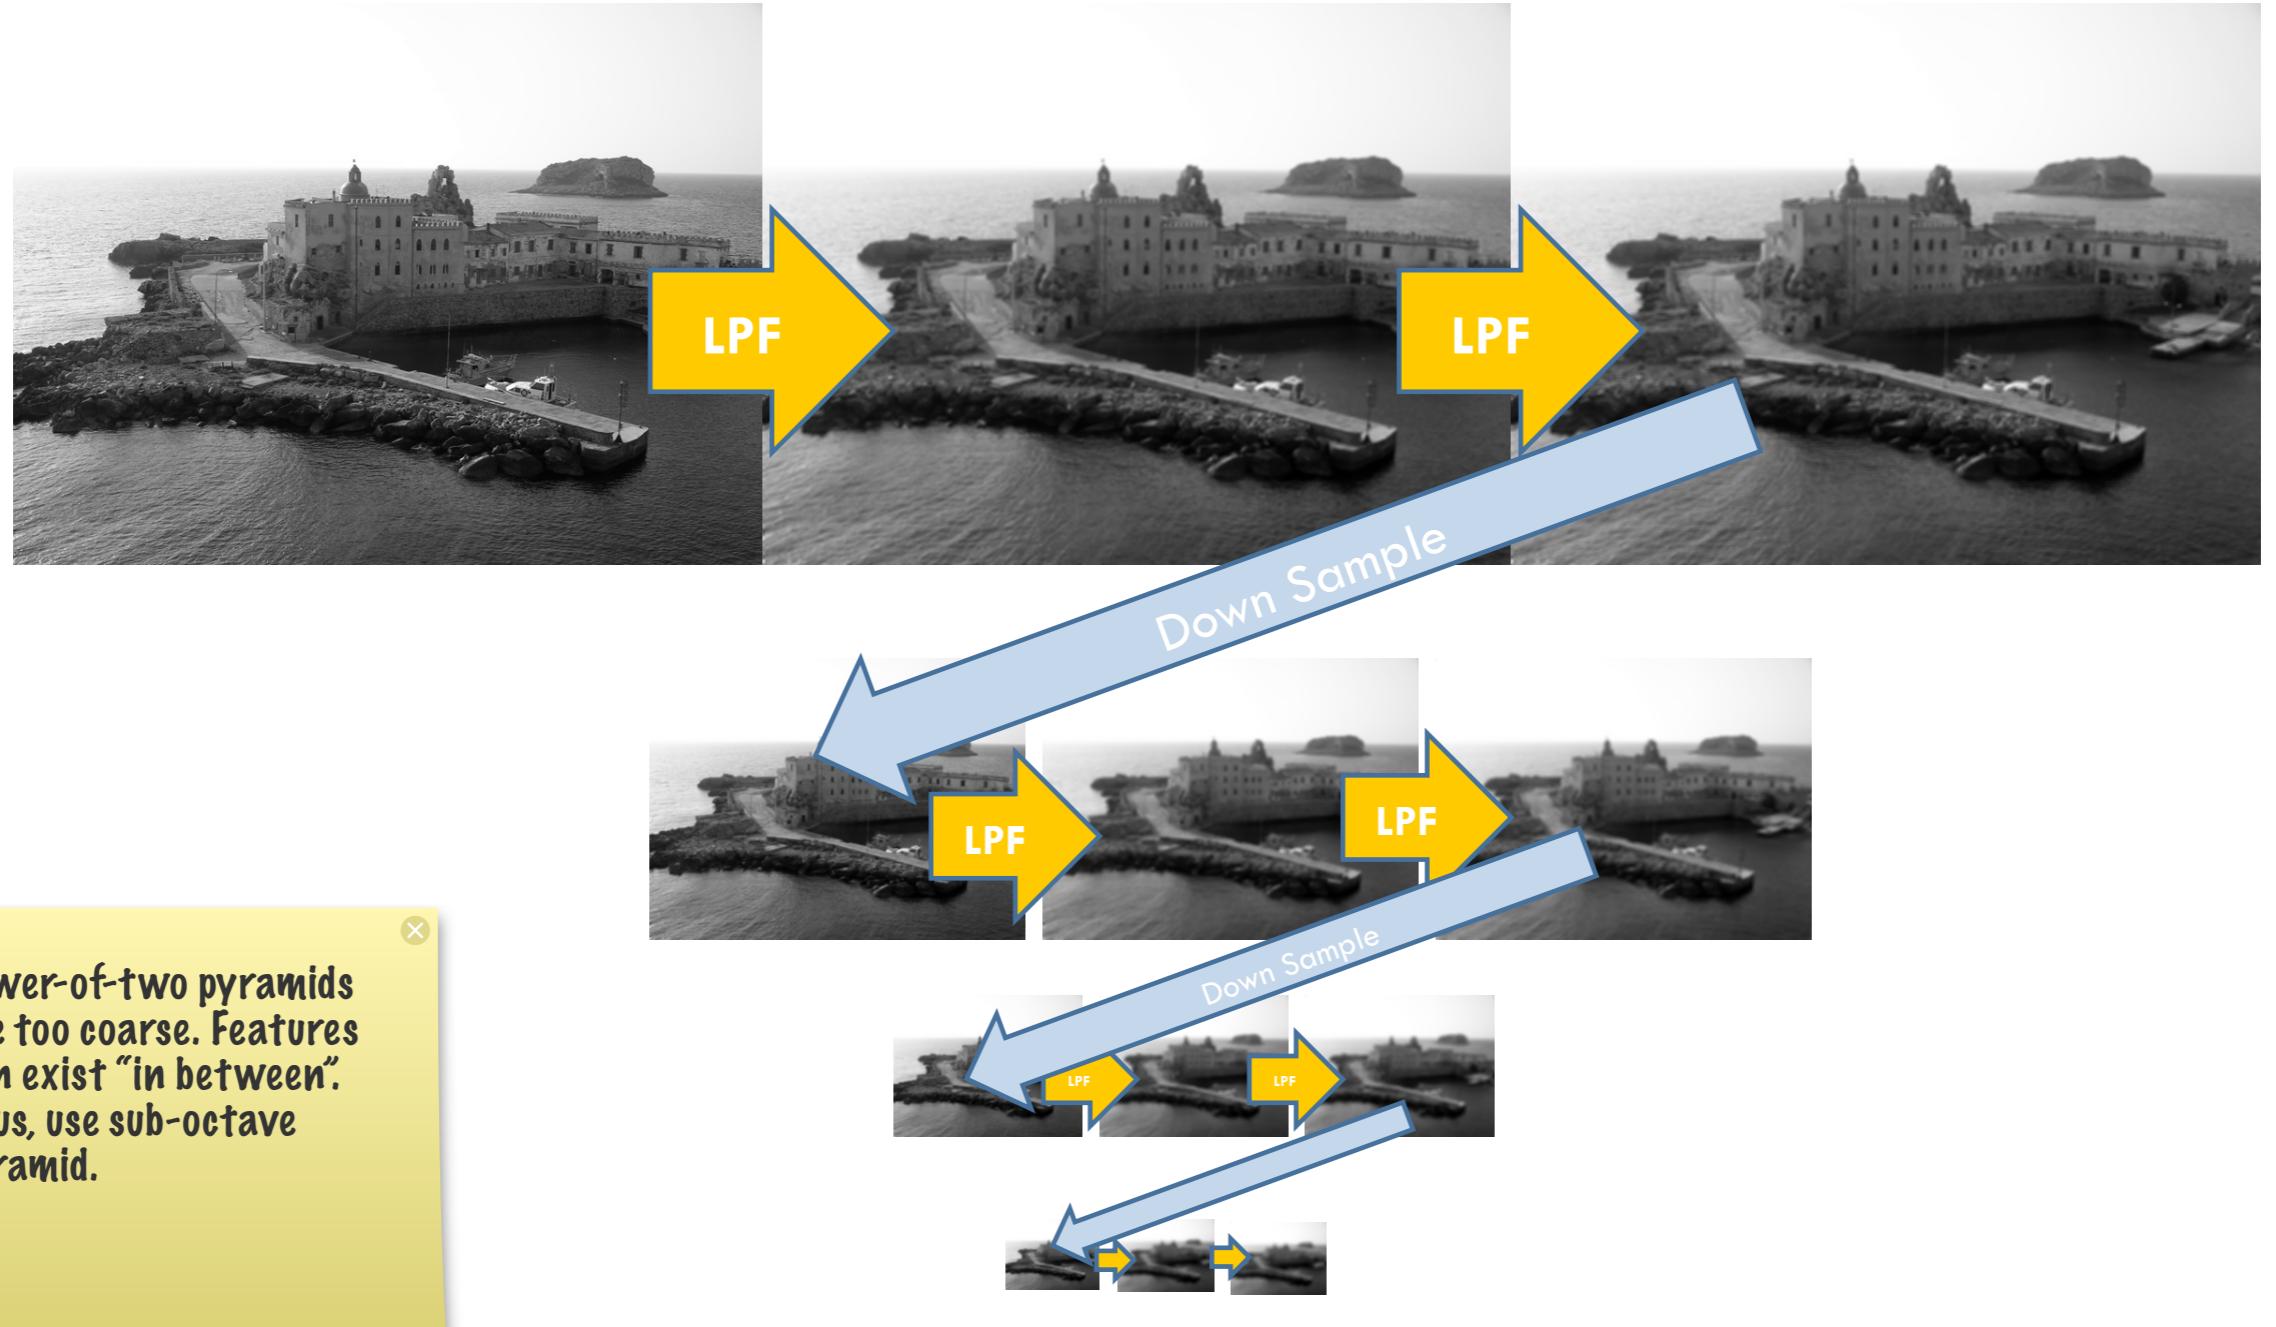
# SIFT: Scale-Invariant Feature Transform

- David Lowe (Univ. British Columbia)
- Probably the single most commonly used tool in computer vision
  - ▶ For better or for worse... often used “reflexively” even if it’s not a good choice!
- Watch out!
  - ▶ Patented, commercial use restricted

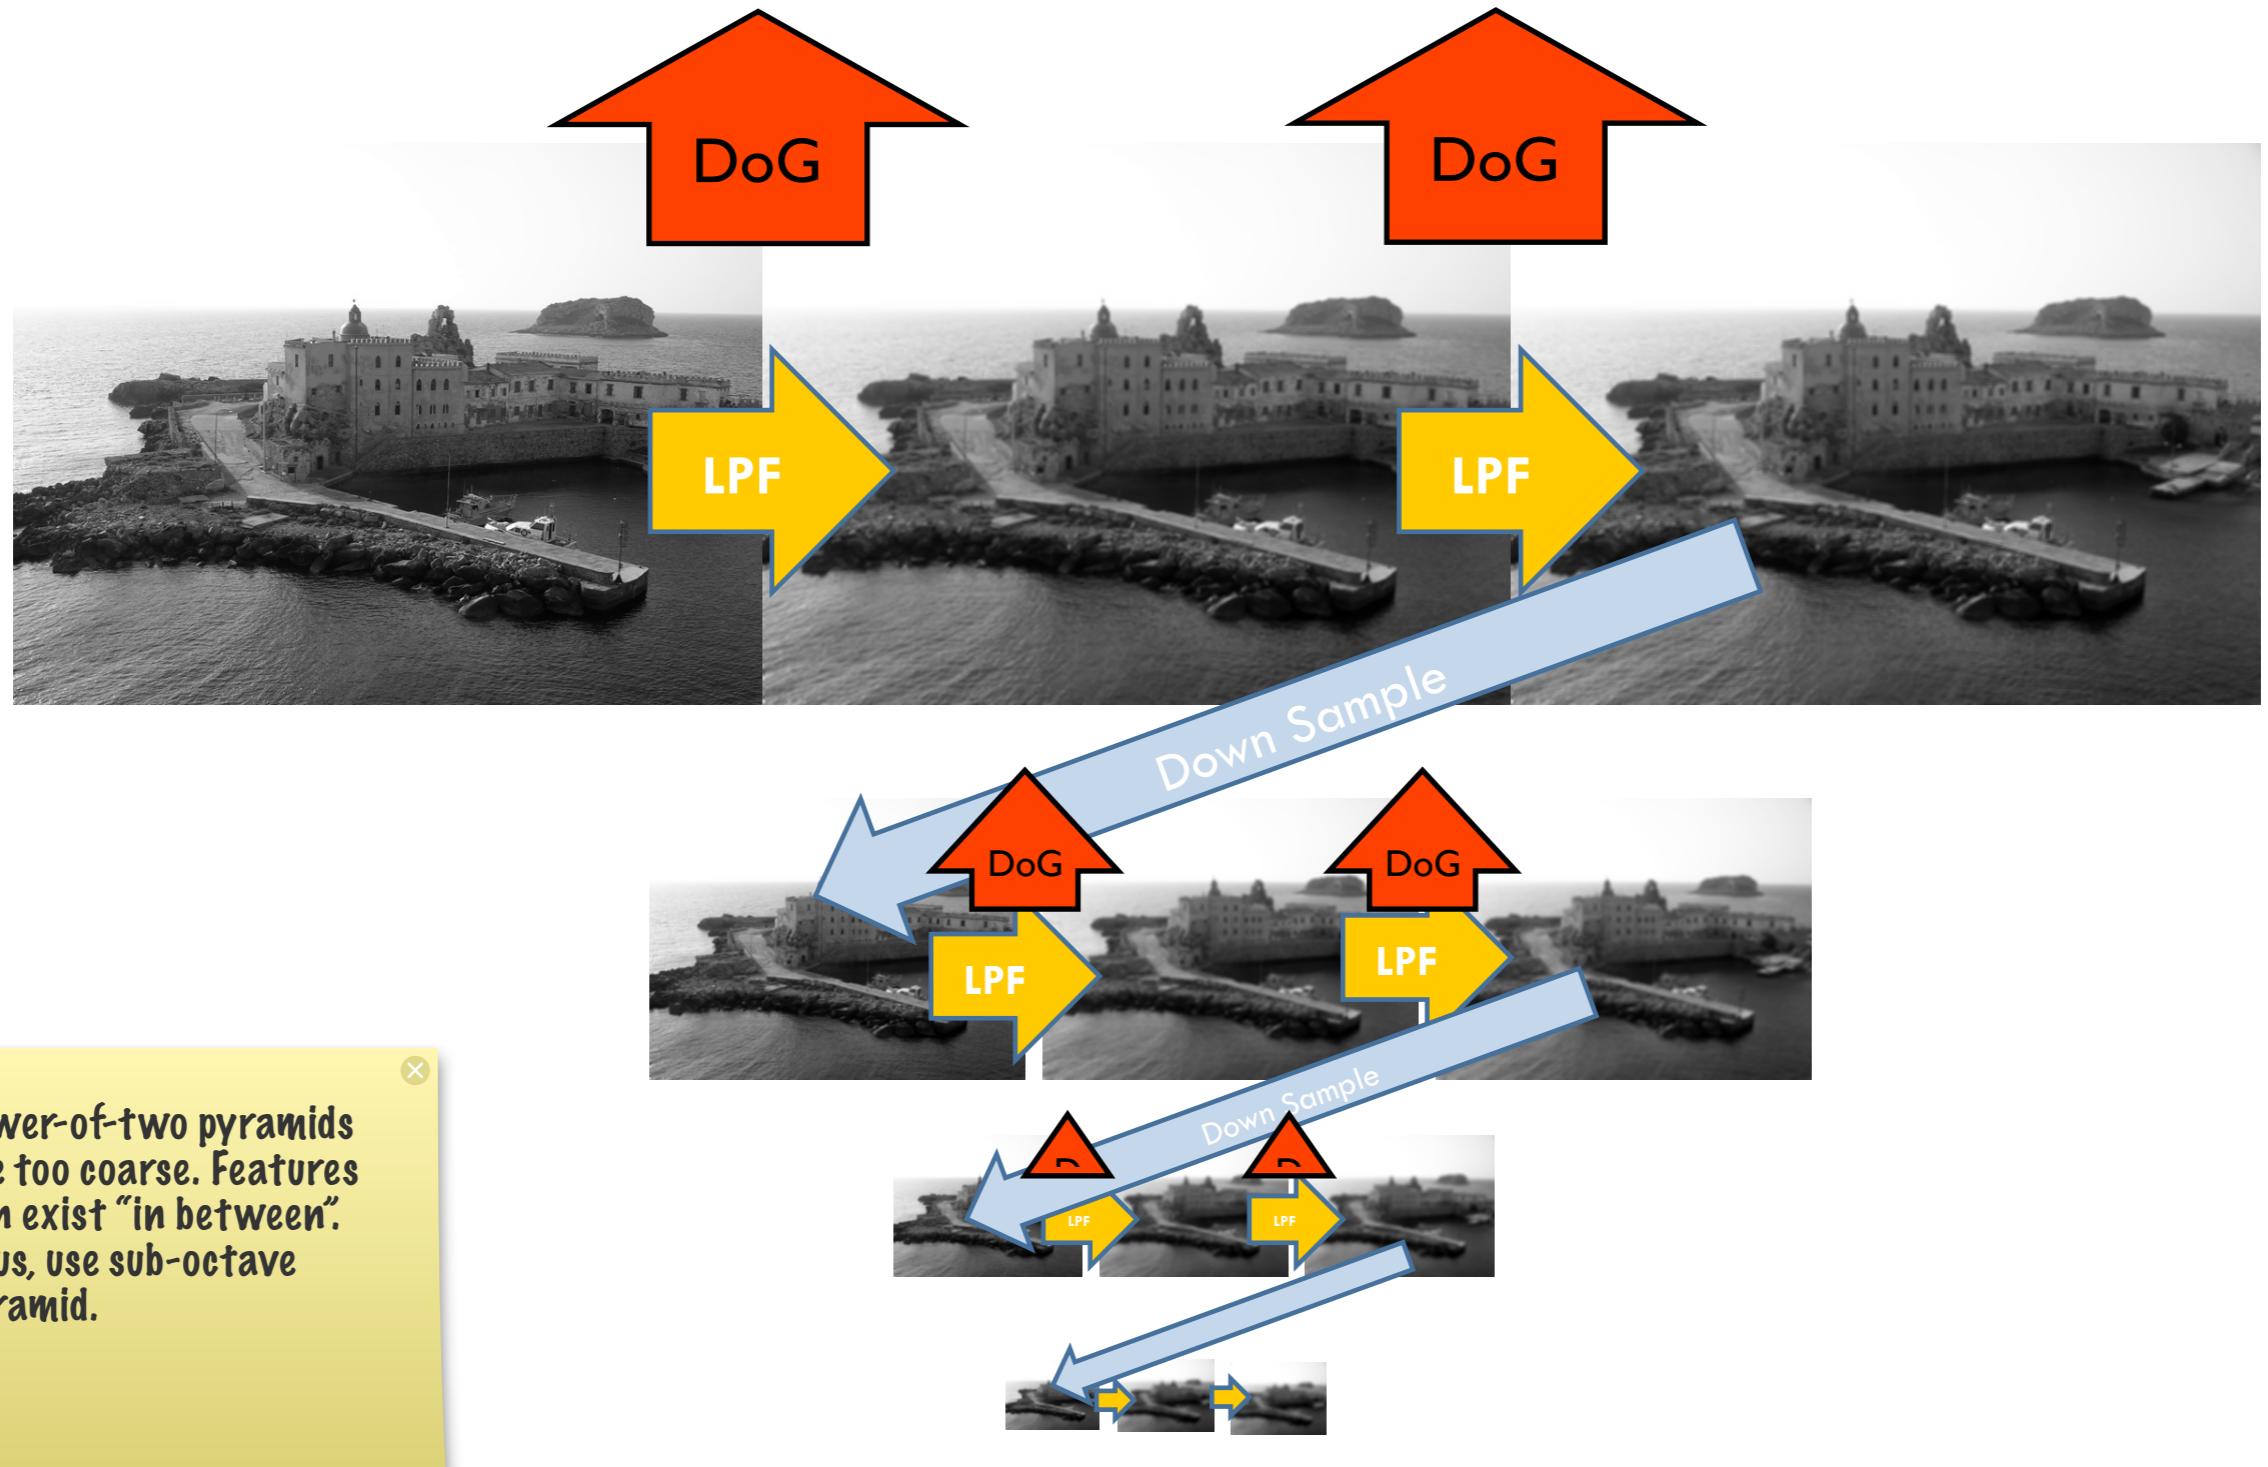
# SIFT

- Detect interest points
  - ▶ Image pyramid using DoG “corners”
  - ▶ Output: corners *and* scale (which level of the pyramid?)
- Output a “descriptor”
  - ▶ Consider pixel match around corner
  - ▶ Compute a histogram of the gradient directions
  - ▶ “Rotate” the histogram so that the dominant direction is first.

# Sub-Octave Image Pyramids

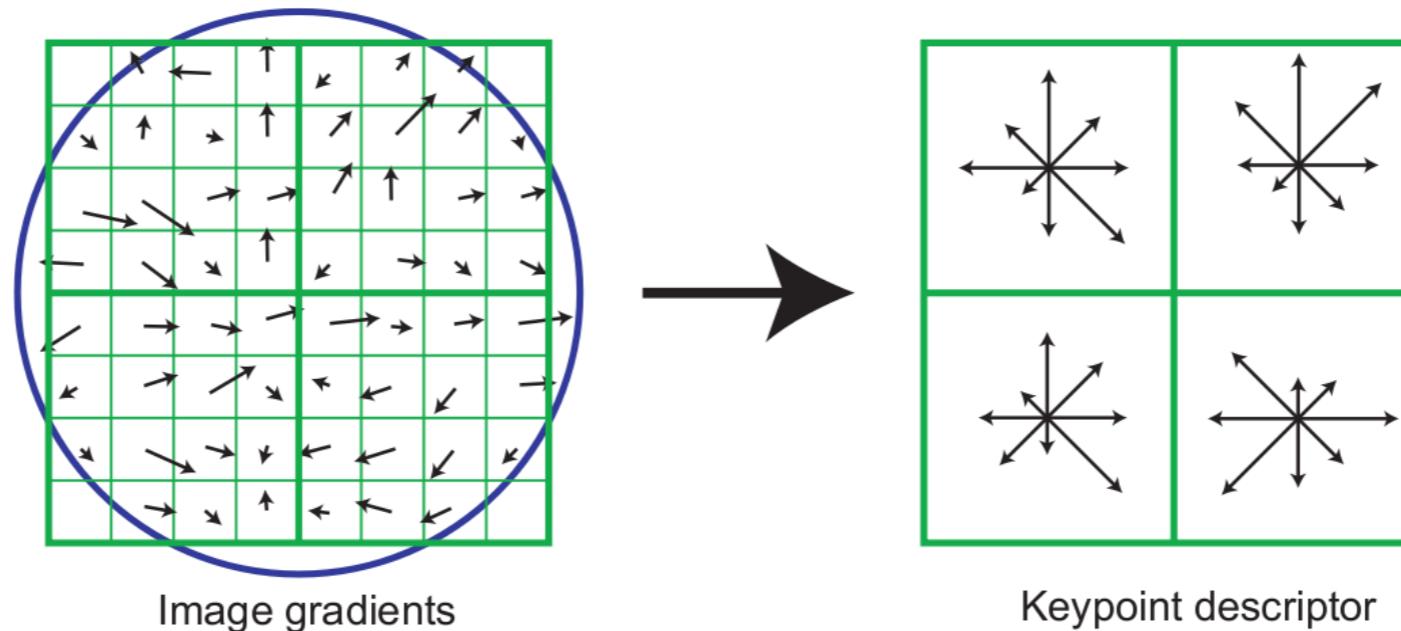


# Sub-Octave Image Pyramids



# SIFT Descriptor

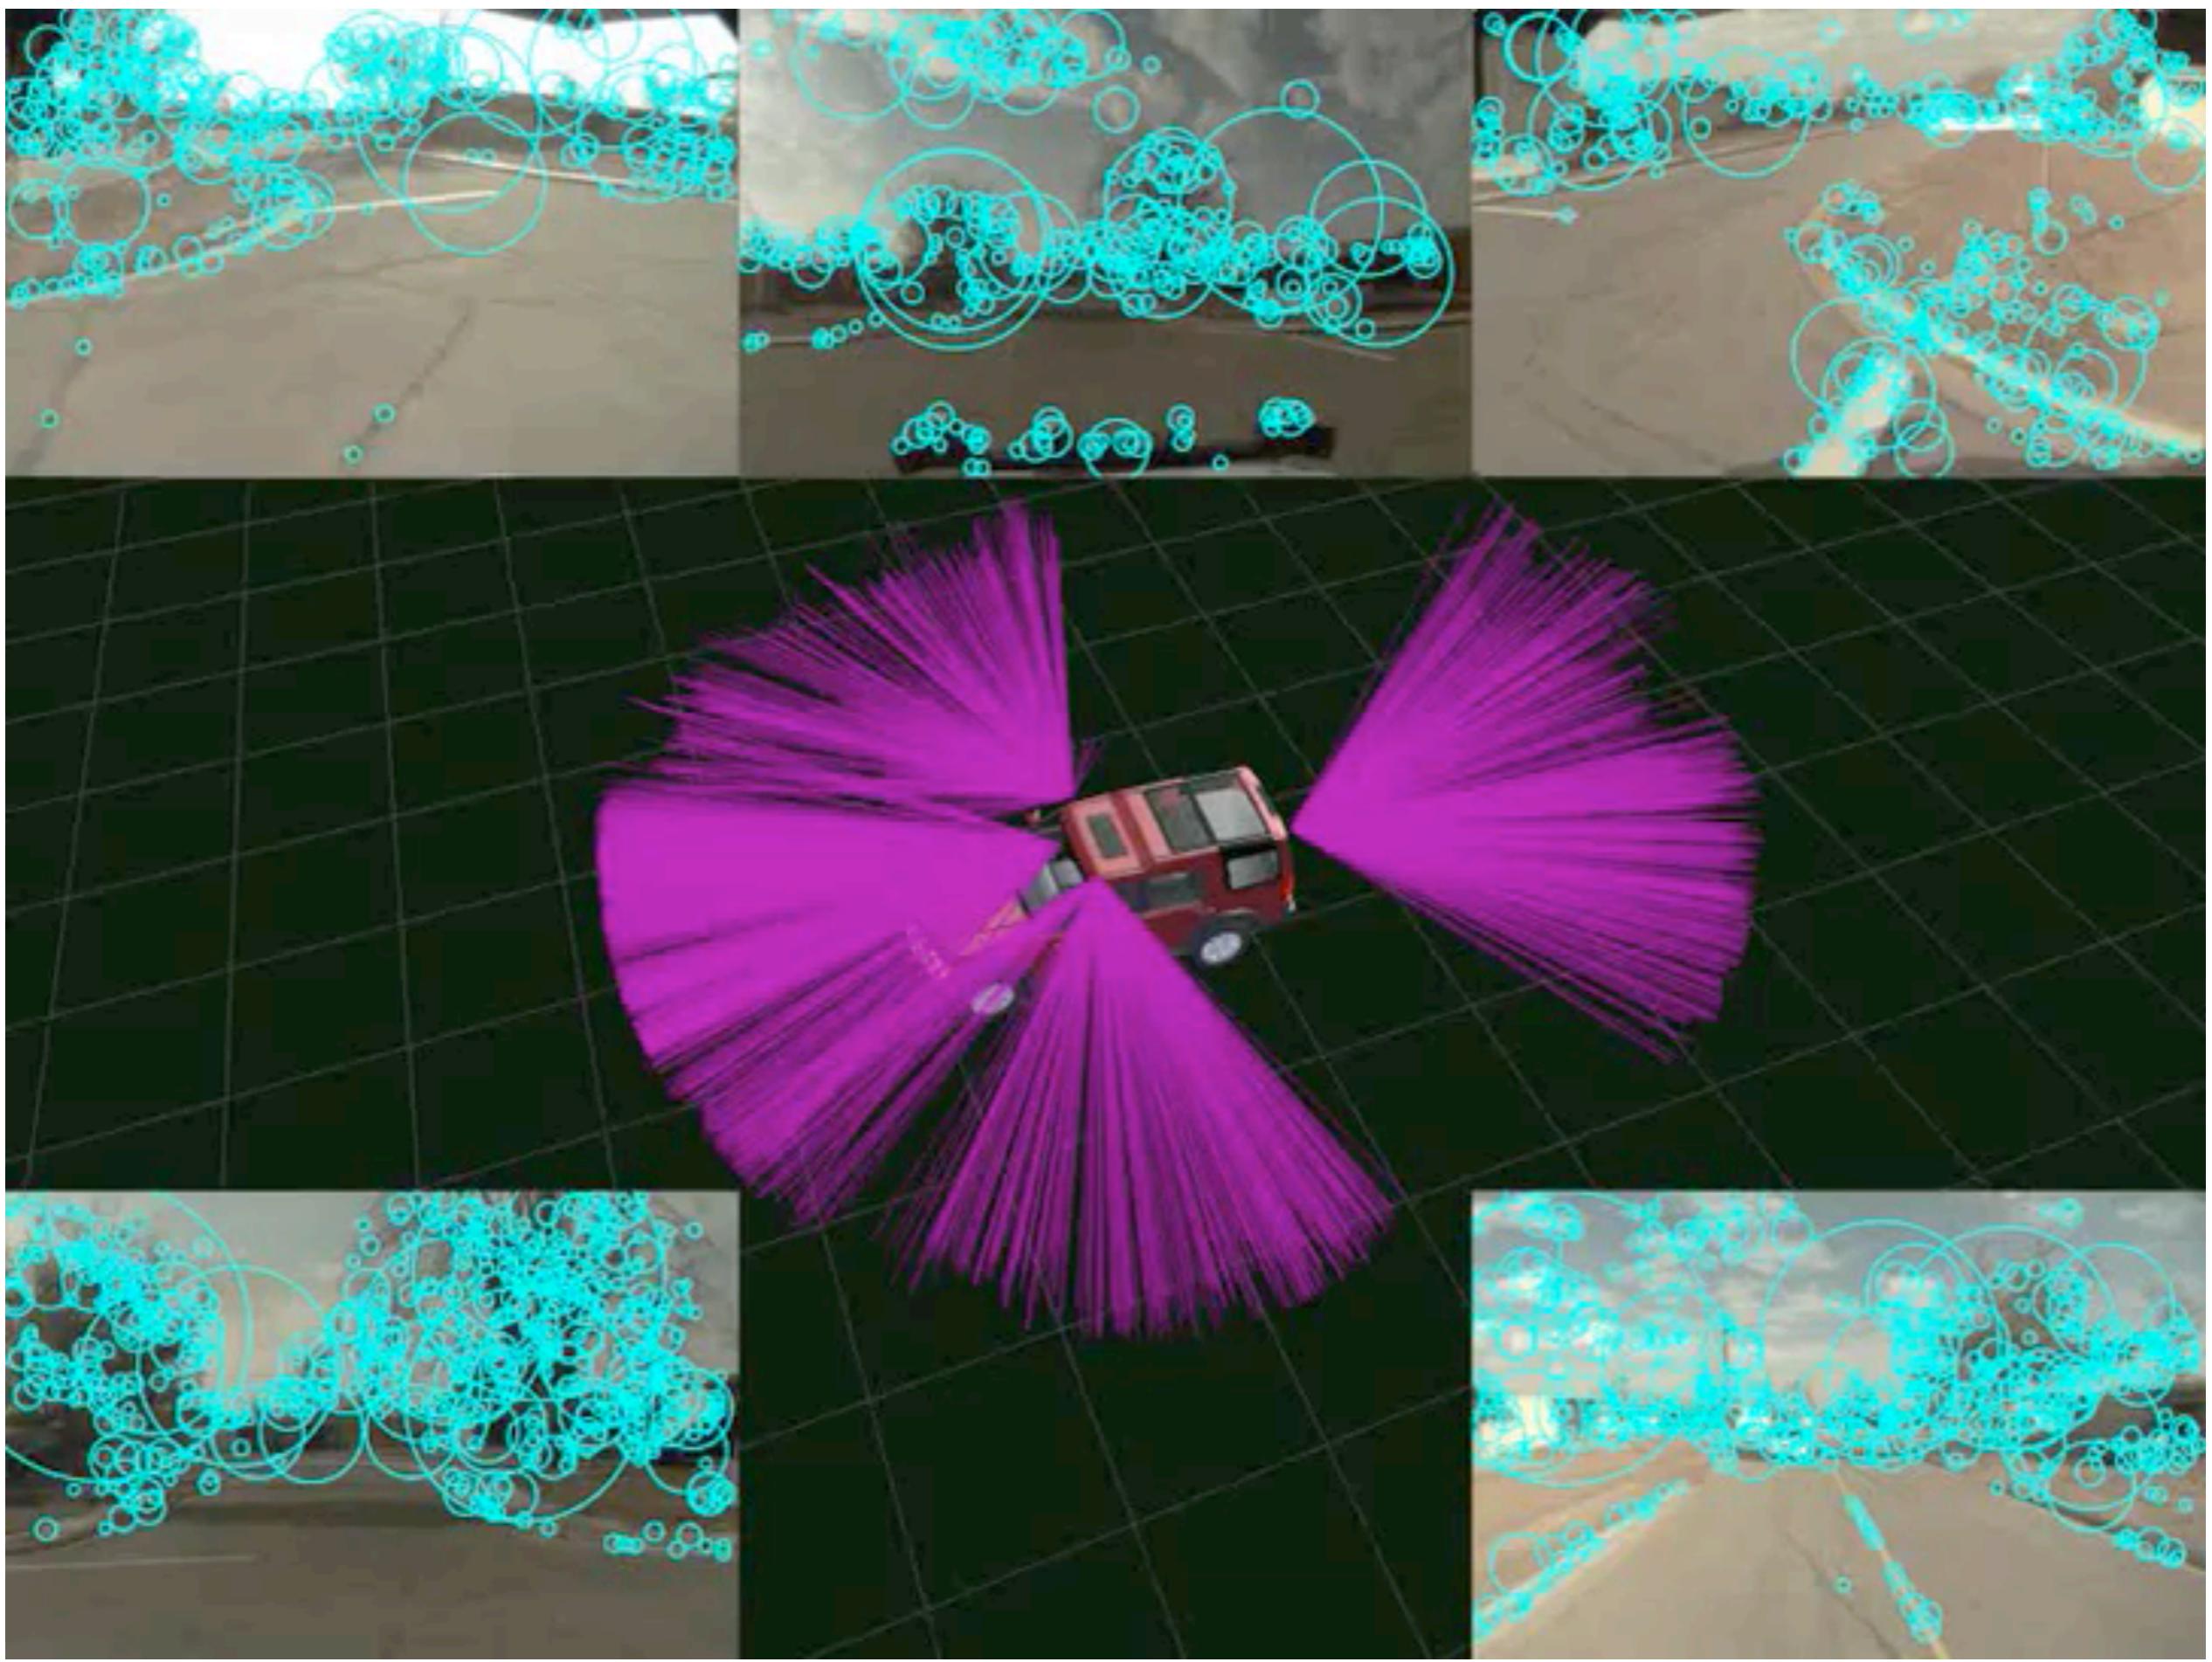
- Histogram of gradients gives good information about a pixel patch
  - ▶ But building just one histogram loses a lot of spatial information.
  - ▶ Idea: For a given interest point, compute a set of histograms; output each.
  - ▶ Shift histograms so dominant direction is first in histogram ==> rotational invariance.



- “Official” SIFT uses 16x16 pixel patches, 4x4 bins, 8 histogram buckets
- How many degrees of freedom in SIFT descriptor?
  - ▶ # bins \* # histogram buckets =  $4*4*8 = 128$

# Matching SIFT Descriptors

- Each SIFT feature:
  - ▶ (x,y,scale) (ignore scale if you want scale invariance!)
  - ▶ descriptor[128]
- Two descriptors can be compared using Euclidean distance...
  - ▶ Small distances = similar descriptors
  - ▶ What if same/similar feature appears more than once? nearest neighbor may not be good enough
- Common approach:
  - ▶ Suppose best match for  $A_i$  is  $B_j$  (with  $d_{ij}$ ).
  - ▶ Suppose next best match for  $A_i$  is  $B_k$  (with  $d_{ik}$ ).
  - ▶ Require  $d_{ij} < \alpha d_{ik}$ . ( $\alpha$  typically 0.8).
- “Marriage” constraint:  $A_i$  and  $B_j$  match only if  $B_j$  is the best feature for  $A_i$  and vice versa.





# Object recognition

- SIFT also used to build object recognition systems



# Artificial Features



# Applications

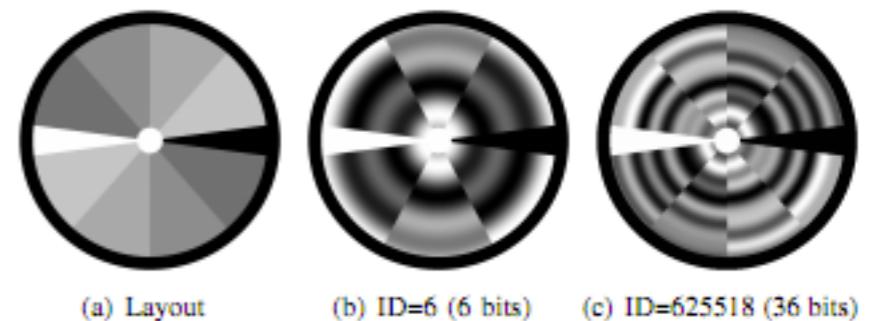
- Ground truthing
- Recognizing robots
- Commanding robots
- Education
  - ▶ Often useful to bypass open-ended perception problems





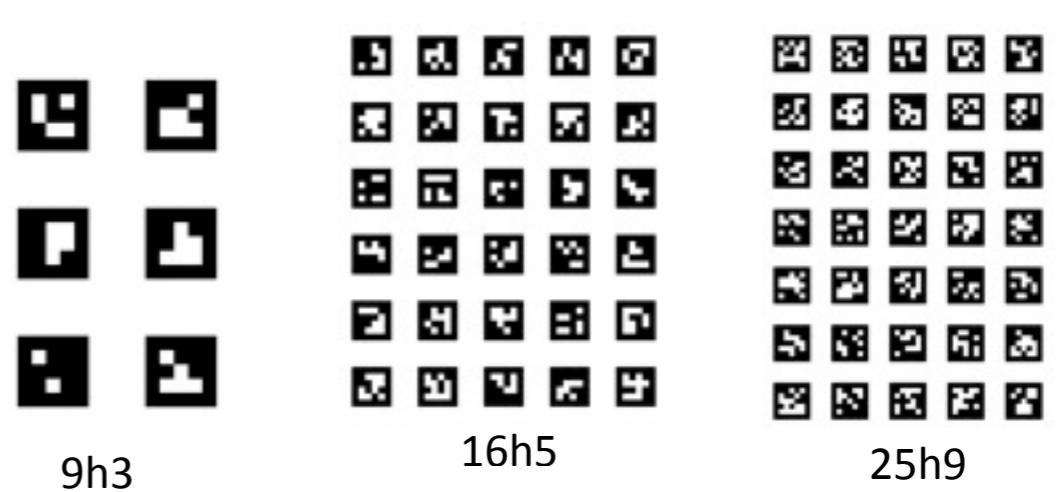
# Related Work

- ARToolkit
  - ▶ Widely used
  - ▶ Primitive binarization scheme => high failure rate in unstructured environments
  - ▶ Weak coding system
  - ▶ Freely available
- ARTag (Fiala, 2005)
  - ▶ Seems to address many shortcomings in ARToolkit
  - ▶ Methods are not well-documented
  - ▶ Source code not available
- Bokode (Mohan et al, 2009)
- Fourier codes (Sattar et al, 2007)
- Quick Response (QR) Tags

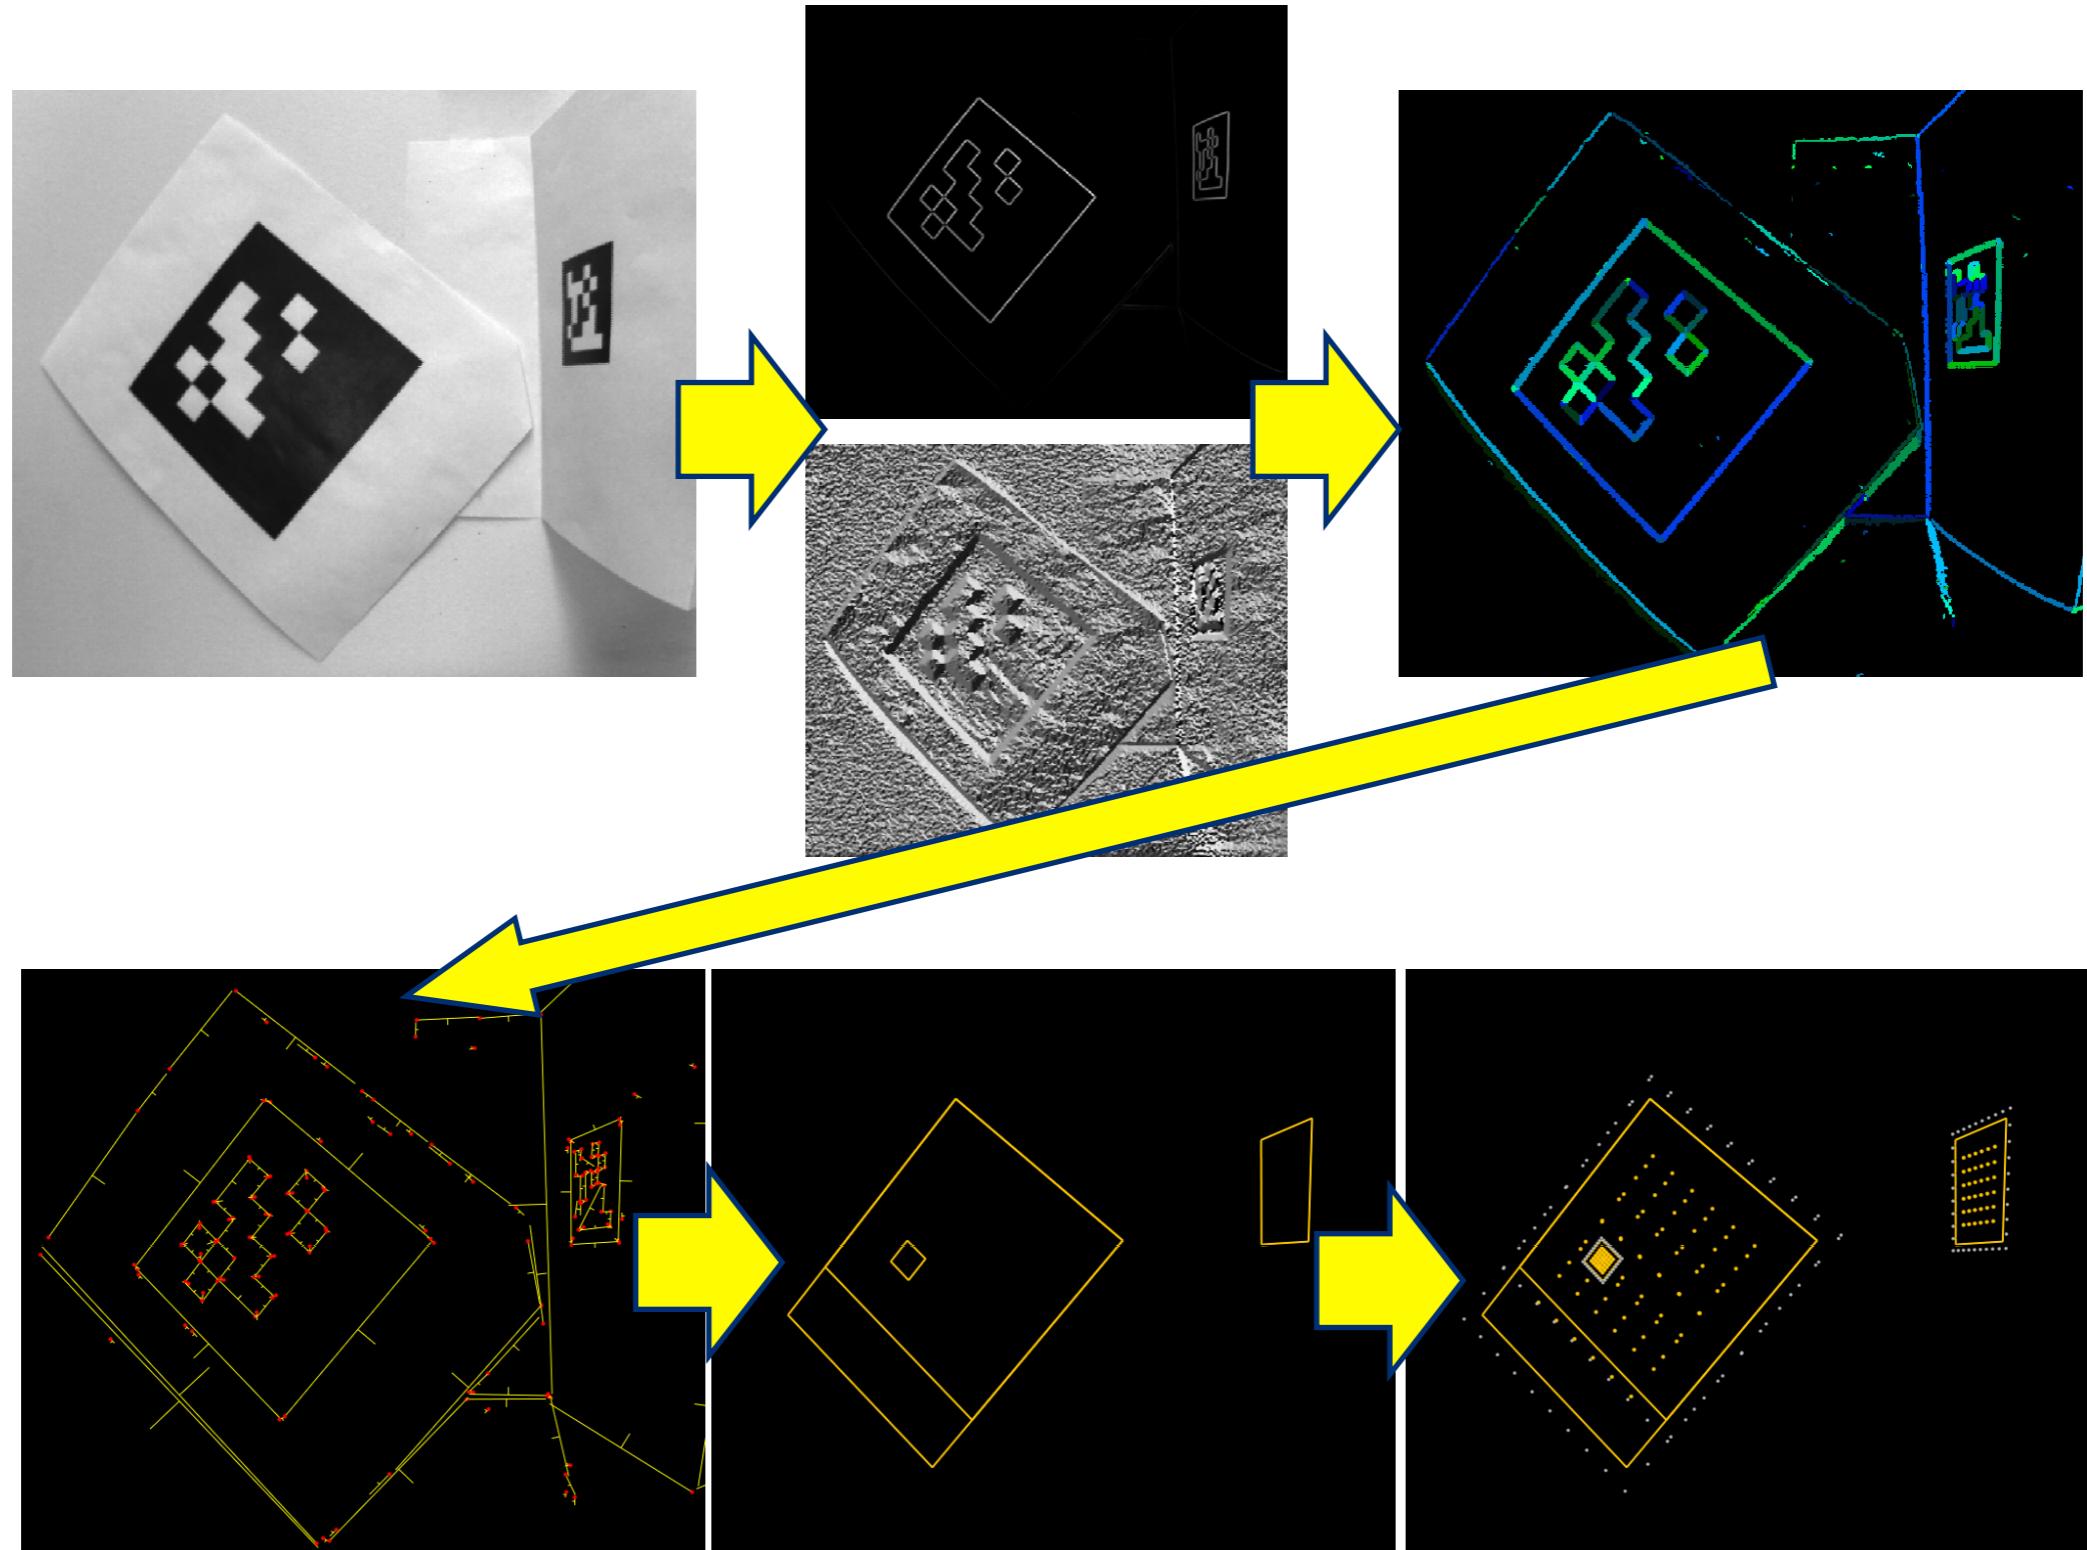


# AprilTags

- Robust detection
  - ▶ Not based on threshold-based binarization scheme
  - ▶ Works better in unstructured environments
  - ▶ Accurate localization
- Strong coding system
  - ▶ Low false positive rate
- Parameterizable
  - ▶ Pick your own tag family



# Detection Approach





# Coding System

- Based on lexicographic code
  - ▶ Nearly optimal coding family

- Simple algorithm:

```
codebook = {}  
for i = 0 : max-codeword  
    if (hamming_distance(i, codebook) > H))  
        codebook = codebook U i  
return codebook
```

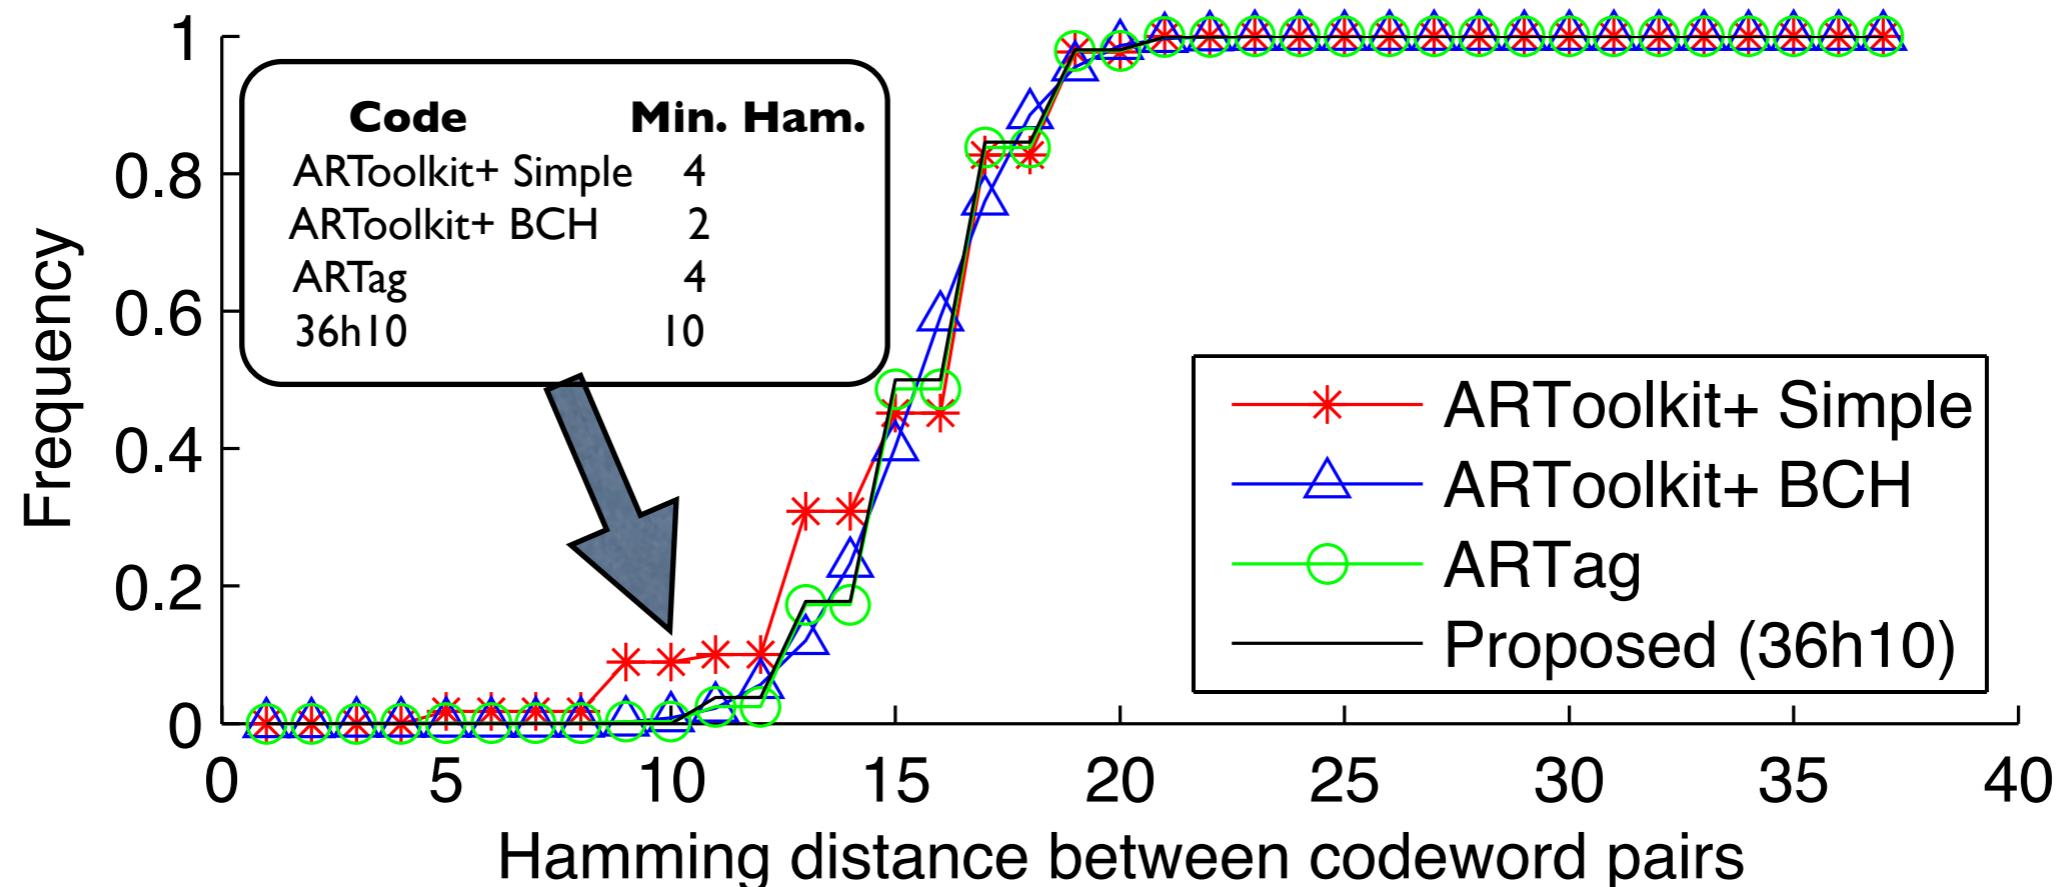
# Coding System

- But we need rotational invariance too!

```
codebook = {}  
for i = 0 : max-codeword  
    if (hamming_distance(i, codebook) > H))  
        codebook = codebook U { i, rot90(i), rot180(i), rot270(i) }  
return codebook
```

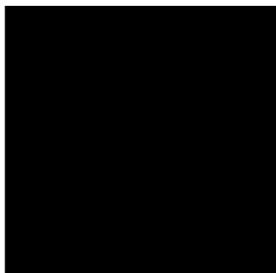
- Note how code generation system can be easily modified to incorporate additional constraints.

# Coding System: Optimality

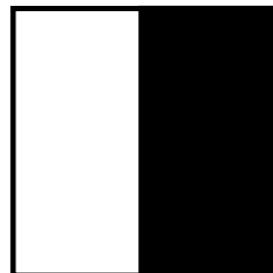


# Tag Complexity

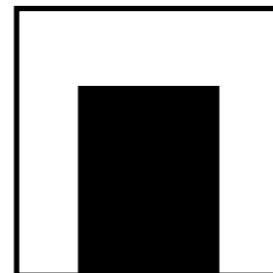
- Coding scheme can generate “perceptually weak” tags
  - ▶ All black tag--- likely to appear by chance in images
  - ▶ ARTag manually identified two “bad” tags and rules them out.
- How can we measure the “badness” of a tag automatically?
- Idea: How many rectangle drawing operations would it take to generate a tag?
  - ▶ Simple greedy search computes upper bound



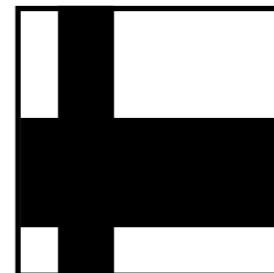
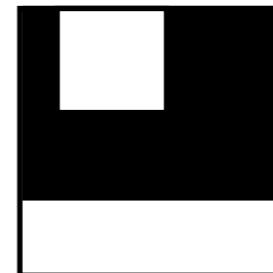
c=1



c=2



c=3



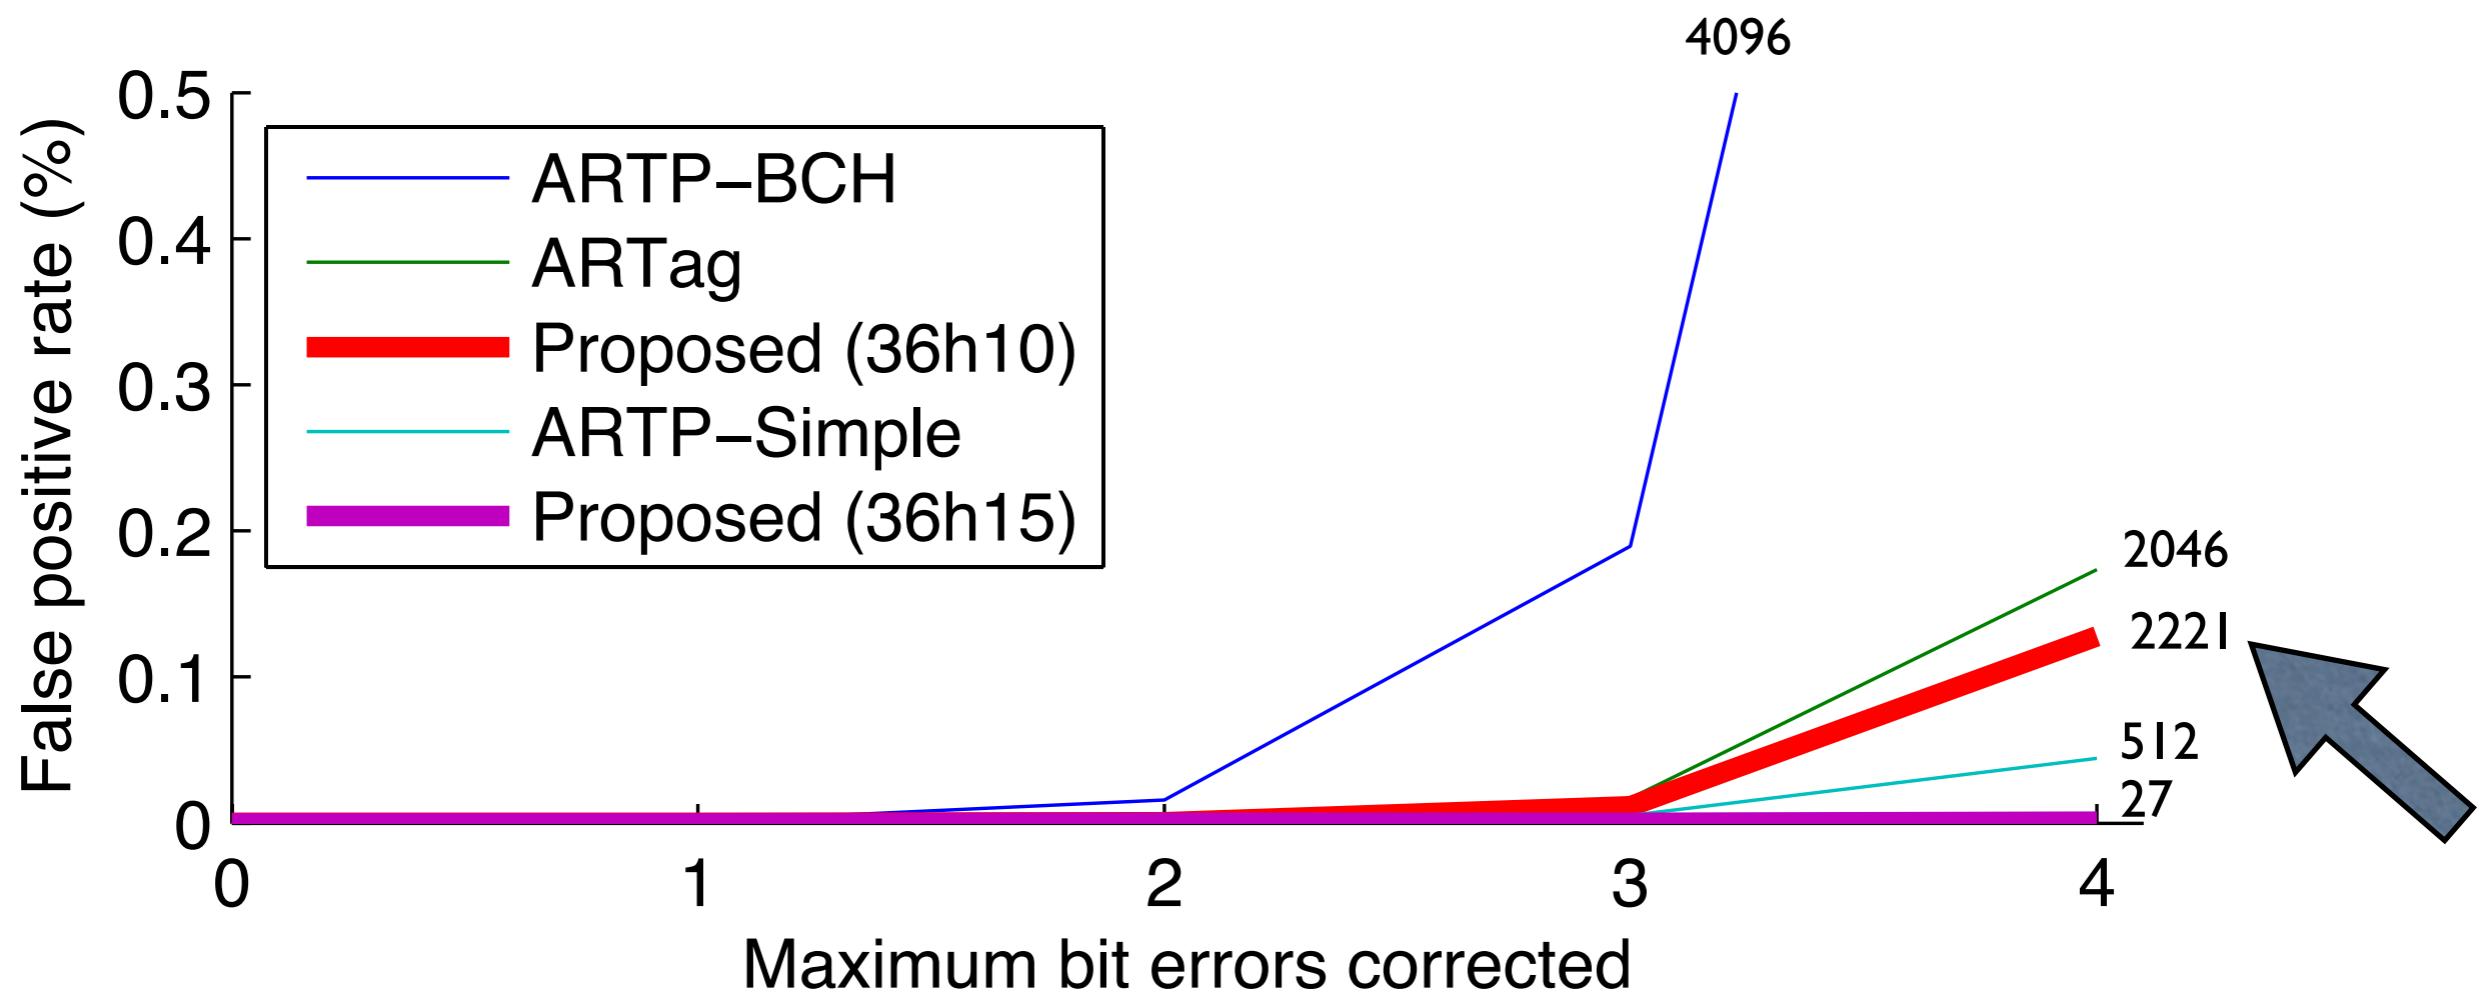
# Results: False Positives

- Use large corpus of images *not* containing AprilTags. Do we detect any?
- Which dataset should I use?
  - ▶ Should convince you that I haven't "cooked" it
  - ▶ Used open "Label Me" dataset
    - Huge! (180,829 images)
    - Wide variety of topics, cameras, image quality



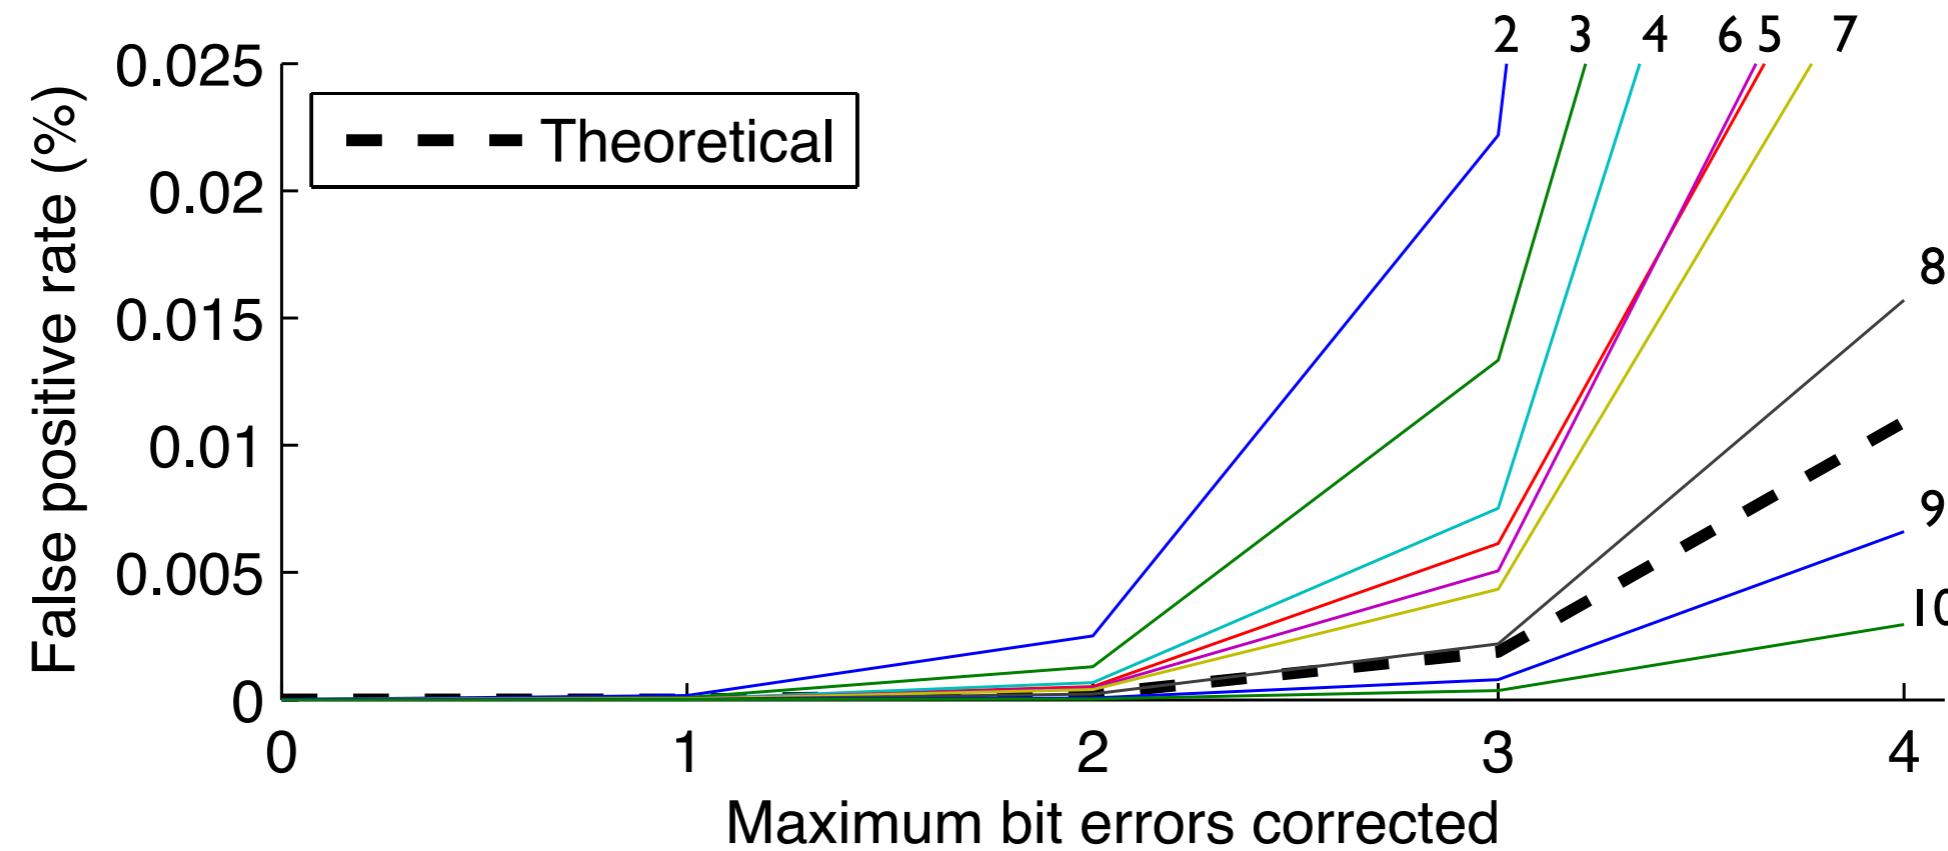
LabelMe (Russell 2005)

# False Positives (Label Me)



AprilTag 36h11 outperforms ARTag: more encodable tags at lower false positive rate!

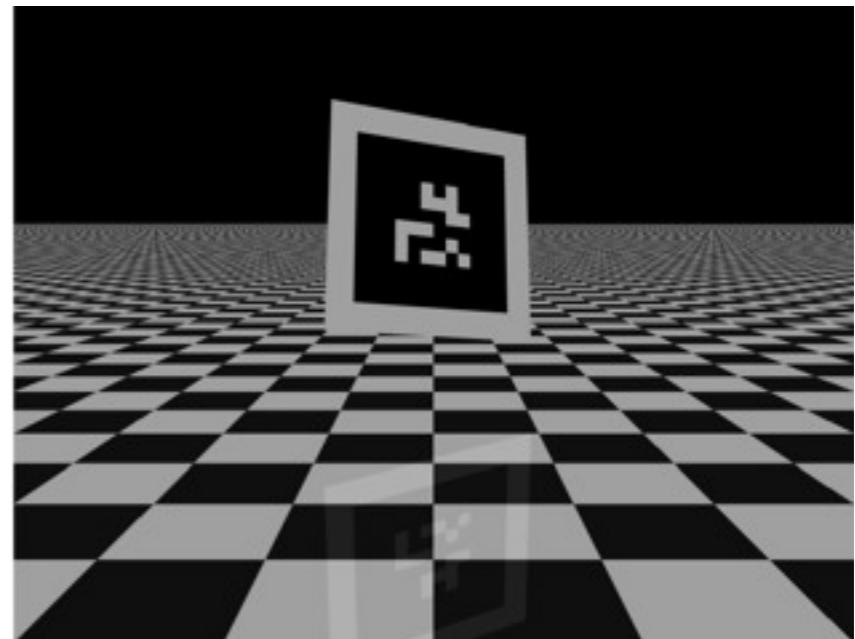
# False positives versus complexity: Label Me



- Conclusion: Our tag complexity heuristic *does* help us reject naturally-occurring patterns
  - At  $c > 8$ , our tags are less likely to appear in real-world images than completely random tags!

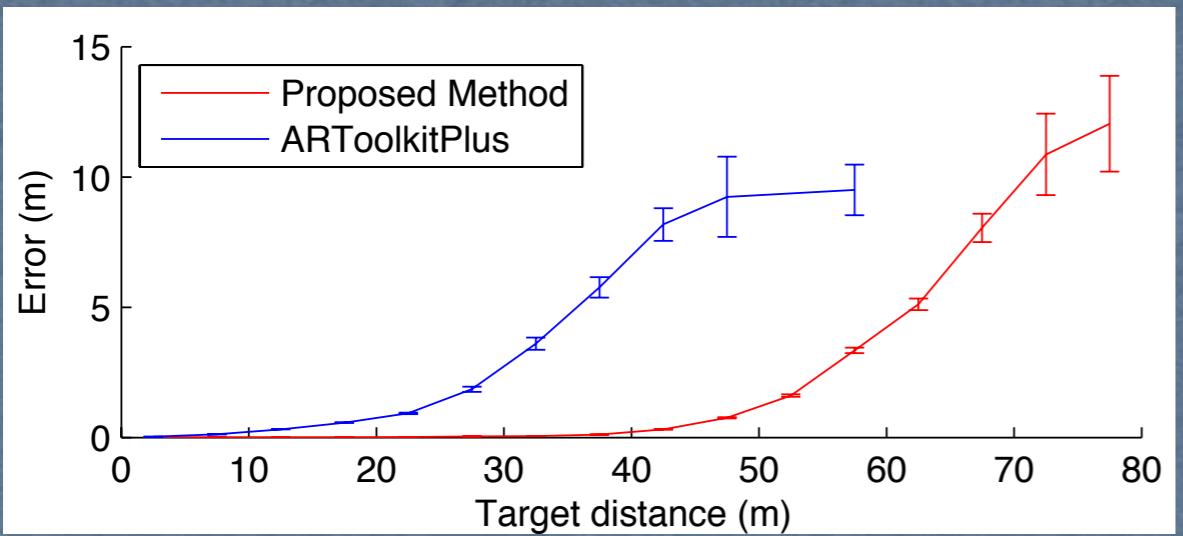
# Evaluation: Accuracy

- For reliable ground-truth, used synthetic ray tracing images
  - ▶ Correct answer known exactly
- Two main factors in detection accuracy
  - ▶ Distance from camera
  - ▶ Angle to camera
- For each particular experiment, we want to know
  - ▶ Accuracy
  - ▶ Detection rate

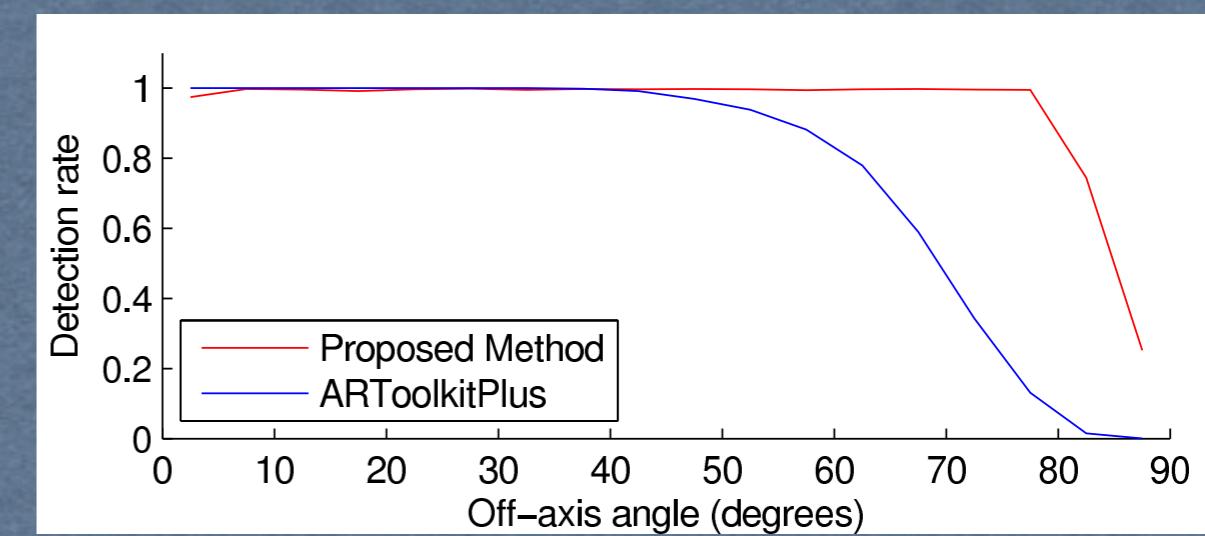
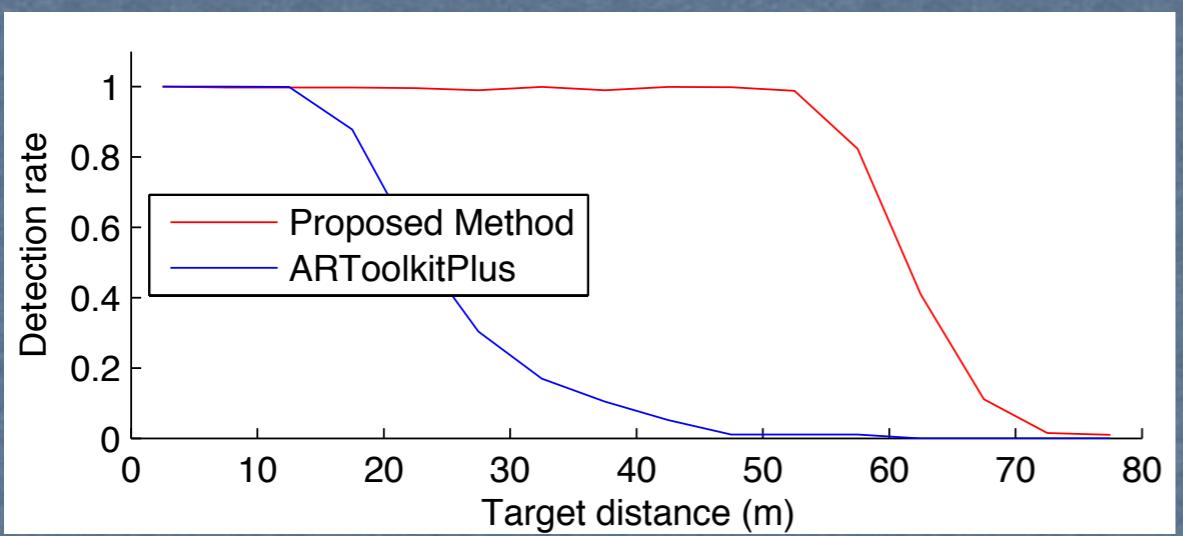
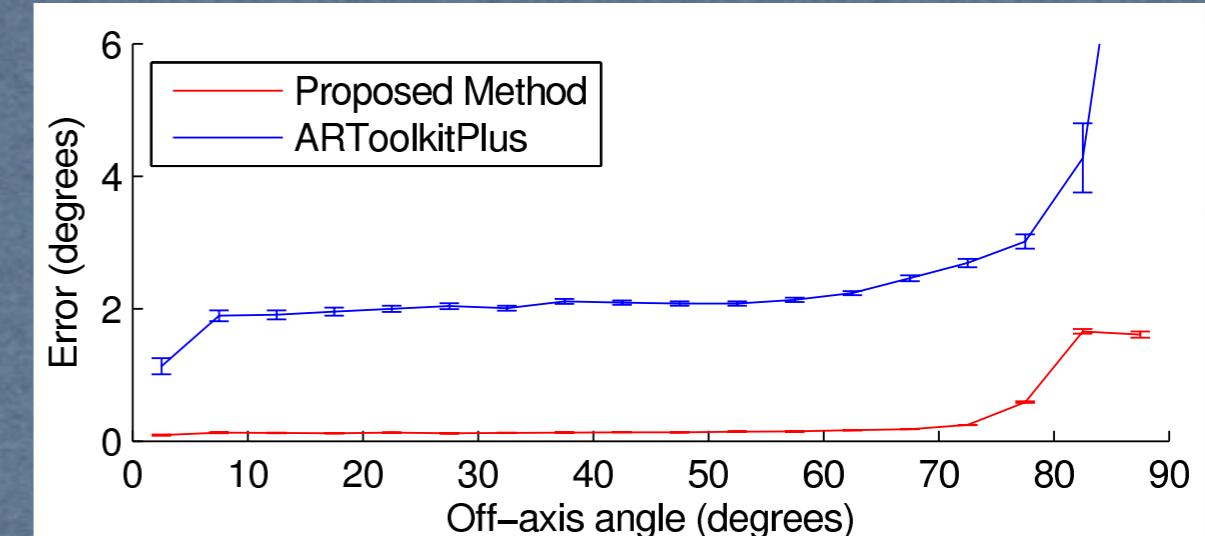


# Results: Accuracy

Distance



Orientation



- AprilTag both more accurate and detects more cases!

# AprilTag: Conclusions

- Very useful for robotics
  - ▶ Ground truthing
  - ▶ Commanding robots
  - ▶ Robots recognizing each other
  - ▶ Education
- Out-performs ARToolkit and ARTag
  - ▶ Better detection method
  - ▶ Better coding system
- Free and Open Source!



<http://april.eecs.umich.edu>