


EECS568 Mobile Robotics: Methods and Principles
Prof. Edwin Olson


L15. Laser Range Finders

Laser Scanners


- Measure range via time-of-flight
- SICK
 - ▶ Industrial safety
 - ▶ 180 samples, 1 degree spacing, 75Hz
 - ▶ Resolution ~1cm, ~0.25 deg.
 - ▶ Interlacing
 - ▶ Max range: "80m" 30m fairly reliable
 - ▶ Intensity
 - ▶ \$4500
- Hokuyo
 - ▶ ~1080samples, 0.25deg spacing, 270 degree FOV, 10-40Hz
 - ▶ \$1100 - \$5000






Laser Scanners: Planar Environments



Planar LIDARS in 3D



3D LIDAR Sensors



Aligning scans = Measuring motion

- Two important cases:
 - ▶ Incremental motion. Use lidar to augment (or replace!) odometry.
 - ▶ Loop closing. Identify places we've been before to over-constrain robot trajectory.



Feature Extraction

- Our plan:
 - ▶ Extract features from two scans
 - ▶ Match features
 - ▶ Compute rigid-body transformation
- Possible features:
 - ▶ Lines
 - ▶ Corners
 - ▶ Occlusion boundaries/depth discontinuities
 - ▶ Trees (!)

Line Extraction

- Given laser scan, produce a set of 2D line segments
- Two basic approaches:
 - ▶ Divisive (“Split N Fit”)
 - ▶ Agglomerative
- Our plan:
 - ▶ If we know which points belong to a line, how do we fit a line to those points?
 - ▶ How do we determine which points belong together?

Line Fitting


- Assume we know a set of points belong to a line. How do we compute parameters of the line?
- How to parameterize the line?
 - ▶ Slope + y intercept?
 - ▶ Endpoints of line
 - ▶ A point on the line + unit vector
 - ▶ R + theta

Line Fitting: Derivation

- Error (for one point):

$$|e| = (p_i - q) \cdot \hat{n}$$

$$\hat{n} = \begin{bmatrix} -\sin(\theta) \\ \cos(\theta) \end{bmatrix}$$



- Cost function
 - ▶ Solution found by minimizing:

$$err^2 = \sum_i ((p_i - q) \cdot \hat{n})^2$$

$$e^2 = \sum_i (x_i^T v)^2 \quad v = \begin{bmatrix} \cos \theta \\ \sin \theta \end{bmatrix}$$

$$= \sum_i (x_i \cos \theta + y_i \sin \theta)^2$$

$$= \sum_i x_i^2 \cos^2 \theta + 2x_i y_i \cos \theta \sin \theta + y_i^2 \sin^2 \theta$$

$$\frac{2e^2}{2\theta} = \sum_i -2x_i^2 \cos \theta \sin \theta + 2x_i y_i (\sin^2 \theta + \cos^2 \theta) + 2y_i^2 \sin \theta \cos \theta = 0$$

$$= \sum_i -x_i^2 \sin 2\theta + 2x_i y_i \cos 2\theta + y_i^2 \sin 2\theta = 0$$

$$= -M_{xx} \sin 2\theta + 2M_{xy} \cos 2\theta + M_{yy} \sin 2\theta = 0$$

$$= (-M_{xx} + M_{yy}) \sin 2\theta + 2M_{xy} \cos 2\theta = 0 \Rightarrow (-M_{xx} + M_{yy}) \sin 2\theta = -2M_{xy} \cos 2\theta$$

$$\Rightarrow \frac{\sin 2\theta}{\cos 2\theta} = \frac{2M_{xy}}{M_{yy} - M_{xx}} \quad \left(\frac{\Delta y}{\Delta x} \right)$$

$$\Rightarrow \theta = \frac{1}{2} \operatorname{atan2}(2M_{xy}, M_{yy} - M_{xx})$$

N.B. There are 2 solutions due to the $\frac{1}{2}$...

$$\theta = \frac{1}{2} \operatorname{atan}() \quad \text{and} \quad \theta = \frac{1}{2} (2\pi + \operatorname{atan}())$$

but both give the same line!

Is this a max or a min? (We want max!)

actually, $\theta = \frac{1}{2} \operatorname{atan} \left(\frac{2M_{xy}}{M_{yy} - M_{xx}} \right)$. Using $\operatorname{atan2}$ destroyed 2 more solutions! (tan is periodic with π) with factor $\frac{1}{2}$, this will lead to $\{\theta, \theta + \frac{\pi}{2}, \theta + \pi, \theta + \frac{3\pi}{2}\}$

Line Fitting: Strategy

$$err^2 = \sum_i ((p_{i_x} - q_x)\hat{n}_x + (p_{i_y} - q_y)\hat{n}_y)^2$$

- Method: work in terms of moments

$$\begin{aligned} M_x &= \sum p_x & C_{xx} &= \frac{1}{N} M_{xx} - \left(\frac{M_x}{N}\right)^2 \\ M_y &= \sum p_y & C_{xy} &= \frac{1}{N} M_{xy} - \frac{M_x}{N} \frac{M_y}{N} \\ M_{xx} &= \sum p_x^2 & C_{yy} &= \frac{1}{N} M_{yy} - \left(\frac{M_y}{N}\right)^2 \\ M_{xy} &= \sum p_x p_y \\ M_{yy} &= \sum p_y^2 \end{aligned}$$

Line Fitting

- Solution:

$$\begin{aligned} q_x &= \frac{1}{N} M_x \\ q_y &= \frac{1}{N} M_y \\ \theta &= \frac{\pi}{2} + \frac{1}{2} \text{atan2}(-2C_{xy}, C_{yy} - C_{xx}) \end{aligned}$$

- How does this answer compare to SVD-based line fitting?
 - ▶ In fact, θ gives the dominant eigenvector!
- *Because all quantities are written in terms of moments, we can compute this quantity incrementally and without storing all the points!*


Which points belong together?

- If we know which points belong together, we can compute the line.
- But we don't know which points belong to a single line!
- For now: make use of fact that lidar points arrive "in order"




Divisive Line Fitting

- Init: All points belong to a single line
- Split-N-Fit(points)
 - ▶ Fit line to points
 - ▶ if line error $<$ thresh then
 - return the line
- else
 - ▶ Which point fits the worst?
 - ▶ Split the line into two lines at that point
 - ▶ return Split-N-Fit(leftpoints) \cup Split-N-Fit(rightpoints)



Divisive Line Fitting



Divisive Line Fitting

- Pros:
 - ▶ Easy to implement
- Cons:
 - ▶ Assumes points are in scan order
 - ▶ Doesn't always generate good answers
 - Can split points that should belong together
 - "Split-N-Fit-N-Merge"
- Complexity?
 - ▶ As many as N divisions, each requiring up to N points. $O(N^2)$

Line Fitting: Agglomerative

- Agglomerative-Fit(points)
 - ▶ Init: create N-1 lines for each pair of adjacent points
 - ▶ do forever
 - for each pair of adjacent lines $i, i+1$
 - Compute error for line that merges those two lines
 - if minimum error $>$ thresh then
 - return lines
 - else
 - merge lines with minimum error

Agglomerative Line Fitting

Agglomerative Line Fitting

- Pros:
 - ▶ Good quality, matches human expectations
- Cons:
 - ▶ Assumes points are in scan order
 - ▶ A little bit tricky to implement *quickly* (use a min-heap)
- Complexity?

RANSAC

- RANdom SAmple Consensus
- best_model = null
- best_consensus = -1
- do forever
 - ▶ Select enough points at random to fit a model M
 - ▶ Compute consensus = # of points that agree with M
 - ▶ if consensus > best_consensus
 - best_model = M
 - best_consensus = c

RANSAC: details


- What is the model for line fitting?
 - ▶ How many points required?
 - ▶ What if we wanted to fit parabolas instead?
- How many iterations do we need?
 - ▶ What is the probability that we pick p inliers?
- How do we compute consensus? How close is “close enough”?
 - ▶ Threshold based on sensor noise
 - ▶ “Soft” consensus

RANSAC

- Pros:
 - ▶ Easy
 - ▶ Does not require points to be in scan order
- Cons:
 - ▶ Hard to exploit points being in scan order
 - ▶ Not obvious how to extract more than one line

Hough Transform


- Each point is evidence for a set of lines
 - ▶ Vote for each of the possible lines.
 - ▶ Lines with lots of evidence get many votes.



- Tends to be slow.

Other Features

- Corners
 - ▶ Intersections of nearby (?) lines.
 - ▶ Easy to extract from lines.
 - ▶ Only need ONE corner correspondence to compute RBT.
- Trees (Victoria Park)
- Depth Discontinuities
 - ▶ Beware of viewpoint effects



Aligning scans using features

- What method?
 - ▶ RANSAC!
 - ▶ Randomly guess enough correspondences to fit a model
 - ▶ Pick the model with the best consensus score.
- Can incorporate *negative* information
 - ▶ Penalize consensus if features are missing.



Aligning scans *without* features

- Can we align the individual lidar points?
 - ▶ Avoids potential for introducing error due to feature extractor problems
 - ▶ Works in any environment (?)
 - Though feature extraction can act as a filter
 - Tree detector rejects ground strikes in Victoria Park

Iterative Closest Point (ICP)

- until converged:
 - ▶ For each point in scan A, find the closest point in scan B.
 - ▶ Compute the RBT that best aligns the correspondences from the previous step.
- Robustness tweaks:
 - ▶ Outliers: If distance between corresponding points is too large, delete the correspondence.
 - ▶ Also match from B to A: if the correspondences are very different in distance, delete the correspondence.

Iterative Closest Line (ICL)


- LIDARs sample space at essentially arbitrary locations.
 - ▶ No reason to think that the exact points sampled in scan A were observed in scan B
- Solution: interpolate lines between points in scan B, find closest line instead of closest point.

Computing RBT from point correspondences

- We need to compute the rigid-body transformation
- 3DOFs:
 - ▶ Translation in x, y
 - ▶ Rotation
- Approaches:
 - ▶ Simple (bad) two point method
 - ▶ Optimal n point method


Naïve two point method (1/2)

- Compute translation component by aligning centroids



Naive two-point method (2/2)

- Now, rotate (so that lines between points have same angle)



General N-point method

- From Horn, “Closed-form solution of absolute orientation using unit quaternions”
- Actually a 3D method, but 2D special case is particularly easy.
- Formulation
 - ▶ Minimize mean squared error between corresponding points
 - ▶ Potential weakness: outliers

Optimal N-point method

- Assume that points x and points y are centered at zero. (This assumption can be patched up later)

$$e_i = y_i - R(x_i) - t$$

$$y \approx R(x) + t \qquad R(x_i) = \begin{bmatrix} \cos(\theta)x_{i_0} - \sin(\theta)x_{i_1} \\ \sin(\theta)x_{i_0} + \cos(\theta)x_{i_1} \end{bmatrix}$$

$$\begin{aligned} \chi^2 &= \sum e_i^T e_i \\ &= \sum y_i^T y_i - R(x_i)^T R(x_i) + t^T t - 2y_i^T R(x_i) - 2y_i^T t + 2t^T R(x_i) \end{aligned}$$

Optimal Translation

- Let's compute the optimal Translation

$$\begin{aligned} \chi^2 &= \sum e_i^T e_i \\ &= \sum y_i^T y_i - R(x_i)^T R(x_i) + t^T t - 2y_i^T R(x_i) - 2y_i^T t + 2t^T R(x_i) \\ \frac{\partial \chi^2}{\partial t} &= \sum 2t - 2y_i + 2R(x_i) = 0 \end{aligned}$$

- Recall: we assumed points x & points y are centered at origin. i.e.,

$$\sum y_i = \sum R(x_i) = 0$$

- And thus:

$$t = 0$$

- (i.e., the optimal translation is the one that aligns their centroids.)

Optimal Rotation

- Let's compute the optimal rotation

$$\begin{aligned}\chi^2 &= \sum e_i^T e_i \\ &= \sum y_i^T y_i - R(x_i)^T R(x_i) + t^T t - 2y_i^T R(x_i) - 2y_i^T t + 2t^T R(x_i)\end{aligned}$$

- Letting $t = 0$ per previous slide:

$$\chi^2 = \sum y_i^T y_i - R(x_i)^T R(x_i) - 2y_i^T R(x_i)$$

- Note that $y^T y$ and $R(x)^T R(x)$ are the lengths of the vectors, and are independent of the rotation. We can thus drop those terms.

Optimal Rotation


- Recall that: $R(x_i) = \begin{bmatrix} \cos(\theta)x_{i_0} - \sin(\theta)x_{i_1} \\ \sin(\theta)x_{i_0} + \cos(\theta)x_{i_1} \end{bmatrix}$

$$\begin{aligned}\chi^2 &= \sum y_i^T R(x_i) \\ &= \sum \cos(\theta)x_{i_0}y_{i_0} - \sin(\theta)x_{i_1}y_{i_0} + \sin(\theta)x_{i_0}y_{i_1} + \cos(\theta)x_{i_1}y_{i_1}\end{aligned}$$

- Collect sin/cos terms, we can write:


$$\begin{aligned}\chi^2 &= M \sin(\theta) + N \cos(\theta) \\ \frac{\partial \chi^2}{\partial \theta} &= M \cos(\theta) - N \sin(\theta) = 0 \\ \theta &= \text{atan2}(M, N)\end{aligned}$$

ICP In Action



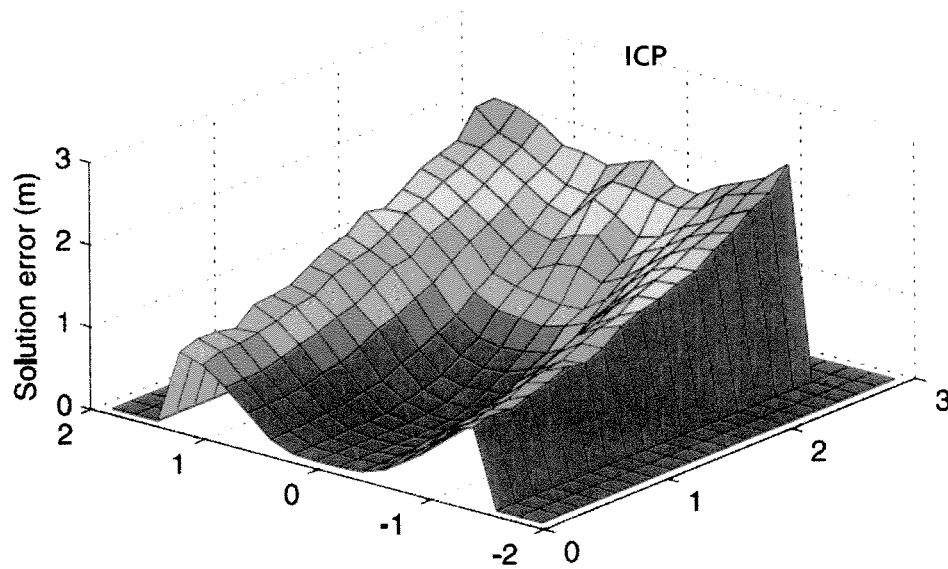
15.2 FPS

ICP In Action (2)



15.1 FPS

ICP Performance



Better point-wise matching

- Next time!

Features vs. Points

- Features (pros)
 - ▶ Data reduction
 - ▶ Noise filtering
 - ▶ Extraction + Matching is often fast
- Non-Feature Matching (pros)
 - ▶ General purpose: don't require world to contain our features
 - ▶ Very robust