


# Scan Alignment




Not-Aligned

Aligned

# Scan Alignment: SVD

- Optimal alignment between corresponding points
  - Assuming that for each source point, we know where the **corresponding** target point is



Slide Credit: [http://www.cse.wustl.edu/~taoju/cse554/lectures/lect05\\_Alignment](http://www.cse.wustl.edu/~taoju/cse554/lectures/lect05_Alignment)

# Scan Alignment: SVD

- SVD algorithm:

- Let  $P$  be a matrix whose  $i$ -th **column** is vector  $p_i - c_S$
- Let  $Q$  be a matrix whose  $i$ -th **column** is vector  $q_i - c_T$
- Consider the **cross-covariance** matrix:

$$M = PQ^T$$

- Find SVD of  $M$ :

$$M = U\Sigma V^T$$

- Find Rotation  $R$ :

$$R = UV^T$$

- Find Translation:

$$\vec{t} = c_T - R * c_S$$

# Scan Alignment: SVD

- SVD algorithm:

- Let P be a matrix whose  $i$ -th column is vector  $p_i - c_S$
- Let Q be a matrix whose  $i$ -th column is vector  $q_i - c_T$
- Consider the cross-covariance matrix:

$$M = PQ^T$$

- Find SVD of M:

$$M = U\Sigma V^T$$

- Find Rotation R:

$$R = UV^T$$

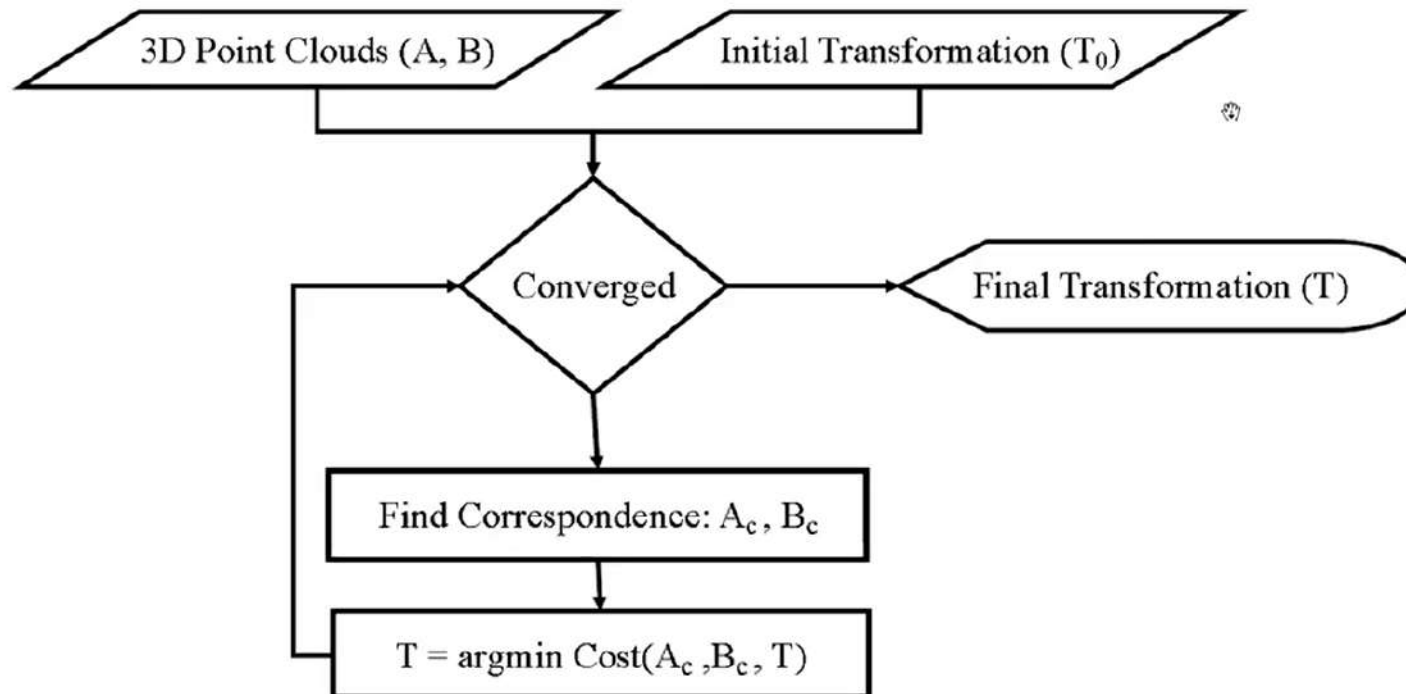
- Find Translation:

$$\vec{t} = c_T - R * c_S$$

## Limitations of SVD

- Requires correct point correspondences
  - Generally correspondences are not known !!

# Iterative Closest Point (ICP)




Besl, P. J. and McKay, N. D. (1992). A Method for Registration of 3-D Shapes. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 14(2):239-255.

# Generalized ICP

- Generalized-ICP is based on attaching a probabilistic model to the minimization step of the standard ICP algorithm.

$$A = \{a_i\}_{i=1,2\dots N} \quad B = \{b_i\}_{i=1,2\dots N} \quad (\text{st } a_i == b_i)$$

$$a_i \sim N(\mu_{ai}, \Sigma_{Ai}) ; \quad b_i \sim N(\mu_{bi}, \Sigma_{Bi})$$



Reference: "Generalized ICP" by A. Segal, D. Haehnel, S. Thrun. RSS 2009



Current Folder: F:\Courses\AE640\2018\EE698G\Assignment 2- 14376\Q4

Editor - F:\Courses\AE640\2018\EE698G\Assignment 2- 14376\Q4\main.m

```
18
19 % Remove zero columns in the scans (if any)
20 scan1(:, ~any(scan1,1)) = [];
21 scan2(:, ~any(scan2,1)) = [];
22
23 %% Perform ICP
24 keyboard;
25 [R,T] = ICP(scan1,scan2,itors,R_init,T_init,max_tresh); % -----> TO DO
26
27 % Note: find R,T such that they register scan2 onto scan1. This is just to maintain uniformity.
28
29 %% Visualizing the output
30 scan2_transformed = R*scan2+repmat(T,1,size(scan2,2));
31
32 figure
33 scatter(scan1(1,:),scan1(2,:), 'r','.');
34 hold on
35 scatter(scan2(1,:),scan2(2,:), 'b','.');
36 axis equal
37 title('Before ICP registration')
38 legend('scan2','scan1');
39
40 figure
41 scatter(scan2_transformed(1,:),scan2_transformed(2,:), 'r','.');
42 hold on
43 scatter(scan1(1,:),scan1(2,:), 'b','.');
44 axis equal
45 title('After ICP registration')
```

Workspace - main

Name	Value
dim	2
itors	500
lidar_scans	1x1 struct
max_tresh	10
R_init	[1,0,0,1]
scan1	2x681 double
scan2	2x681 double
T_init	[0,0]

Command Window

New to MATLAB? See resources for [Getting Started](#).

find\_correspondance.m (Function)

Finds the correspondances between the original scan p and transformed scan q

find\_correspondance(p, q, max\_tresh)

Waiting for input

Zoom: 100%

UTF-8


CRLF

script

Ln 24

Col 1





Current Folder: F:\Courses\AE640\2018\EE698G\Assignment 2- 14376\Q4

Editor: F:\Courses\AE640\2018\EE698G\Assignment 2- 14376\Q4\ICP.m

```

4 align_thresh=5.53e-04;
5 R=R_init;
6 T=T_init;
7
8 for j=1:iters
9     % Transform scan2 frame with respect to scan1 frame
10    scan2t=R*scan2+repmat(T,1,size(scan2,2));
11
12    % Find correspondences between points in scan1 and scan2
13    scan2t=find_correspondance(scan1, scan2t, max_tresh);
14
15    % Apply SVD Algorithm to compute transformations incrementally
16    mu1=mean(scan1,2);
17    mu2=mean(scan2t,2);
18    A=scan2t-repmat(mu2,1,size(scan2t,2));
19    B=scan1-repmat(mu1,1,size(scan1,2));
20    [U,~,V]=svd(B*A');
21    R=U*V'*R;
22    T=mu1-(U*V')*mu2+T;
23
24    % Computer alignment error
25    error_align=sqrt(sum((scan1(1,:)-scan2t(1,:)).^2+(scan1(2,:)-scan2t(2,:)).^2)/length(scan1));
26    disp(error_align);
27    if error_align<align_thresh
28        break;
29    end
30 end
31 end

```

Command Window

New to MATLAB? See resources for [Getting Started](#).

```

K =
    0.9965    0.0837
   -0.0837    0.9965

>> T

T =

   -0.0106
   -0.0005

fx >>

```

Workspace

Name	Value
dim	2
iters	500
lidar_scans	1x1 struct
max_tresh	10
R	[0.9965, 0.0837; -0.0837, 0.9965]
R_init	[1, 0; 0, 1]
scan1	2x681 double
scan2	2x681 double
scan2_transformed	2x681 double
T	[-0.0106; -5.1657e-04]
T_init	[0, 0]

Zoom: 100% UTF-8 CRLF ICP Ln 23 Col 6

Type here to search 23°C Light rain 14:27 03-02-2022


# Introduction to Computer Vision for Robotics

**AE640A** Autonomous Navigation

Mangal Kothari

Credit: Harsh Sinha

# What is Computer Vision?



## What is Computer Vision?

- *Vision* is about discovering from images what is present in the scene and where it is.
- In *Computer Vision* a camera (or several cameras) is linked to a computer. The computer interprets images of a real scene to obtain information useful for tasks such as navigation, manipulation and recognition.

# Overview of the field




Image Credits: Karpathy, CVPR'15

What kind of Information?




Image Credits: <https://tinyurl.com/lxuex6o>



# Overview of the field

Primary themes in Computer Vision are:

1. Object Detection
2. Segmentation



Segmentation: Which pixels  
belong to which object?  
Credits: Own Work

# Overview of the field

Primary themes in Computer Vision are:

1. Object Detection
2. Segmentation
3. Image Modifications/Enhancements




Image Colorization: From  
Grayscale to Colored Images  
Credits: Richard Zhang, CVPR 2016



# Overview of the field

Primary themes in Computer Vision are:

1. Object Detection
2. Segmentation
3. Image Modifications/Enhancements
4. Image to Text
5. Image Generation
6. Motion Estimation
7. 3D reconstruction from Images



3D Reconstruction: REMODE,  
Real Time Reconstruction  
Credits: Matia Pizzoli, ICRA 2014

# A look at history

- **Robert Nathan** started writing computer programs for enhancing images from NASA's spacecraft's at Jet Propulsion Lab, NASA.
- The Summer Vision Project: Project at MIT to **solve a significant part of visual system**. Primary Objective was to divide the image into object, background and chaos regions, **over the course of a summer**.



Credits: EE604, nasa.gov

MASSACHUSETTS INSTITUTE OF TECHNOLOGY  
PROJECT MAC

Artificial Intelligence Group  
Vision Memo. No. 100.

July 7, 1966

## THE SUMMER VISION PROJECT

Seymour Papert

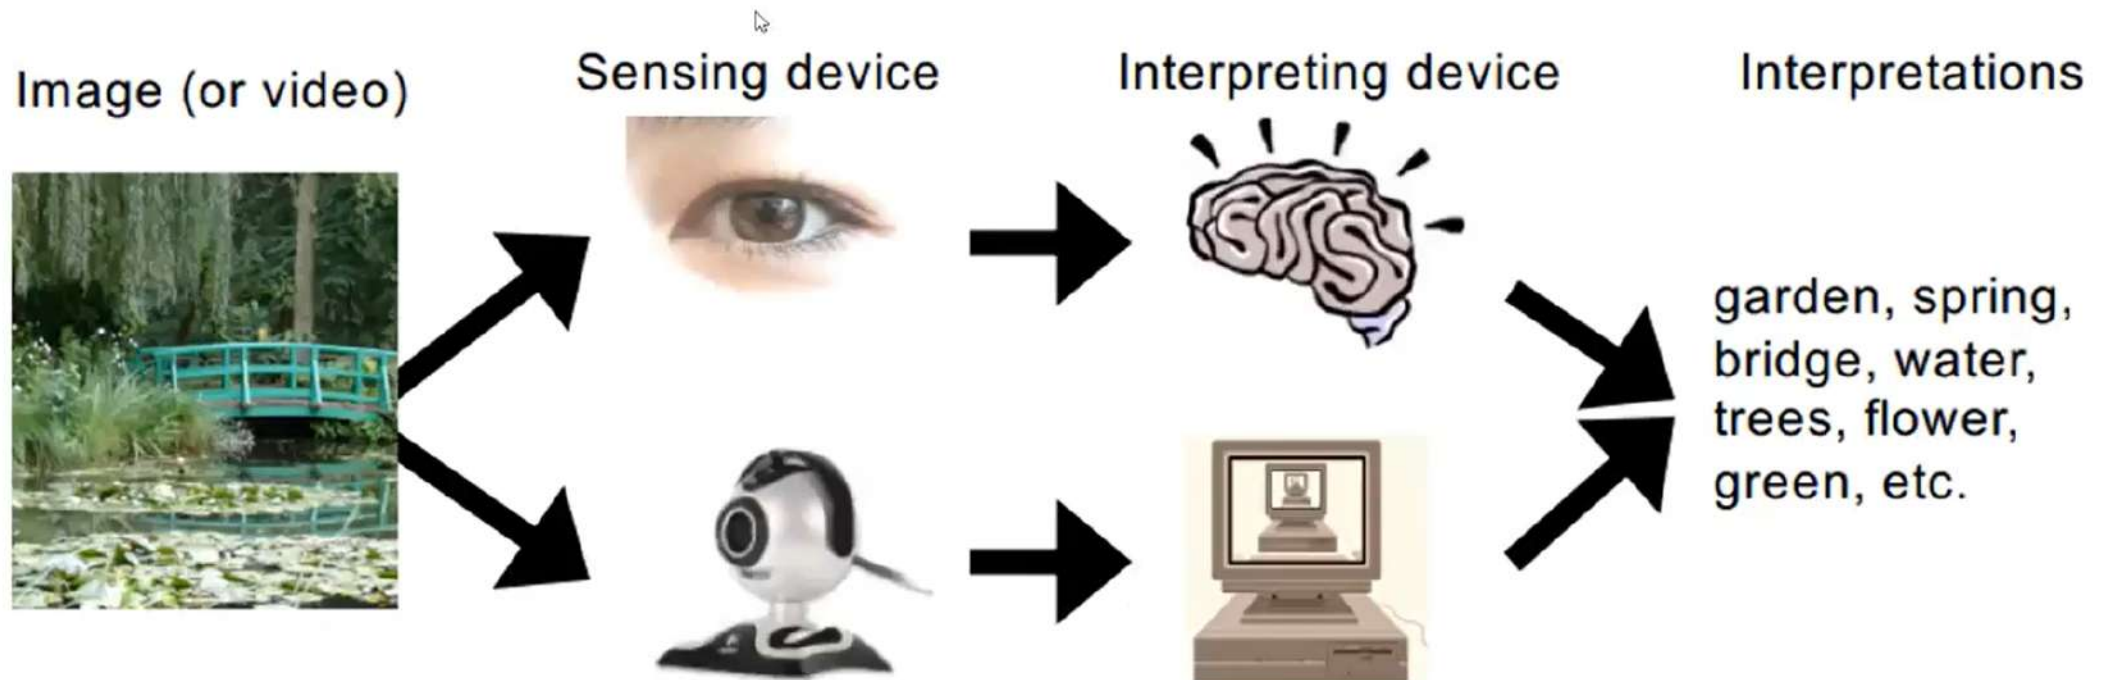
The summer vision project is an attempt to use our summer workers effectively in the construction of a significant part of a visual system. The particular task was chosen partly because it can be segmented into sub-problems which will allow individuals to work independently and yet participate in the construction of a system complex enough to be a real landmark in the development of "pattern recognition".

Credits: <https://tinyurl.com/y6bpo4nk>

# Hard Problem?

- Why are we still working on roughly the same problem as the “summer vision project”?
  - Why is it that creating 3D models of chairs is easier than identifying them?
- 
- There is a large gap between some  $\sim 1920 \times 1080 \times 3$  numbers and the high-level abstract meaning we associate with them.
  - Images are **2D** representation of information from **3D** world.

# The ~~(human)~~ ~~(computer)~~ vision system



Credits: CS131, Stanford

# Camera Models



Credits: <https://tinyurl.com/y6qen2vb>



# Camera Models

- Like so many things in engineering, we create a simple “model” of a camera to which is easy to understand and can approximate the actual functioning of a camera to a good degree.
- There are different models:
  - Pinhole camera model
  - Lens model
  - ...



# Cameras

- Demosaicing interpolates RGB subsamples to get colour image




This filter is  
called  
Bayer filter






# World Simplest Camera?



- Just hold up a piece of film
- Do we get an image on the film?
  - ▶ For each piece of the film, where do the photons come from?


Slide Credit: Edwin Olson, University of Michigan

# Let's add an aperture



- An aperture blocks all but a small subset of the rays
  - ▶ Causes the image to appear in focus!


# Aperture Size



- Why not make the aperture super small?
  - ▶ A “pin-hole” lens.
  - ▶ Not enough light to “register” on our film
- What happens when the aperture is bigger?
  - ▶ More rays can fit through--- blurrier image
- Is there any way of getting a sharp image, but allow more light through?
  - ▶ Yes! A *lens*.

Slide Credit: Edwin Olson, University of Michigan



# Lenses



- A lens collects rays with a particular divergence and refocuses them to a point.
  - ▶ But points at the “wrong” distance won’t be refocused exactly.
- *Depth of field*: how much of the scene is in focus
- We’re going to ignore this today, however--- we’re going to assume a “pin-hole” model.

Slide Credit: Edwin Olson, University of Michigan

# Pinhole camera model



Credits: Wikipedia, Pinhole Camera Model

# The Perspective Matrix

- Suppose we write a point in the world (like the position of the candle flame) as a vector:

$$p = \begin{bmatrix} x \\ y \\ z \end{bmatrix}$$

- Can we write a matrix so that  $p' = Mp$ ?

$$p' = \begin{bmatrix} fx/z + c_x \\ fy/z + c_y \end{bmatrix}$$

Slide Credit: Edwin Olson, University of Michigan

# Do homogeneous coordinates help?

- Eureka!

$$p' = \begin{bmatrix} fx/z + c_x \\ fy/z + c_y \\ 1 \end{bmatrix} = \begin{bmatrix} fx + c_x z \\ fy + c_y z \\ z \end{bmatrix} = \begin{bmatrix} f & 0 & c_x & 0 \\ 0 & f & c_y & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

2d  
homogeneous  
coordinates2d  
homogeneous  
coordinatesperspective  
transform3d  
homogeneous  
coordinates

Slide Credit: Edwin Olson, University of Michigan



# Rigid-Body Transformations

- The product of two rigid-body transformations is **always** another rigid-body transformation!
- So no matter how the object has been translated or rotated, we can describe its position with a single 4x4 matrix, which has the structure:

$T_x$   
 $T_y$   
 $T_z$

$$\begin{bmatrix} x' \\ y' \\ z' \\ 1 \end{bmatrix} = \begin{bmatrix} R_{00} & R_{01} & R_{02} & T_x \\ R_{10} & R_{11} & R_{12} & T_y \\ R_{20} & R_{21} & R_{22} & T_z \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

Slide Credit: Edwin Olson, University of Michigan

# Putting it all together

$$\begin{bmatrix} x' \\ y' \\ s \end{bmatrix} = \begin{bmatrix} f & 0 & c_x & 0 \\ 0 & f & c_y & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} R_{00} & R_{01} & R_{02} & T_x \\ R_{10} & R_{11} & R_{12} & T_y \\ R_{20} & R_{21} & R_{22} & T_z \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

2d  
homogeneous  
pixel (camera)  
coordinates

Focal length and  
focal center of camera  
  
“Intrinsics”

Rigidly move every object in  
the world to simulate the  
camera’s true position  
  
“Extrinsics”

3d  
homogenous  
(world)  
coordinates

Slide Credit: Edwin Olson, University of Michigan