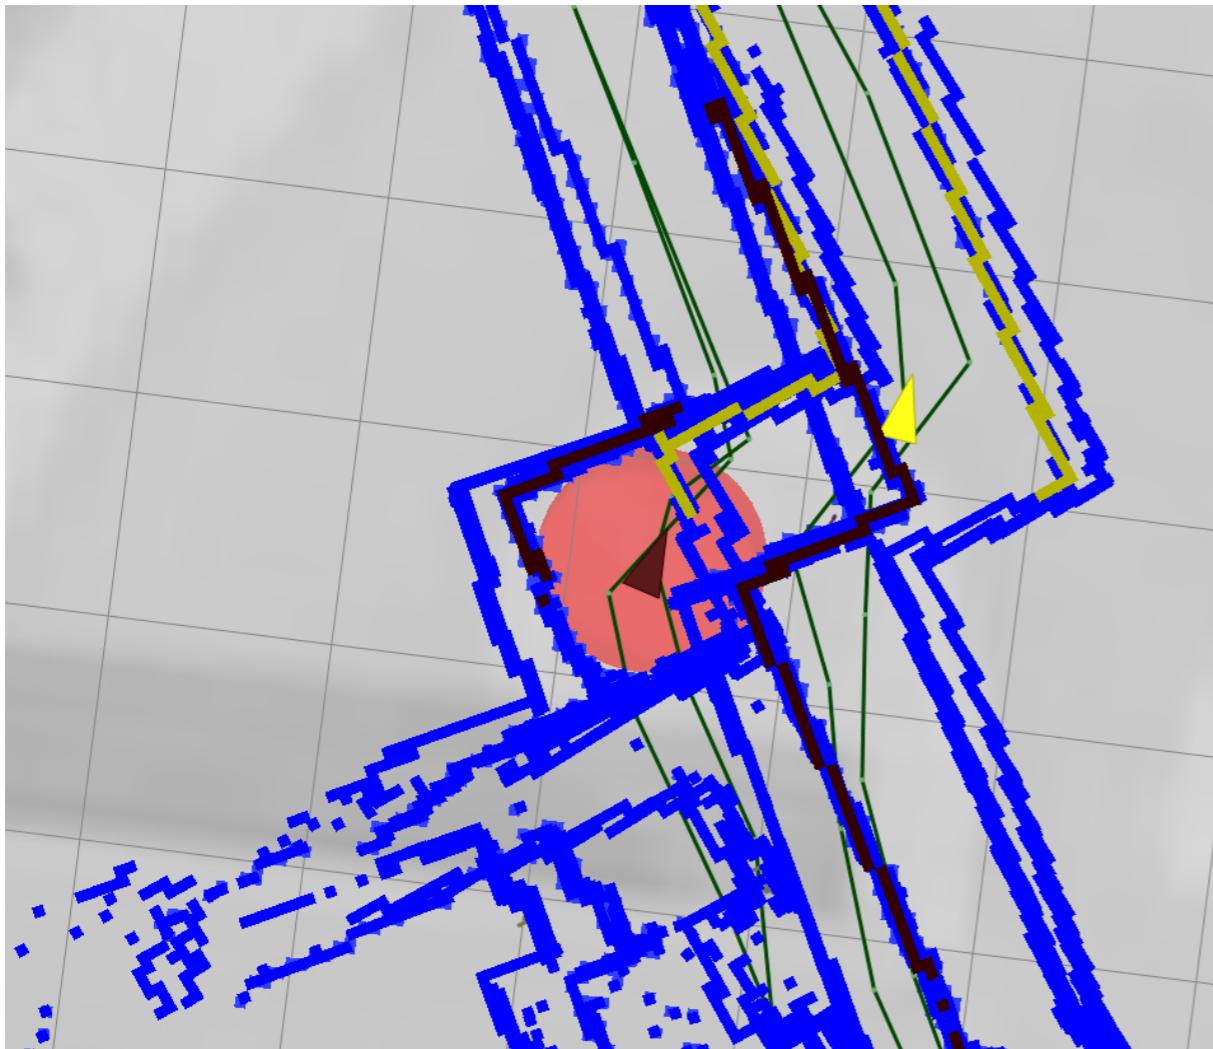


EECS568 Mobile Robotics: Methods and Principles  
Prof. Edwin Olson

# L16. Scan Matching and Image Formation

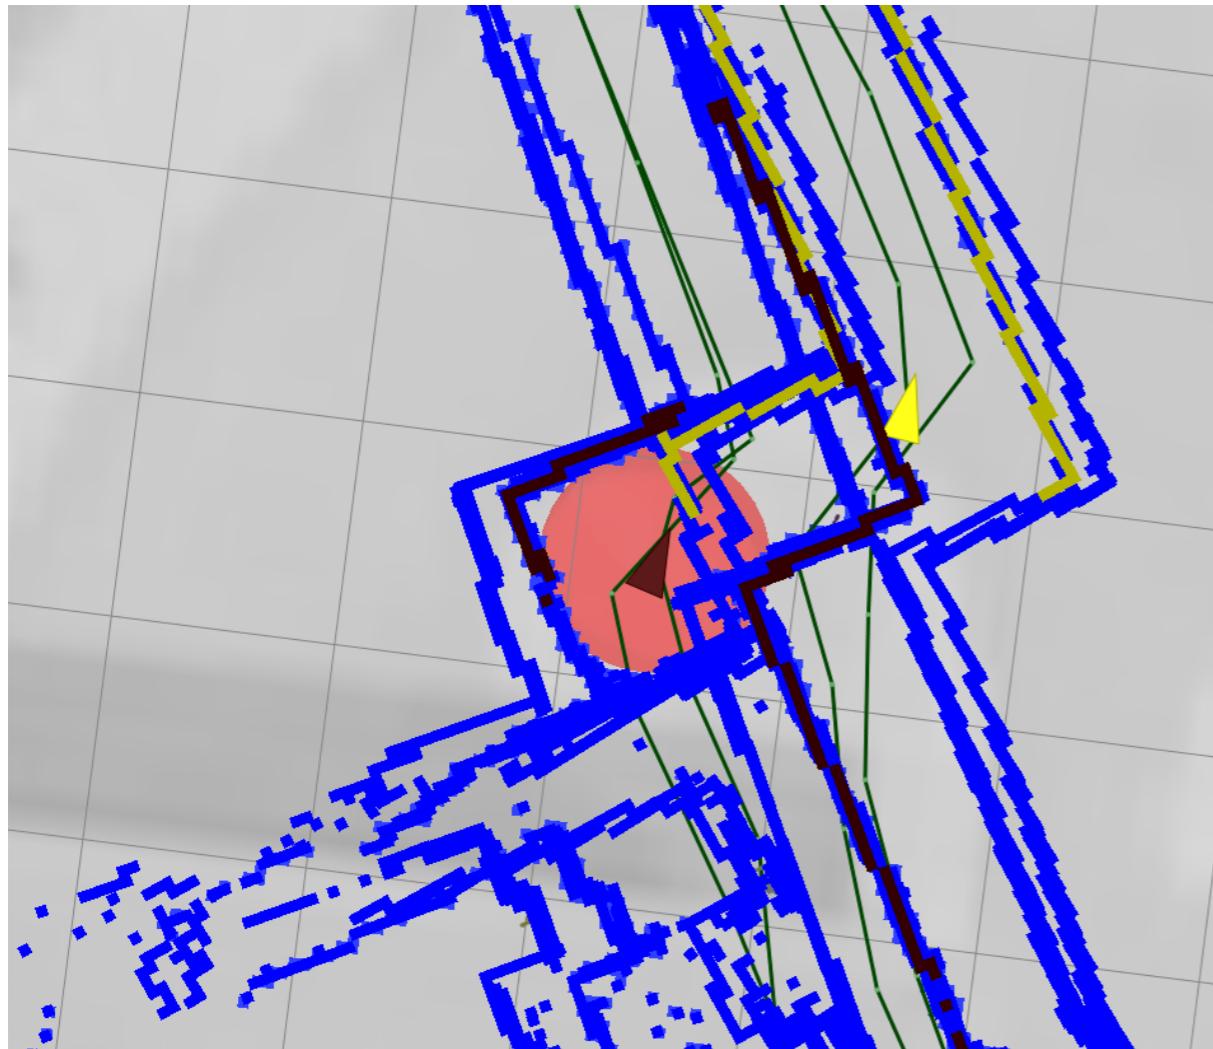
# Scan Matching



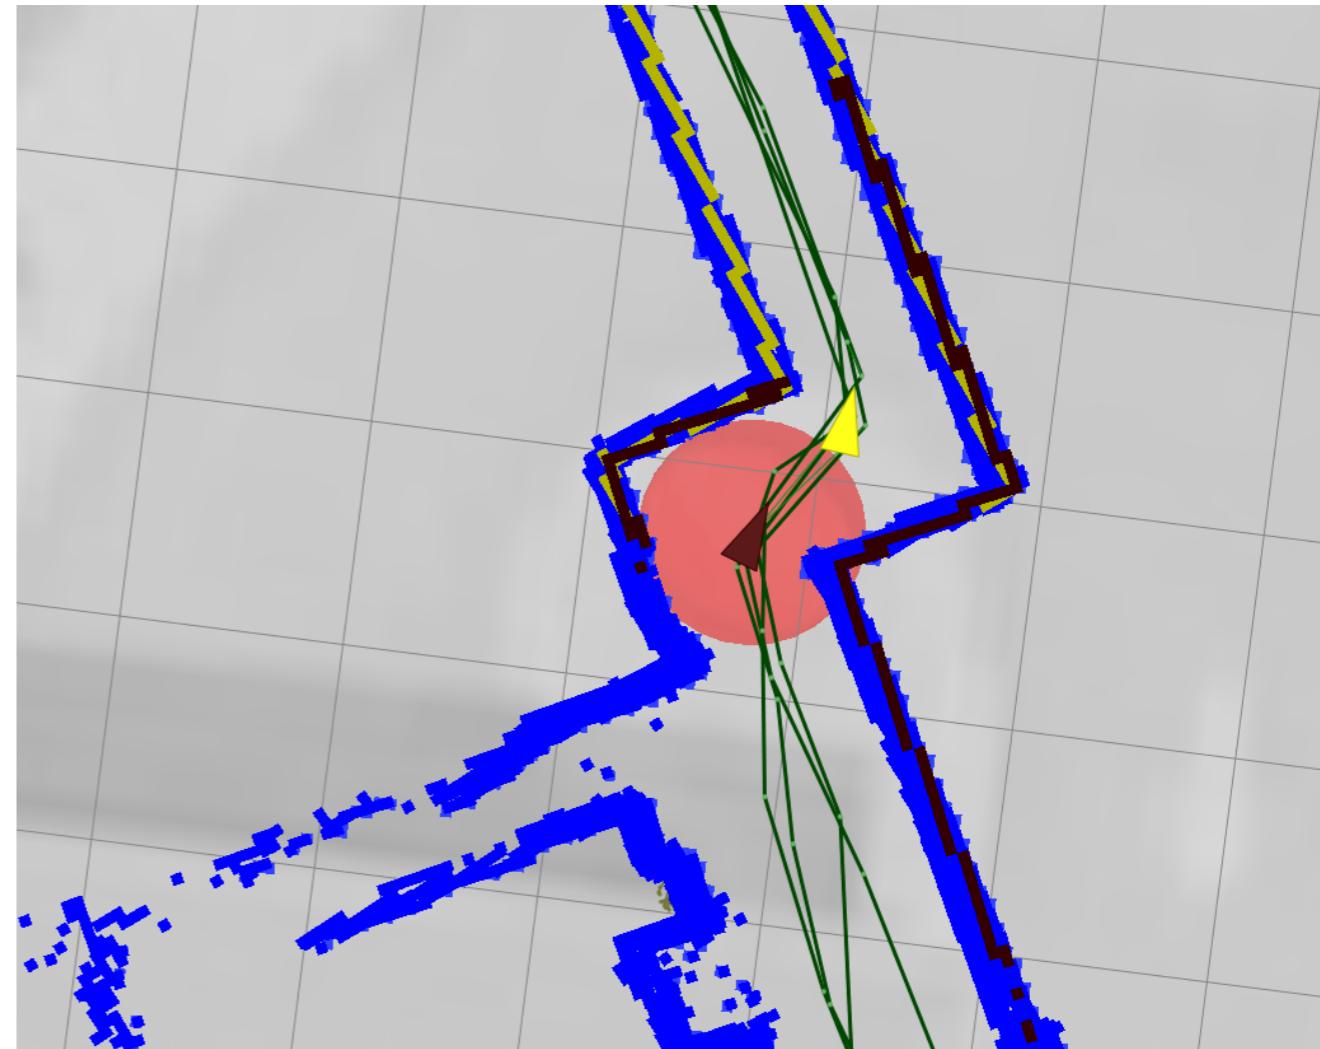
Before

After

# Scan Matching



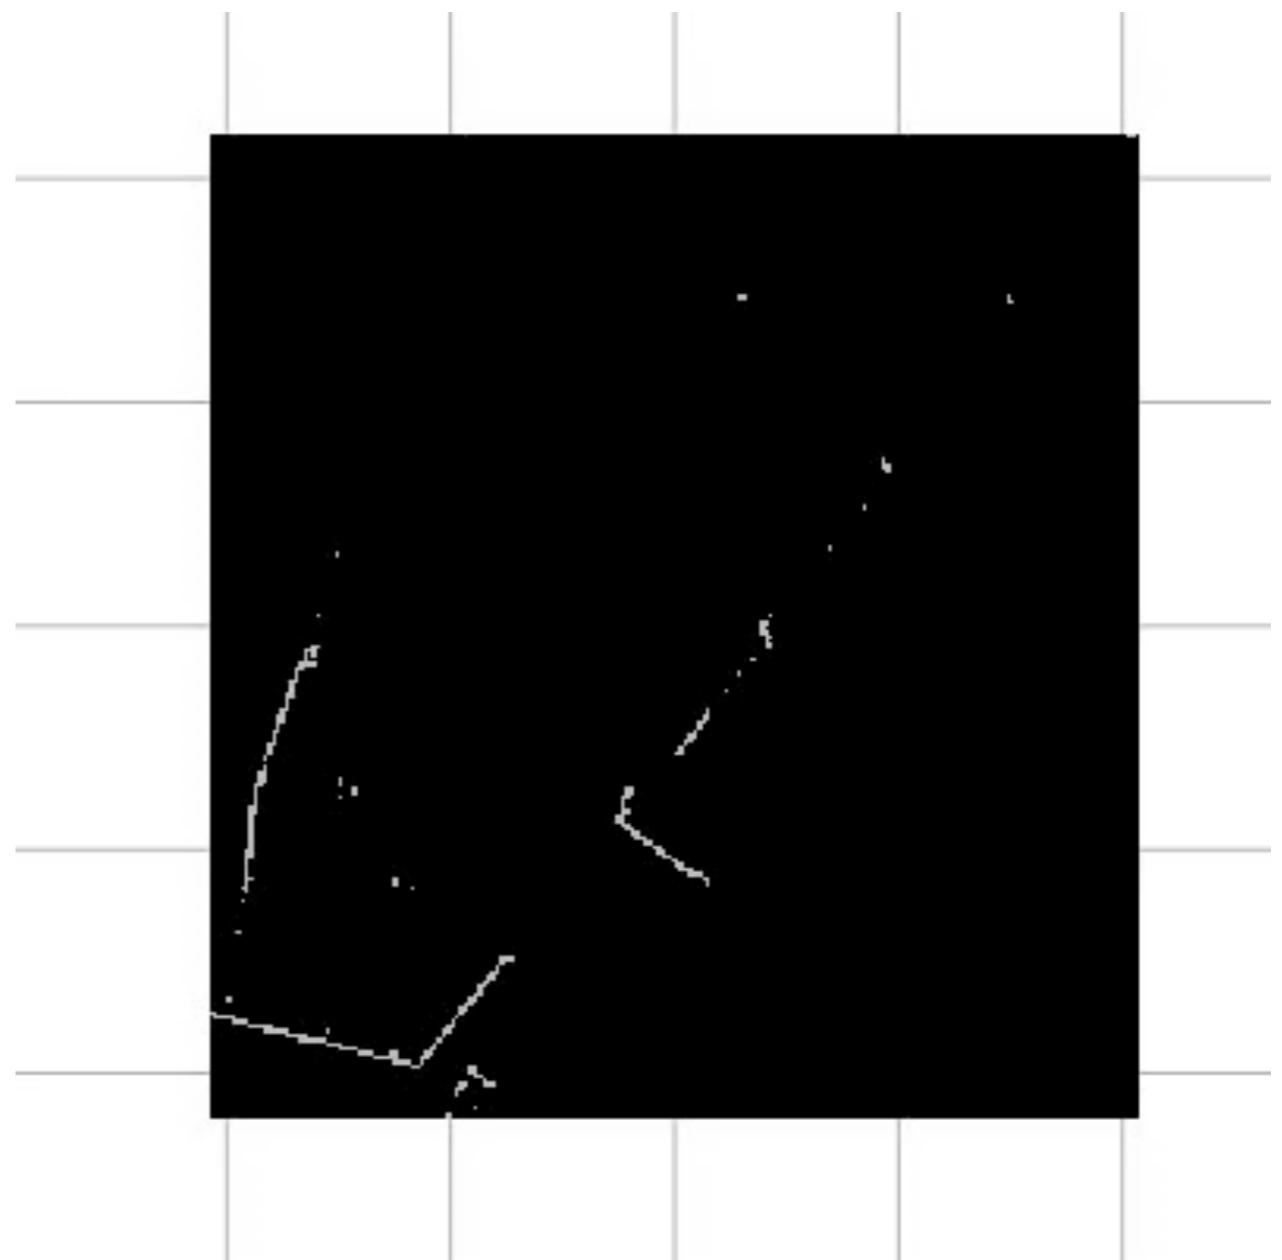
Before



After

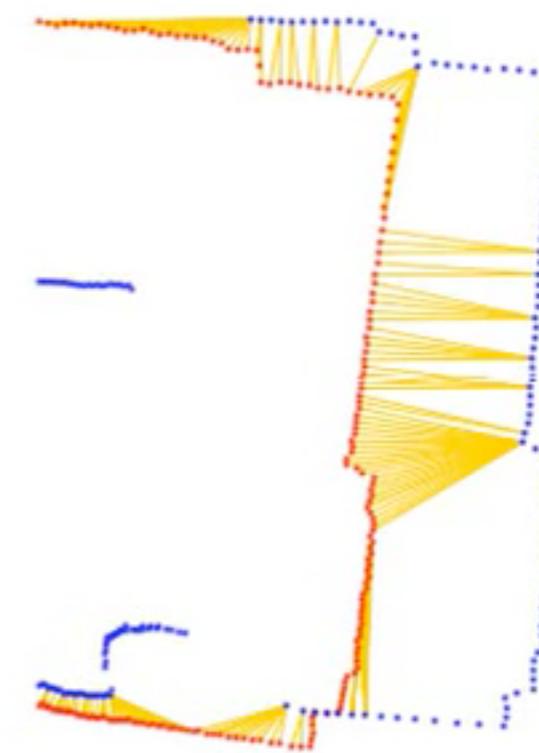
# Map matching has to be *fast*

- 14 robots generate new maps at ~5 Hz.
- Must attempt matches against as many other possibly-overlapping maps as possible
- Our previous best methods manage ~50Hz.
  - Fast, but not fast enough!



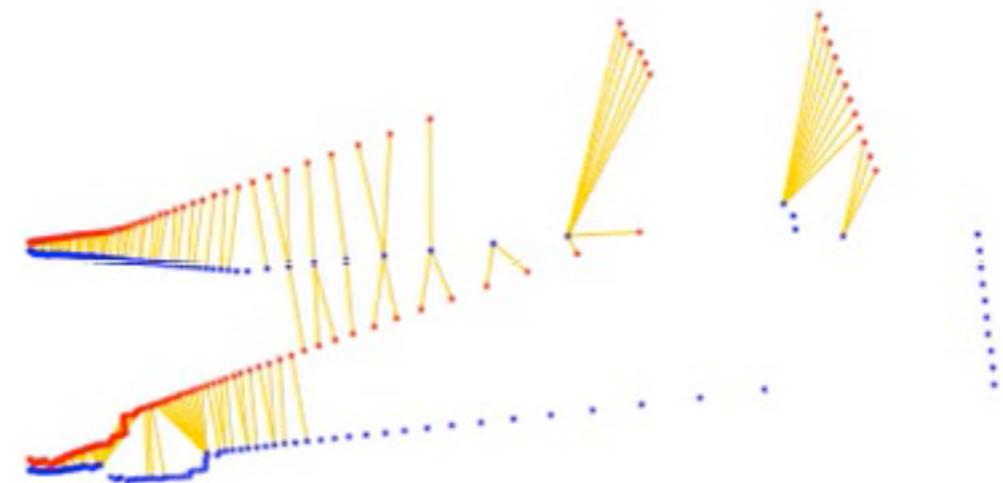
# Classic Method: ICP

- Associate each point in scan A with its nearest neighbor in scan B
  - Compute rigid-body transformation
- Repeat.



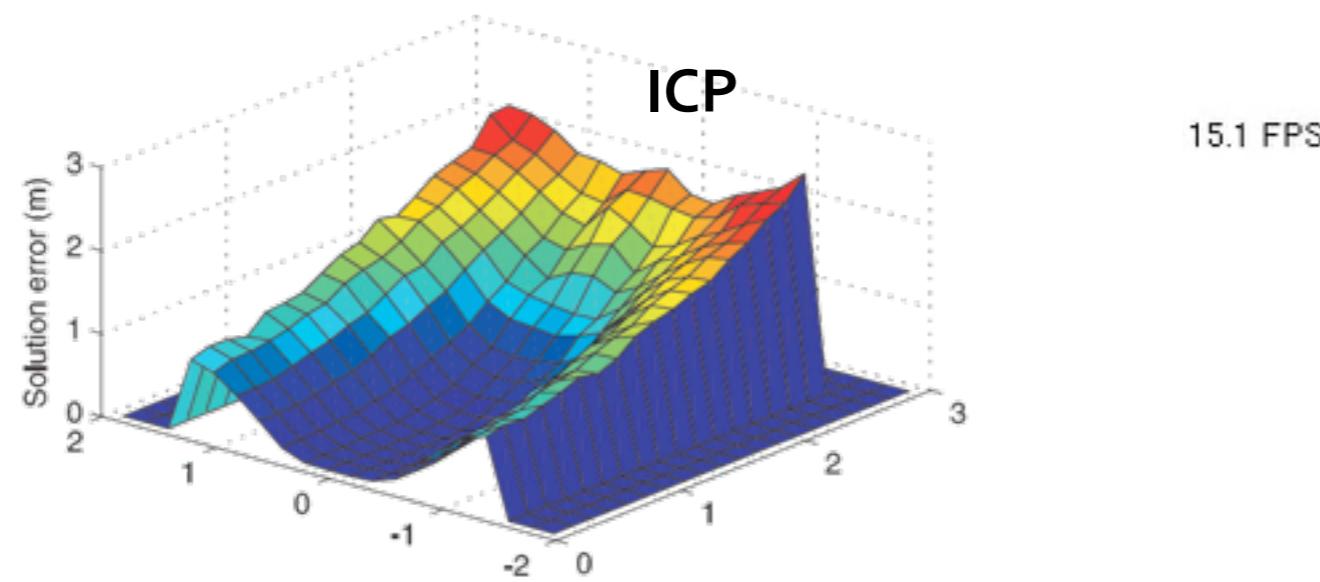
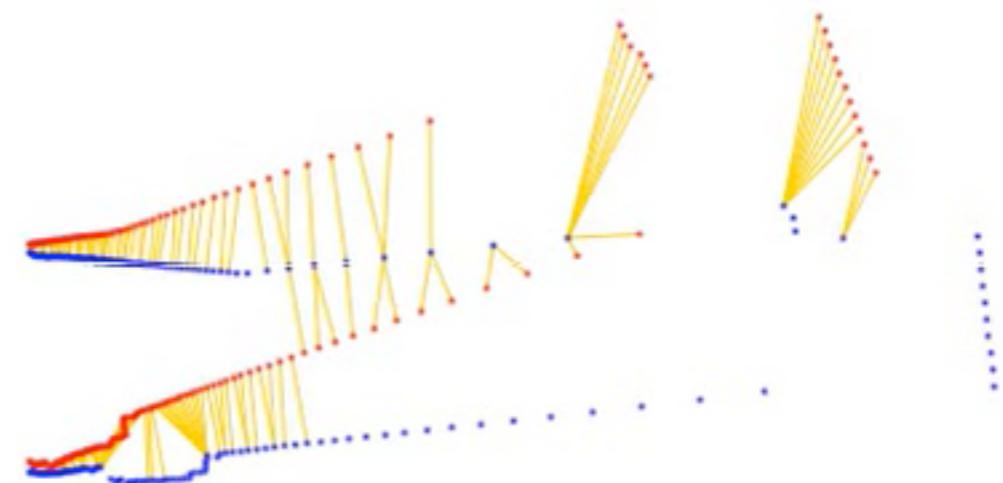
# Classic Method: ICP

- Not probabilistically grounded
- Often fails to converge to even a “sensible” solution.



# Classic Method: ICP

- Not probabilistically grounded
- Often fails to converge to even a “sensible” solution.



# Probabilistic Formulation

- What we want:

$$p(x_i | x_{i-1}, u, z, m)$$

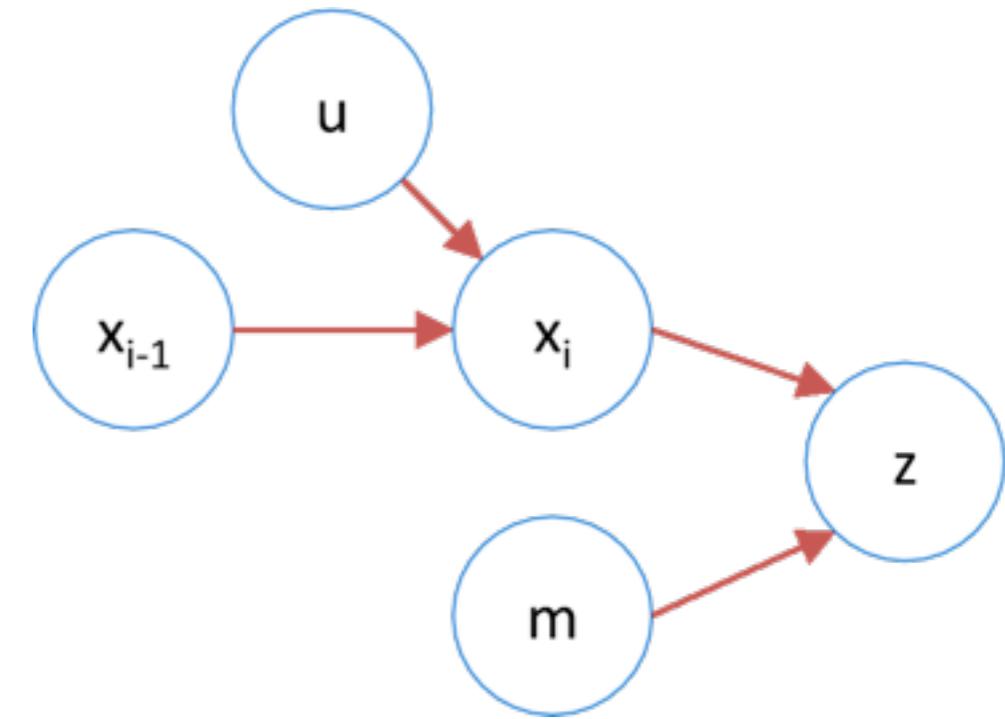
- Apply Bayes' law and exploit conditional independencies:

$$p(x_i | x_{i-1}, u, z, m) \propto p(z|x_i, m)p(x_i|x_{i-1}, u)$$

*observation  
model*

*motion  
model*

- Basic idea: search for  $x_i$  that maximizes probability.



# Observation model

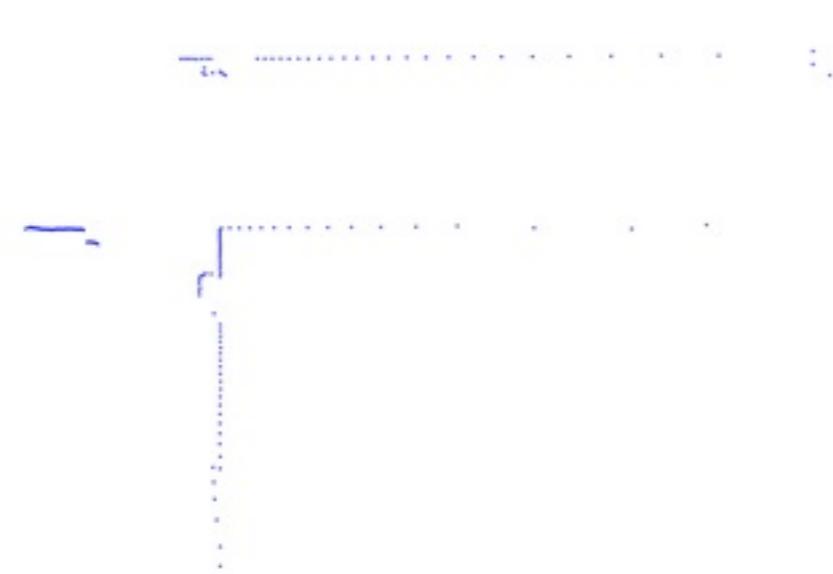
- How do we compute  $p(z|x_i, m)$  ?
- Lie #1: Assume independent laser measurements

$$p(z|x_i, m) = \prod_j p(z_j|x_i, m)$$

- Lie #2: Ignore occlusion/visibility effects
- Lie #3: Assume radially symmetric Gaussian noise model
- We can use a simple 2D lookup table for the probability of each point!

# Lookup Table

- Precompute log-likelihood of observing a point at each location
  - Approximately a convolution of a Gaussian with the “reference” scan



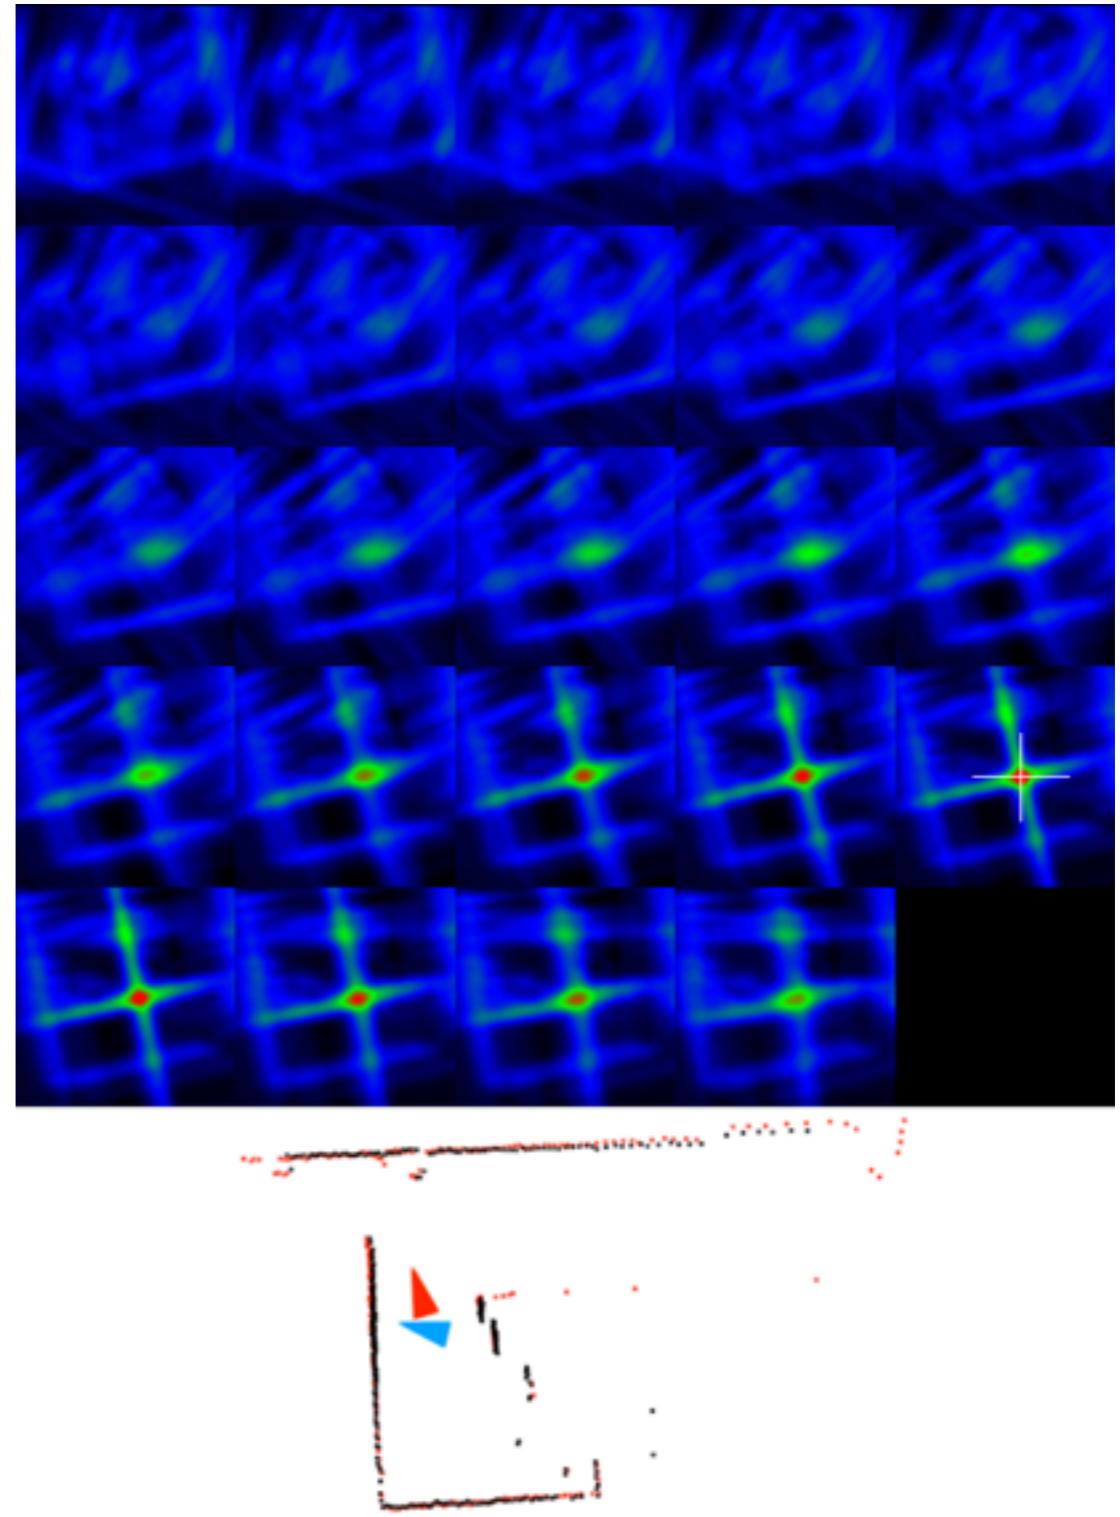
Reference Scan



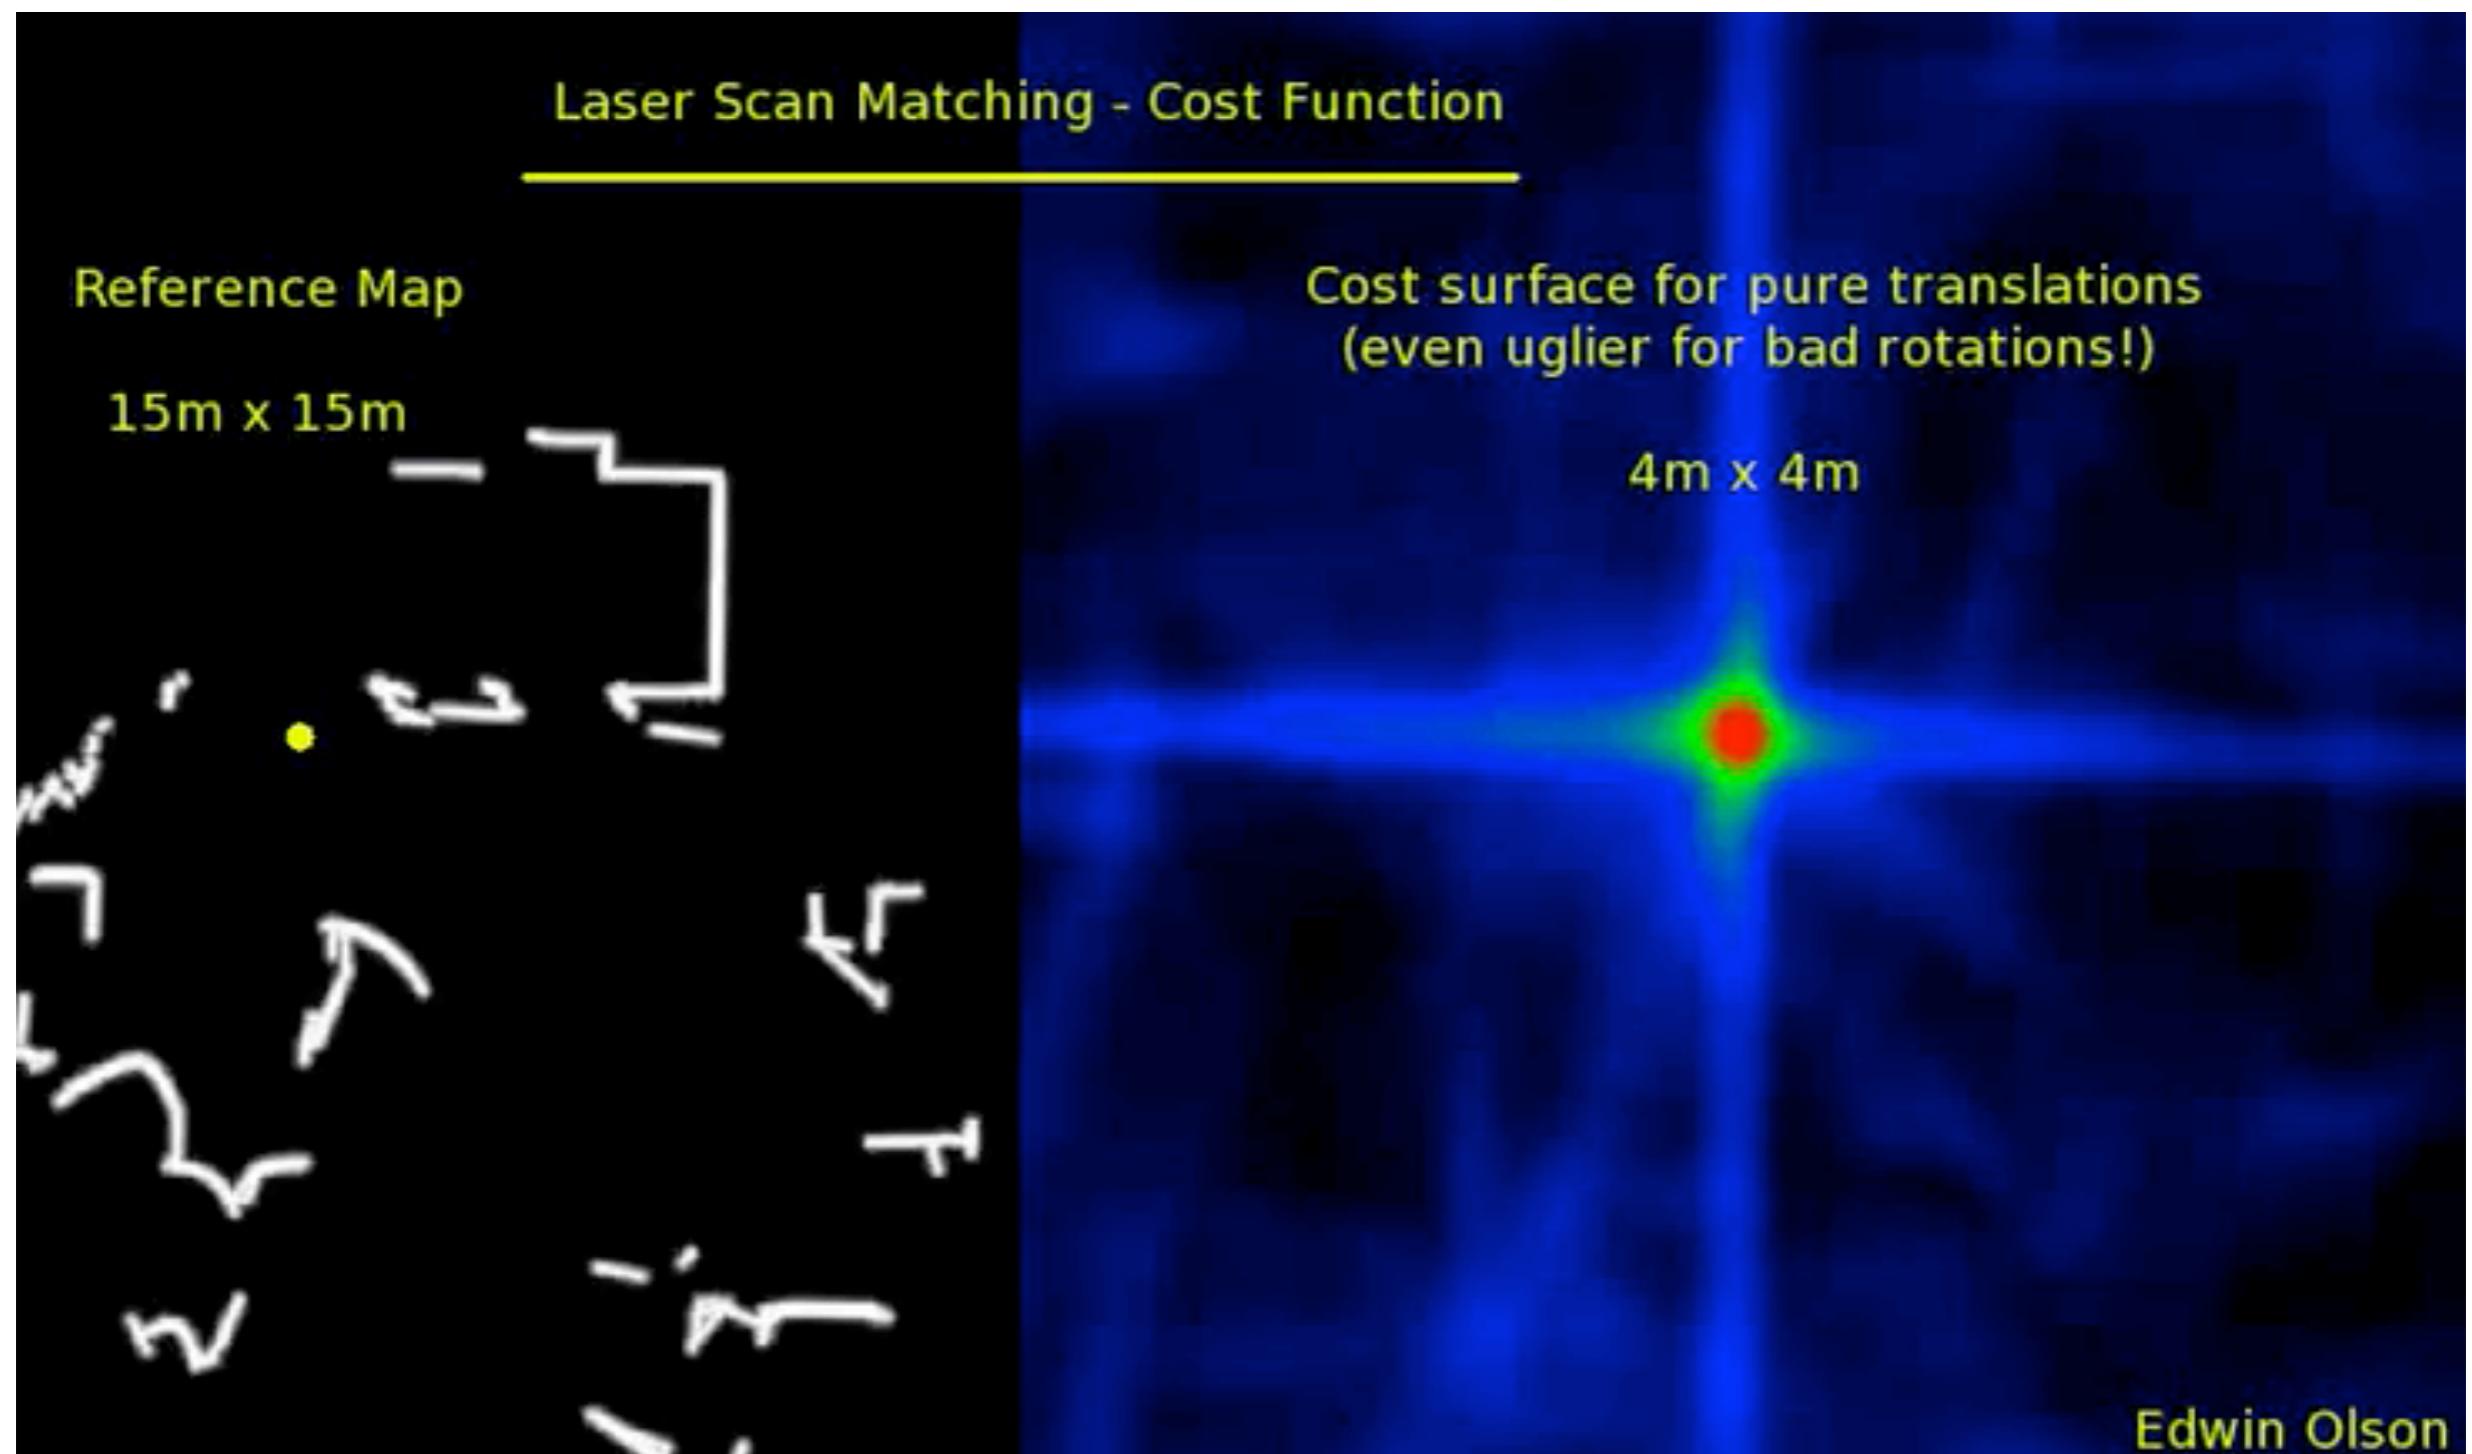
Lookup Table

# Correlative Scan Matching

- If we exhaustively search over  $x_i$ 's, we just pick out the maximum likelihood solution.
- This is quite slow!
  - Runtime grows rapidly with amount of uncertainty



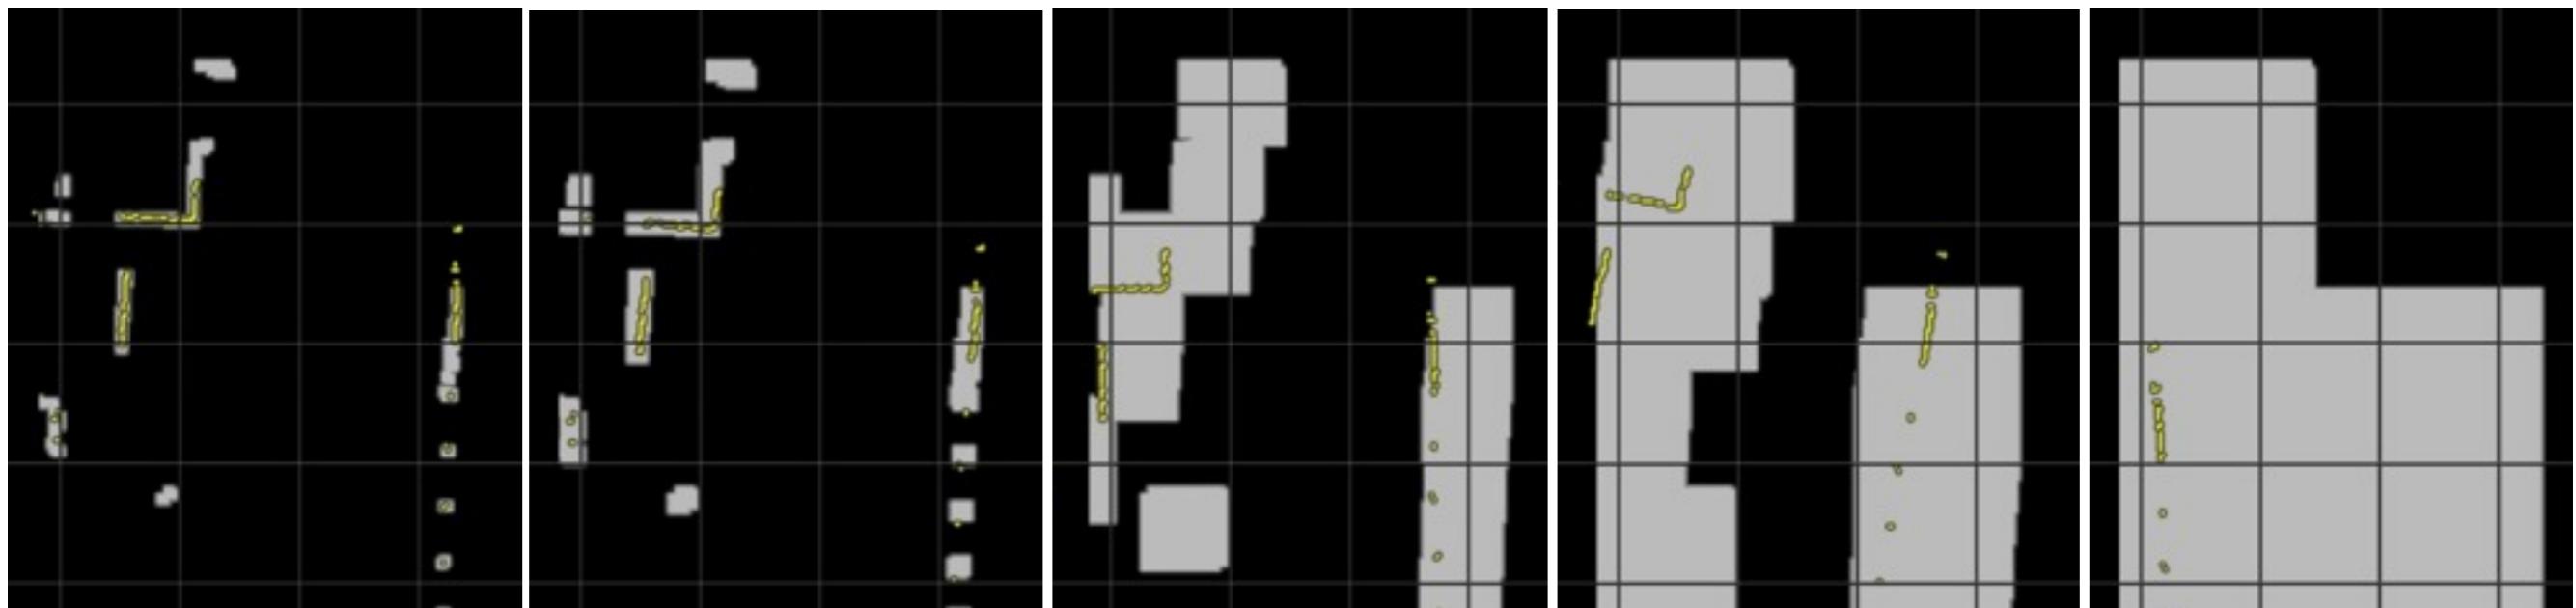
# Scan Matching in Action



# Multi-Resolution scan matching

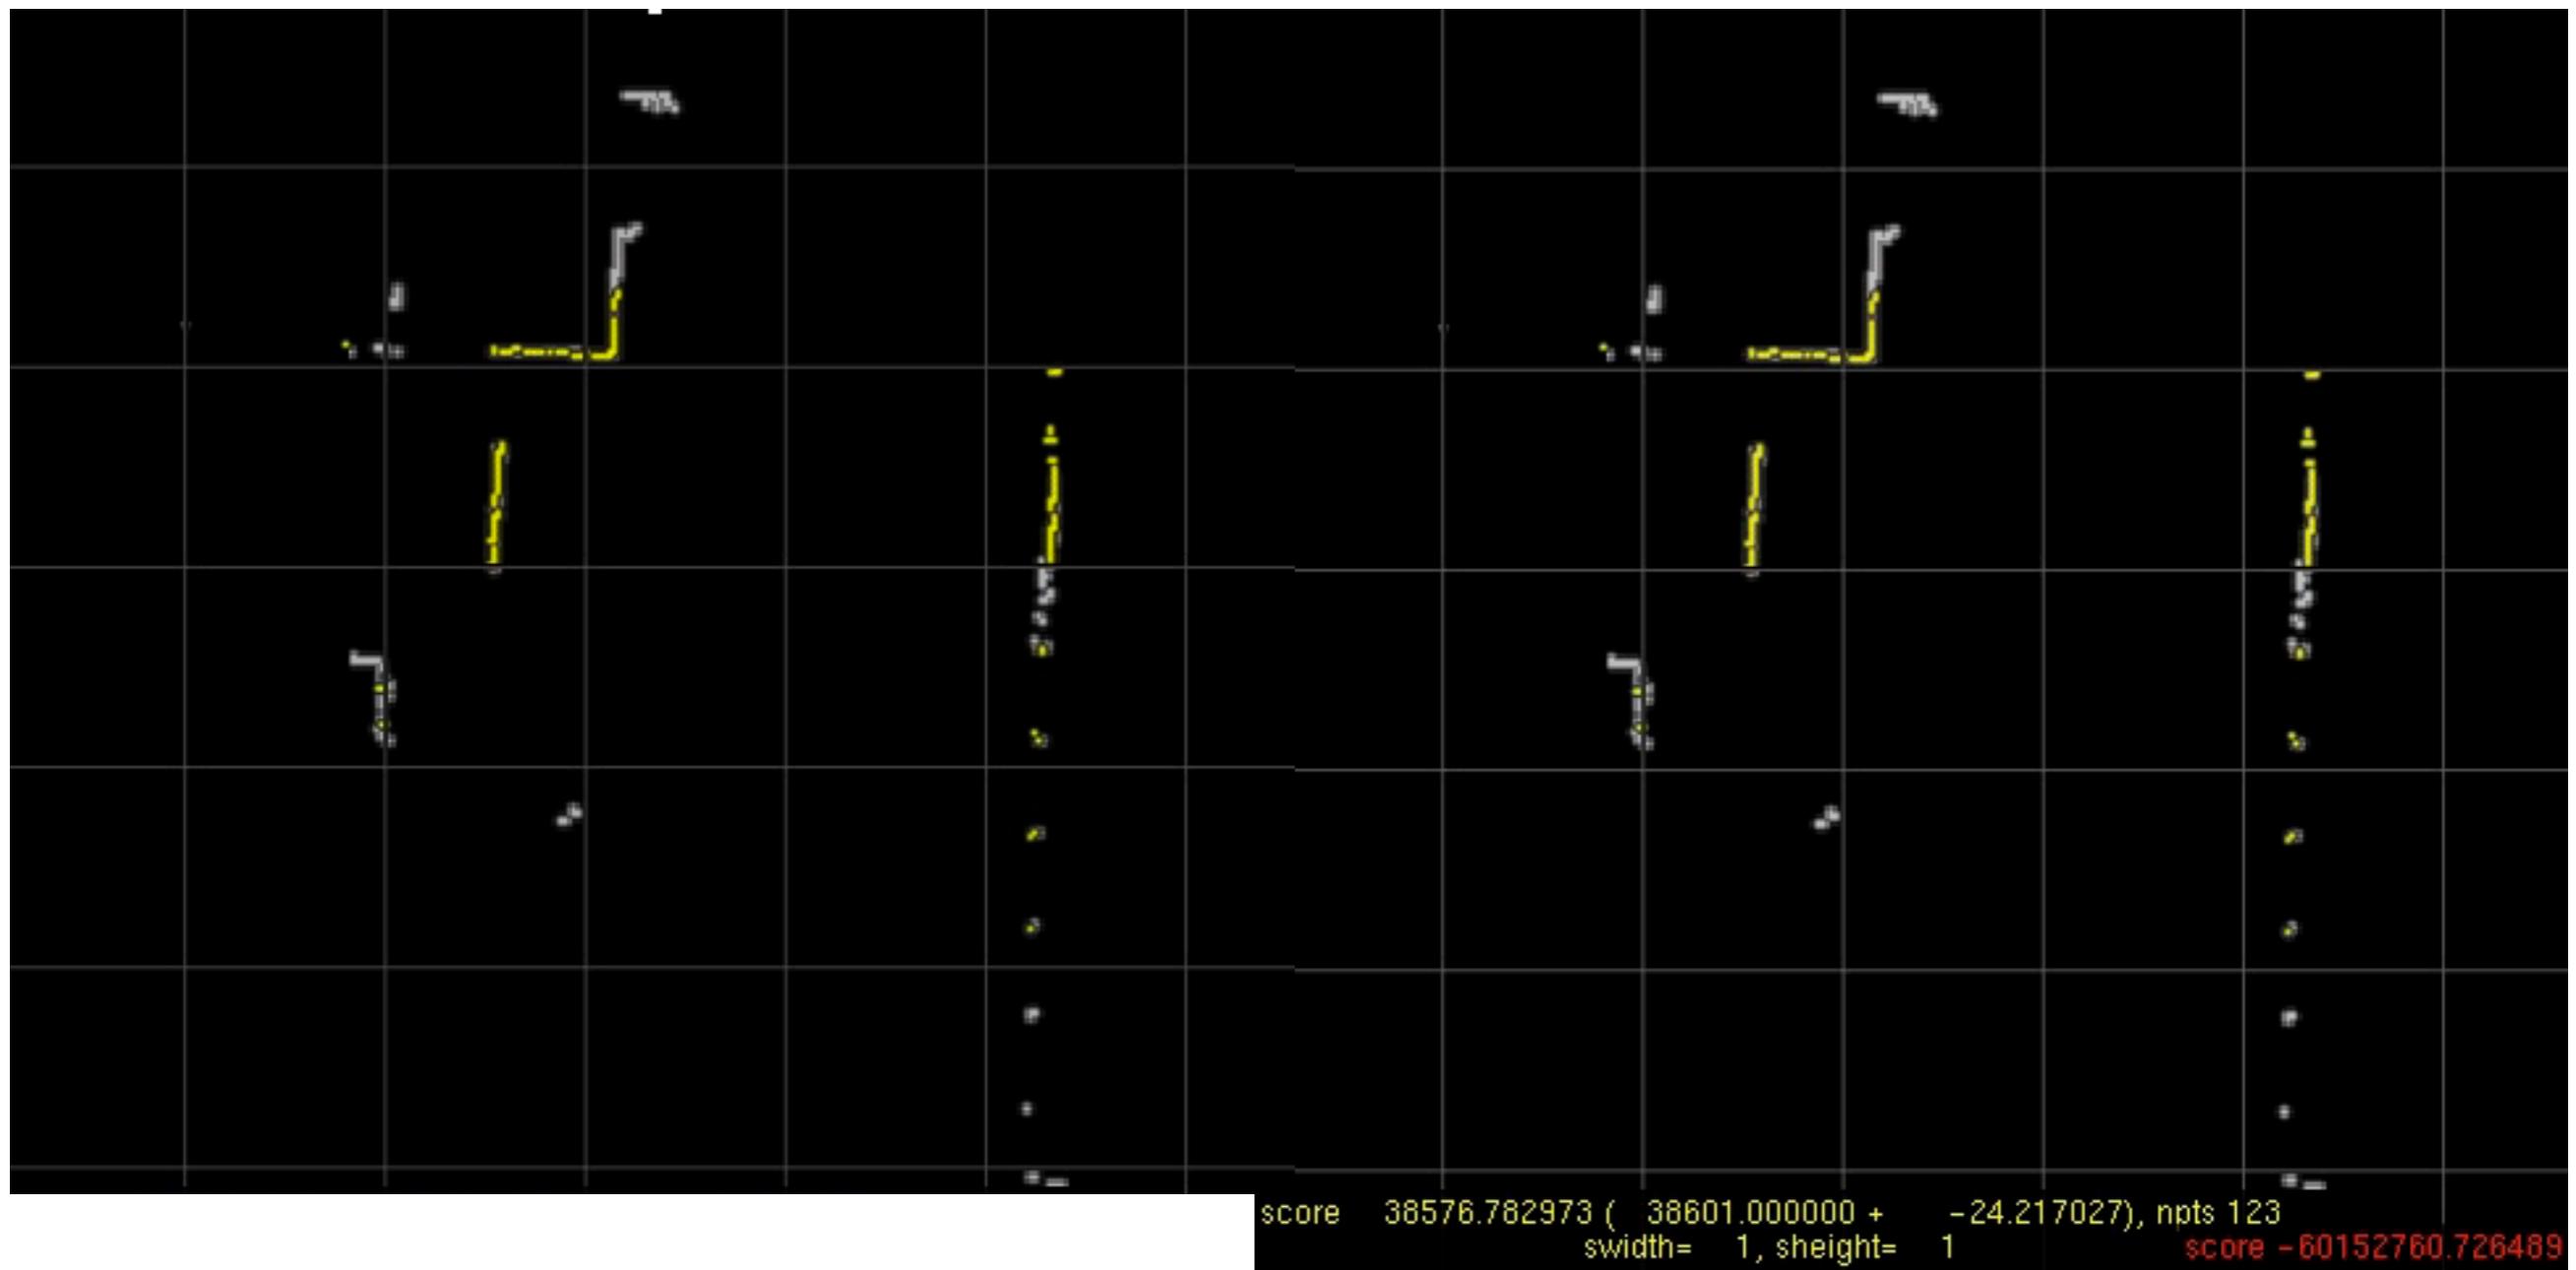
- Initialize a heap
- Populate heap with an exhaustive search at very coarse resolution
  - Use low-resolution versions of lookup tables so that we don't suffer from aliasing, and so we don't miss the optimum
- while true
  - Extract most promising solution,  $s$
  - if full-resolution( $s$ )
    - return  $s$
  - else
    - Resample  $s$  at finer resolution, add to heap.

# Multi-resolution matching



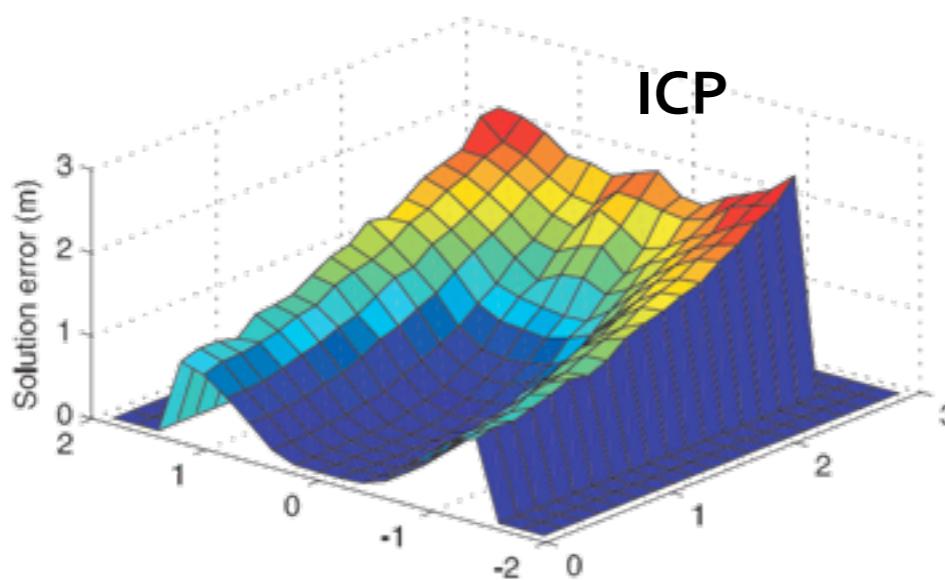
- Result of max-convolutions
  - Guarantees that score computed by low-resolution version is at least as good as any higher-resolution version in the same area.

# Multi-resolution matching



# Matching Results

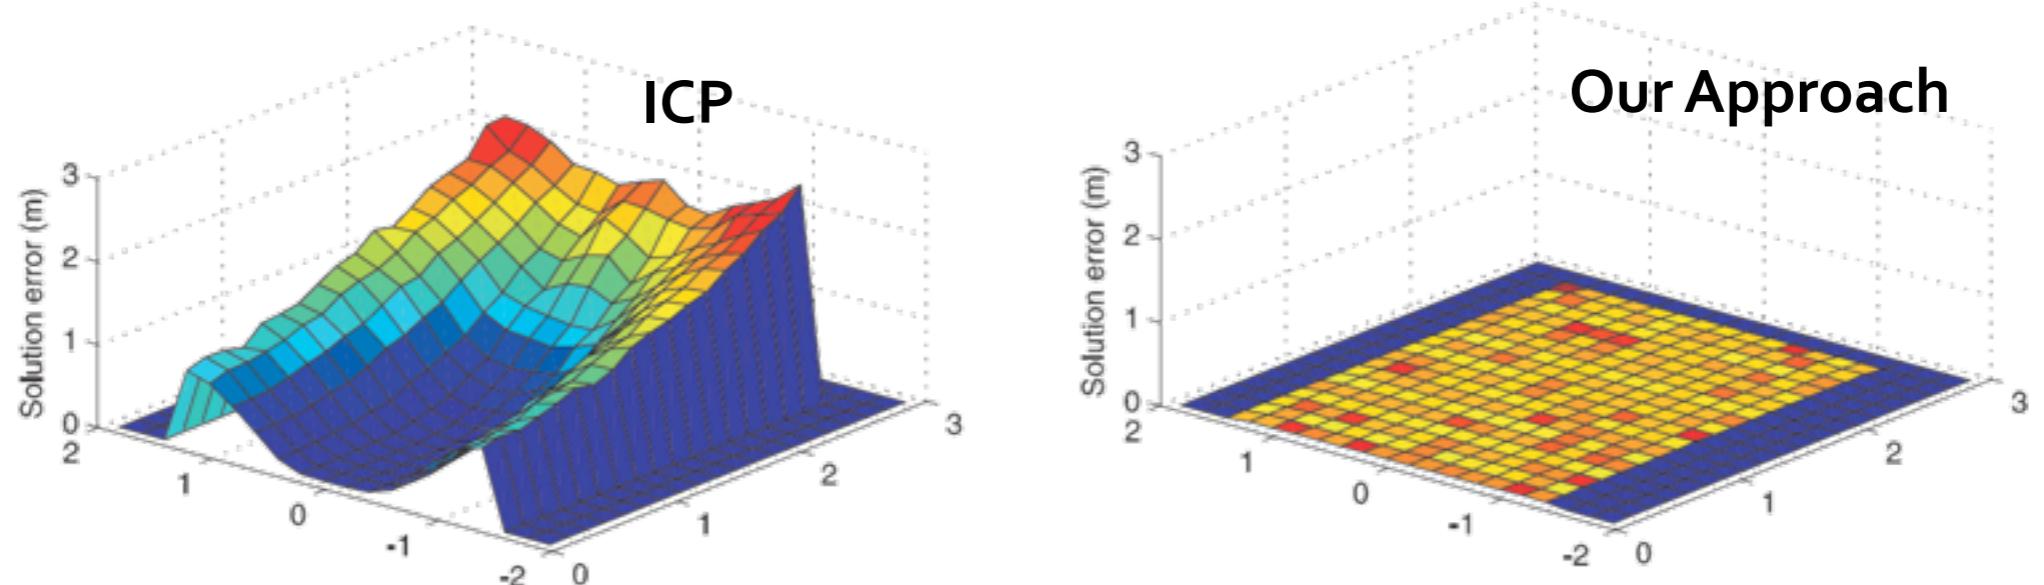
- Immunity to local minima (versus ICP)



- Search time is largely independent of search window
  - Low-probability alignments are quickly ruled out
  - Typical matching in  $\sim 2$  ms

# Matching Results

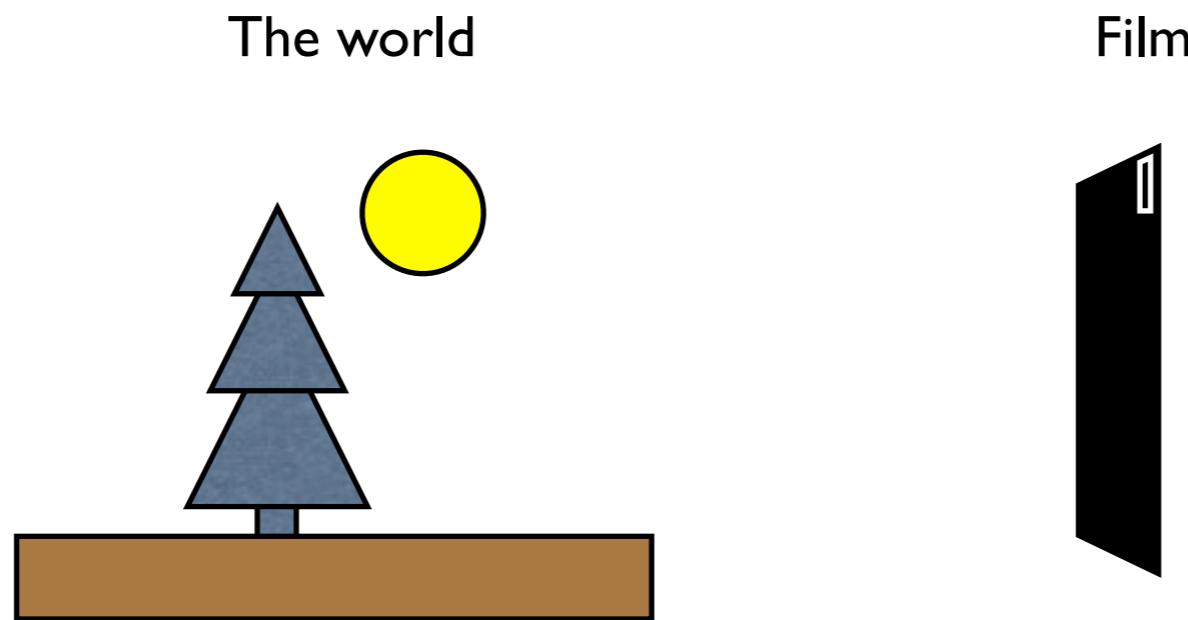
- Immunity to local minima (versus ICP)



- Search time is largely independent of search window
  - Low-probability alignments are quickly ruled out
  - Typical matching in  $\sim 2$  ms

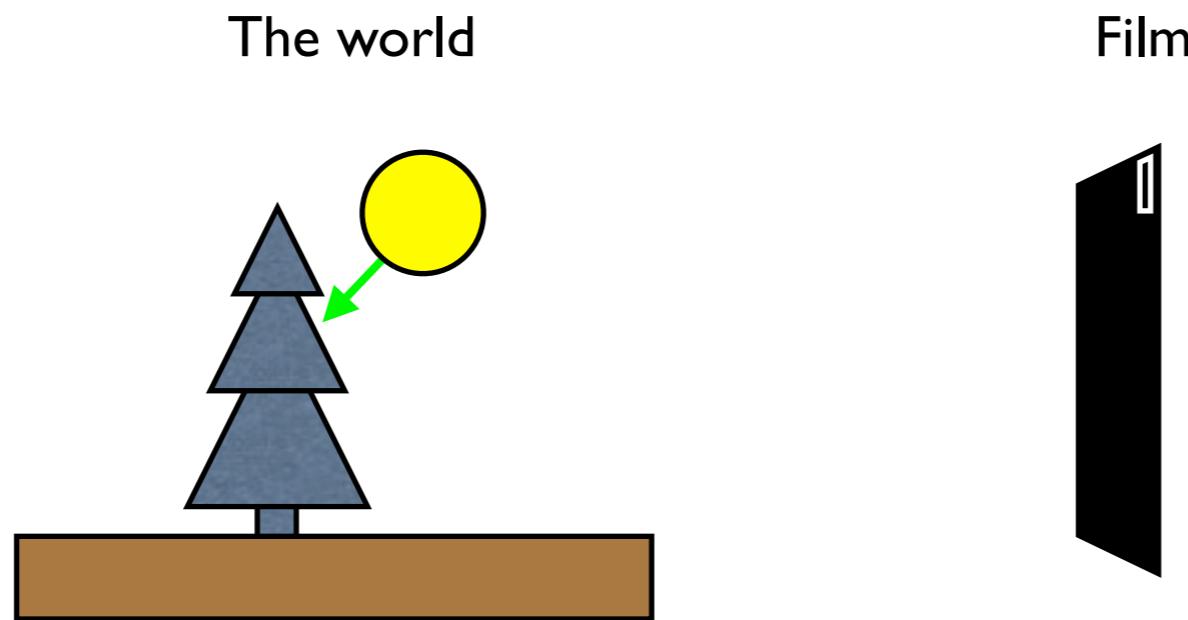
# Cameras and Image Formation

# World Simplest Camera?



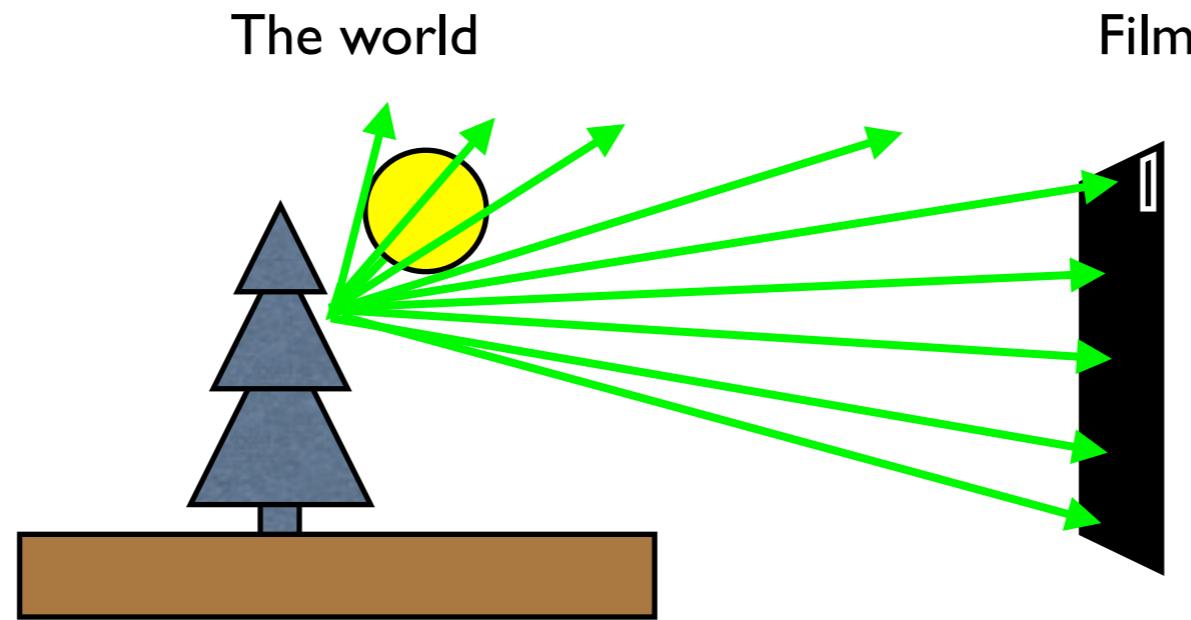
- Just hold up a piece of film
- Do we get an image on the film?
  - ▶ For each piece of the film, where do the photons come from?

# World Simplest Camera?



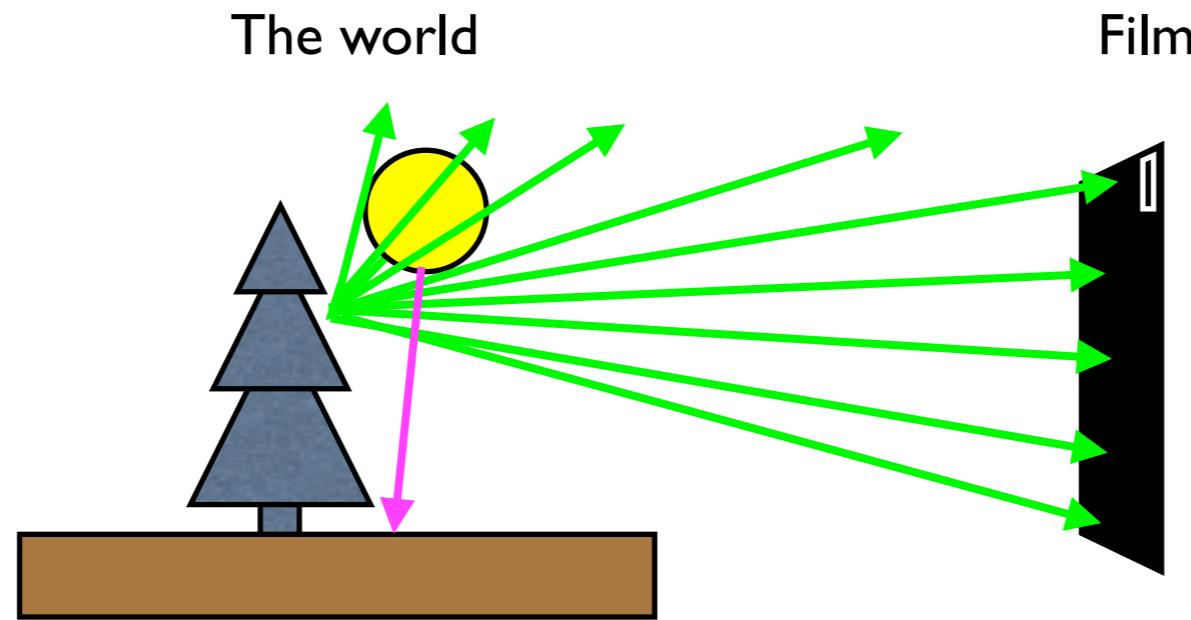
- Just hold up a piece of film
- Do we get an image on the film?
  - ▶ For each piece of the film, where do the photons come from?

# World Simplest Camera?



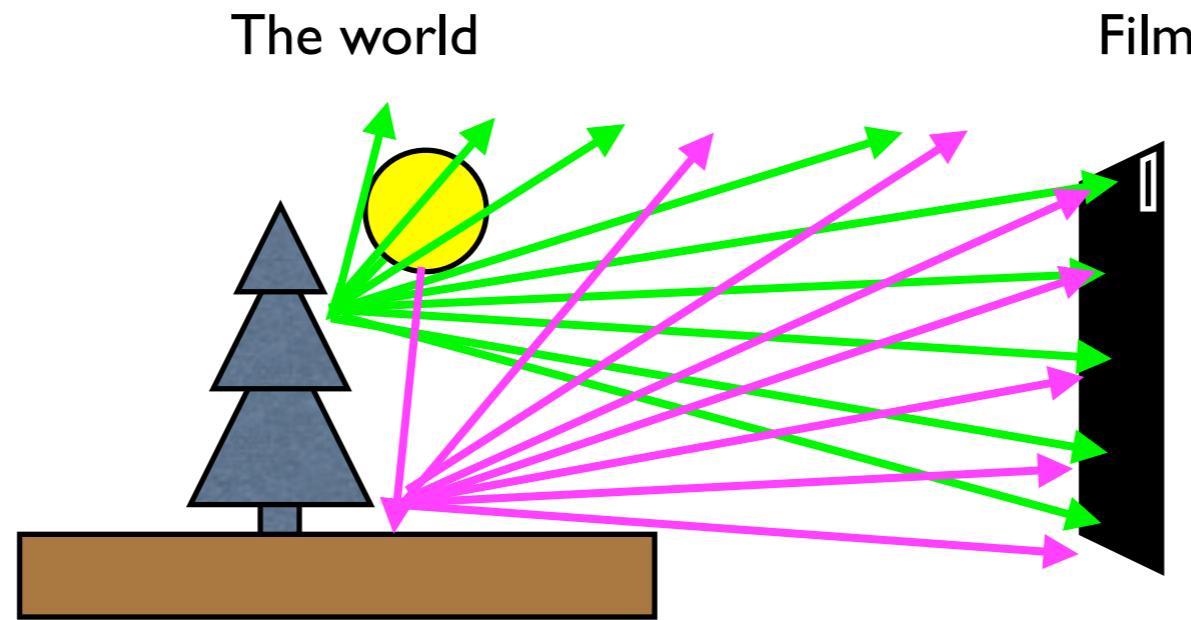
- Just hold up a piece of film
- Do we get an image on the film?
  - ▶ For each piece of the film, where do the photons come from?

# World Simplest Camera?



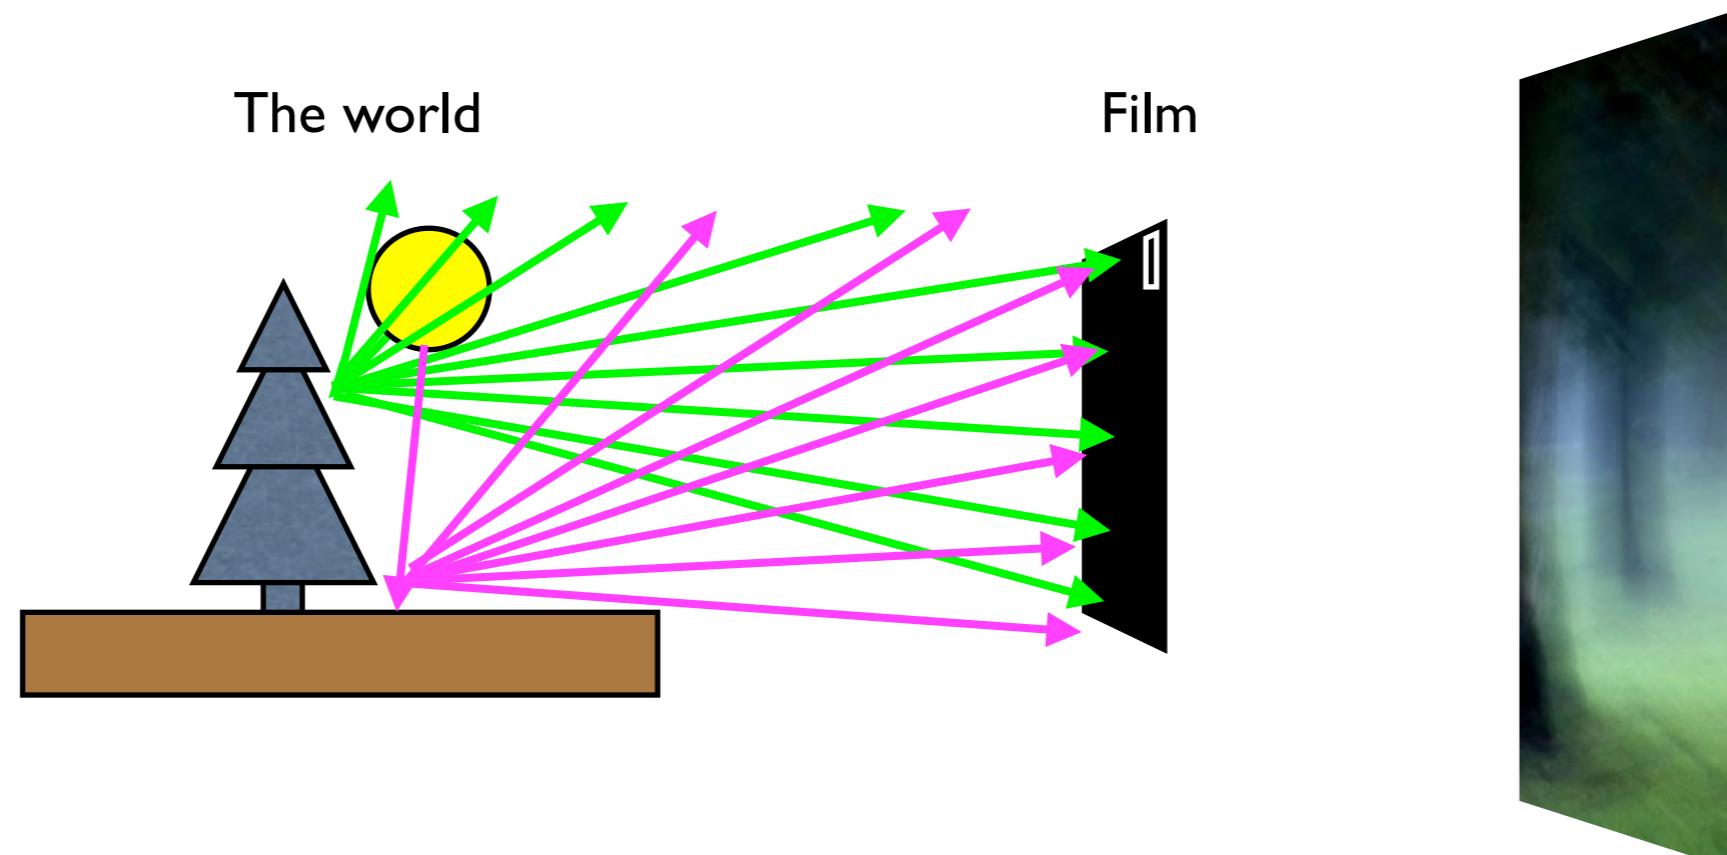
- Just hold up a piece of film
- Do we get an image on the film?
  - ▶ For each piece of the film, where do the photons come from?

# World Simplest Camera?



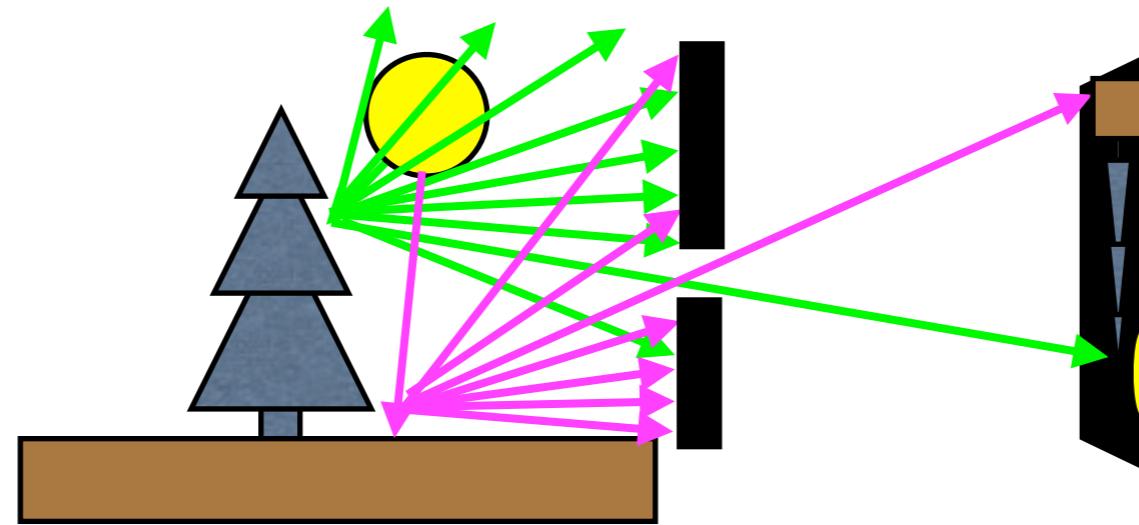
- Just hold up a piece of film
- Do we get an image on the film?
  - ▶ For each piece of the film, where do the photons come from?

# World Simplest Camera?



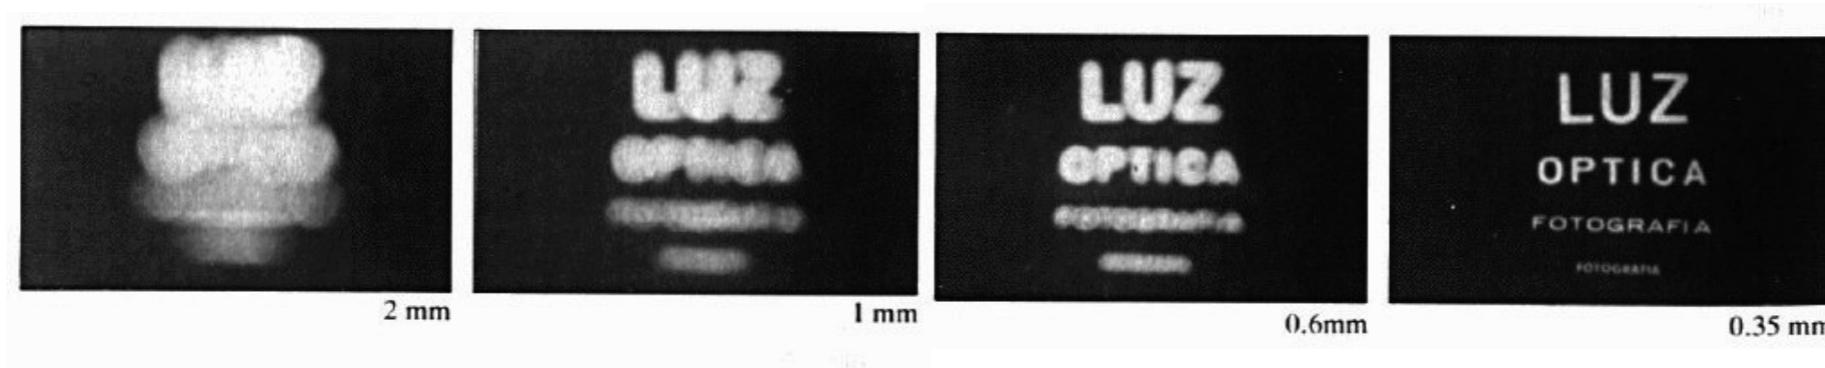
- Just hold up a piece of film
- Do we get an image on the film?
  - ▶ For each piece of the film, where do the photons come from?

# Let's add an aperture



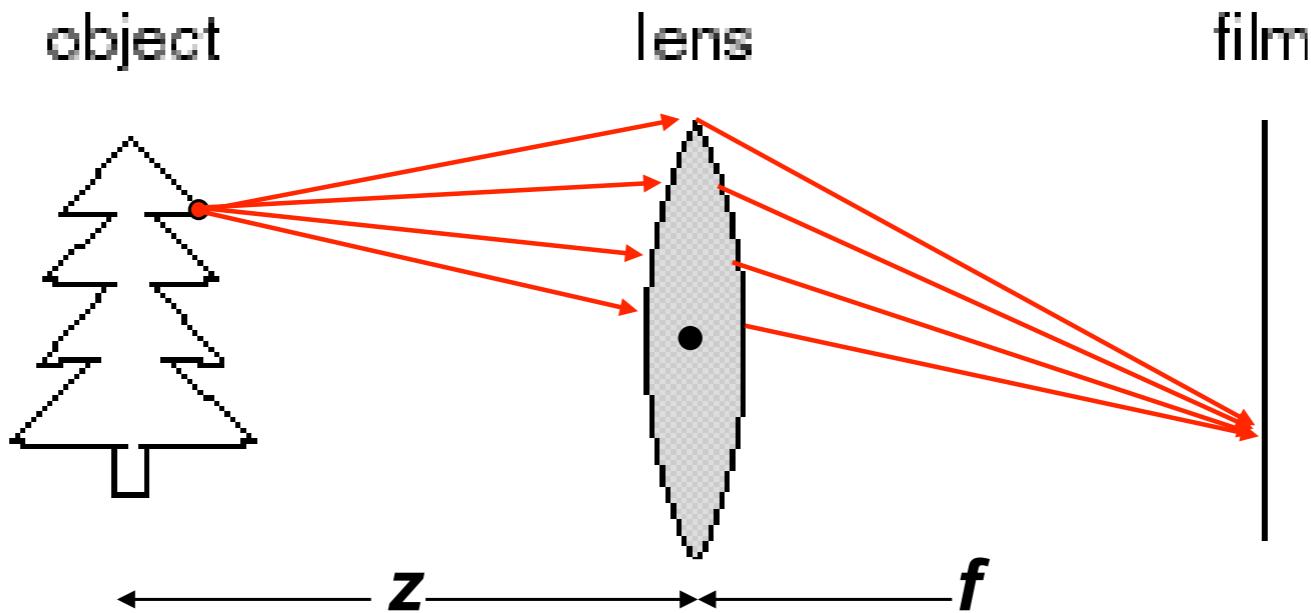
- An aperture blocks all but a small subset of the rays
  - ▶ Causes the image to appear in focus!

# Aperture Size



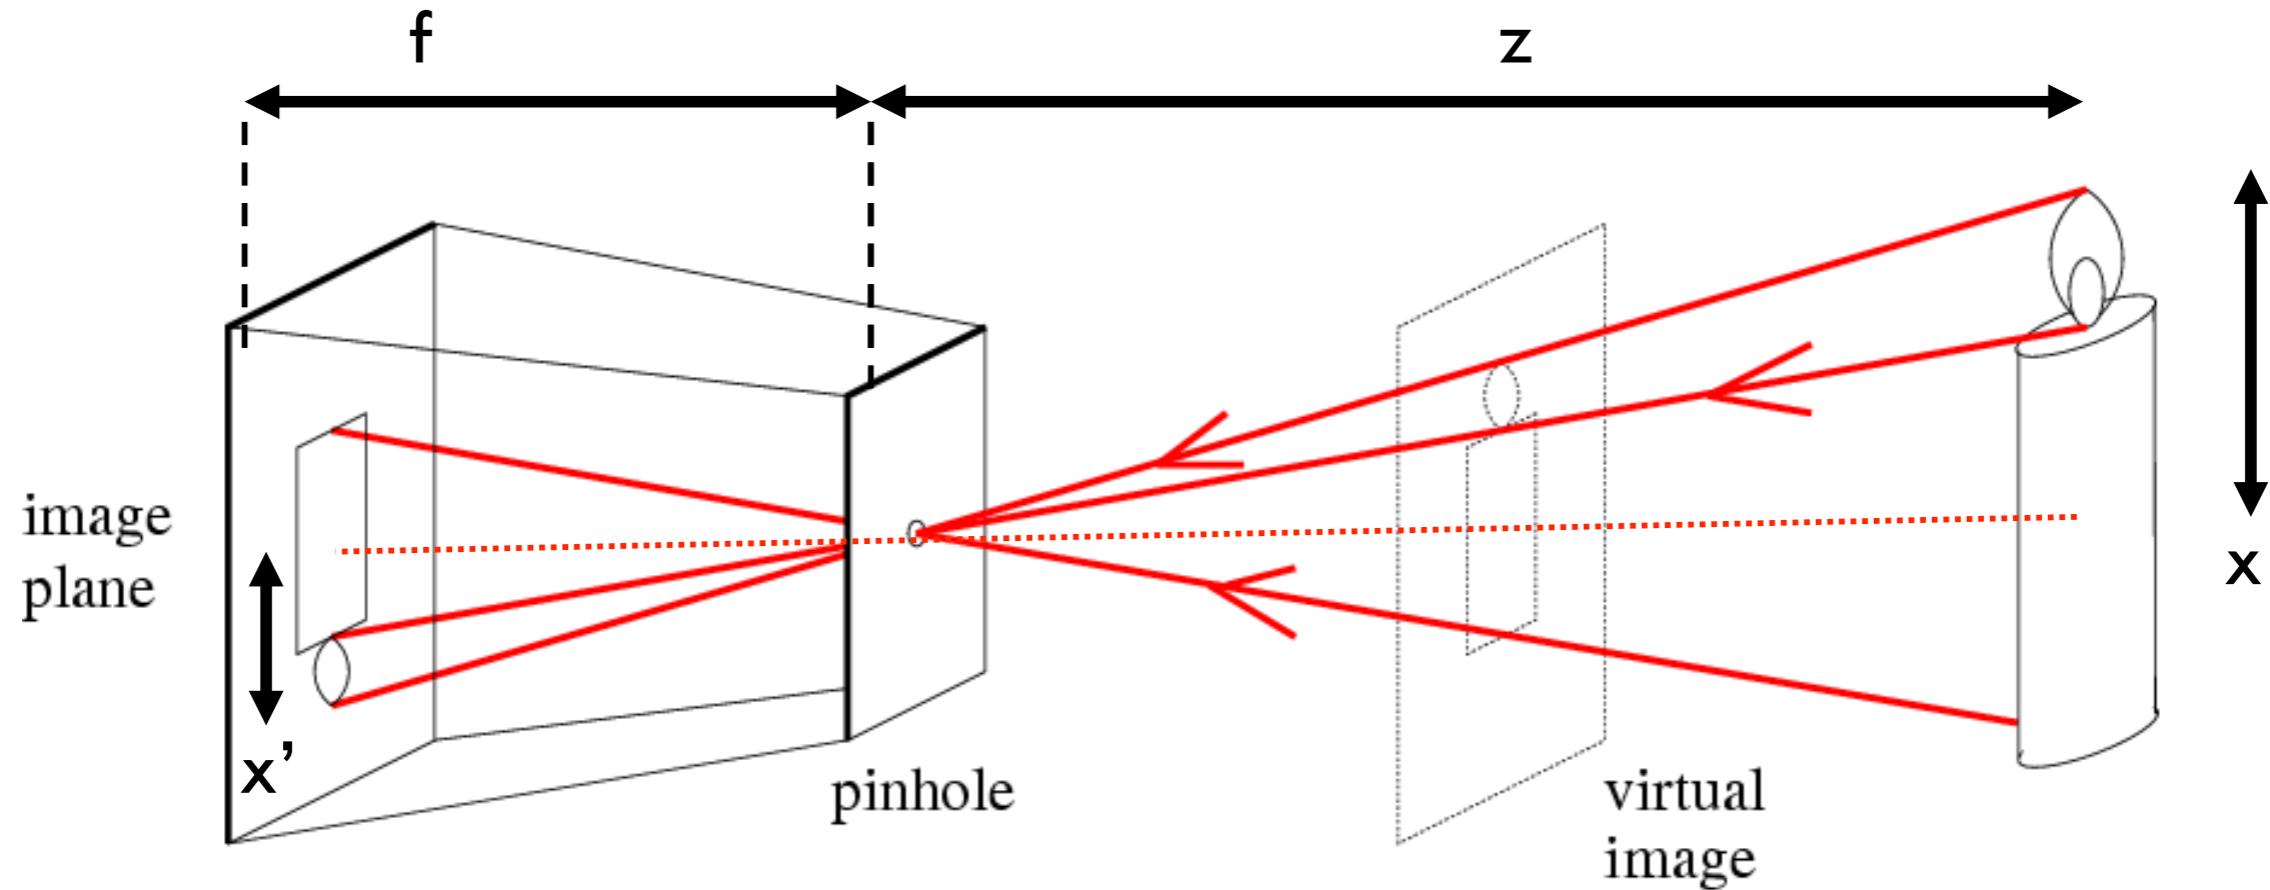
- Why not make the aperture super small?
  - ▶ A “pin-hole” lens.
  - ▶ Not enough light to “register” on our film
- What happens when the aperture is bigger?
  - ▶ More rays can fit through--- blurrier image
- Is there any way of getting a sharp image, but allow more light through?
  - ▶ Yes! A lens.

# Lenses



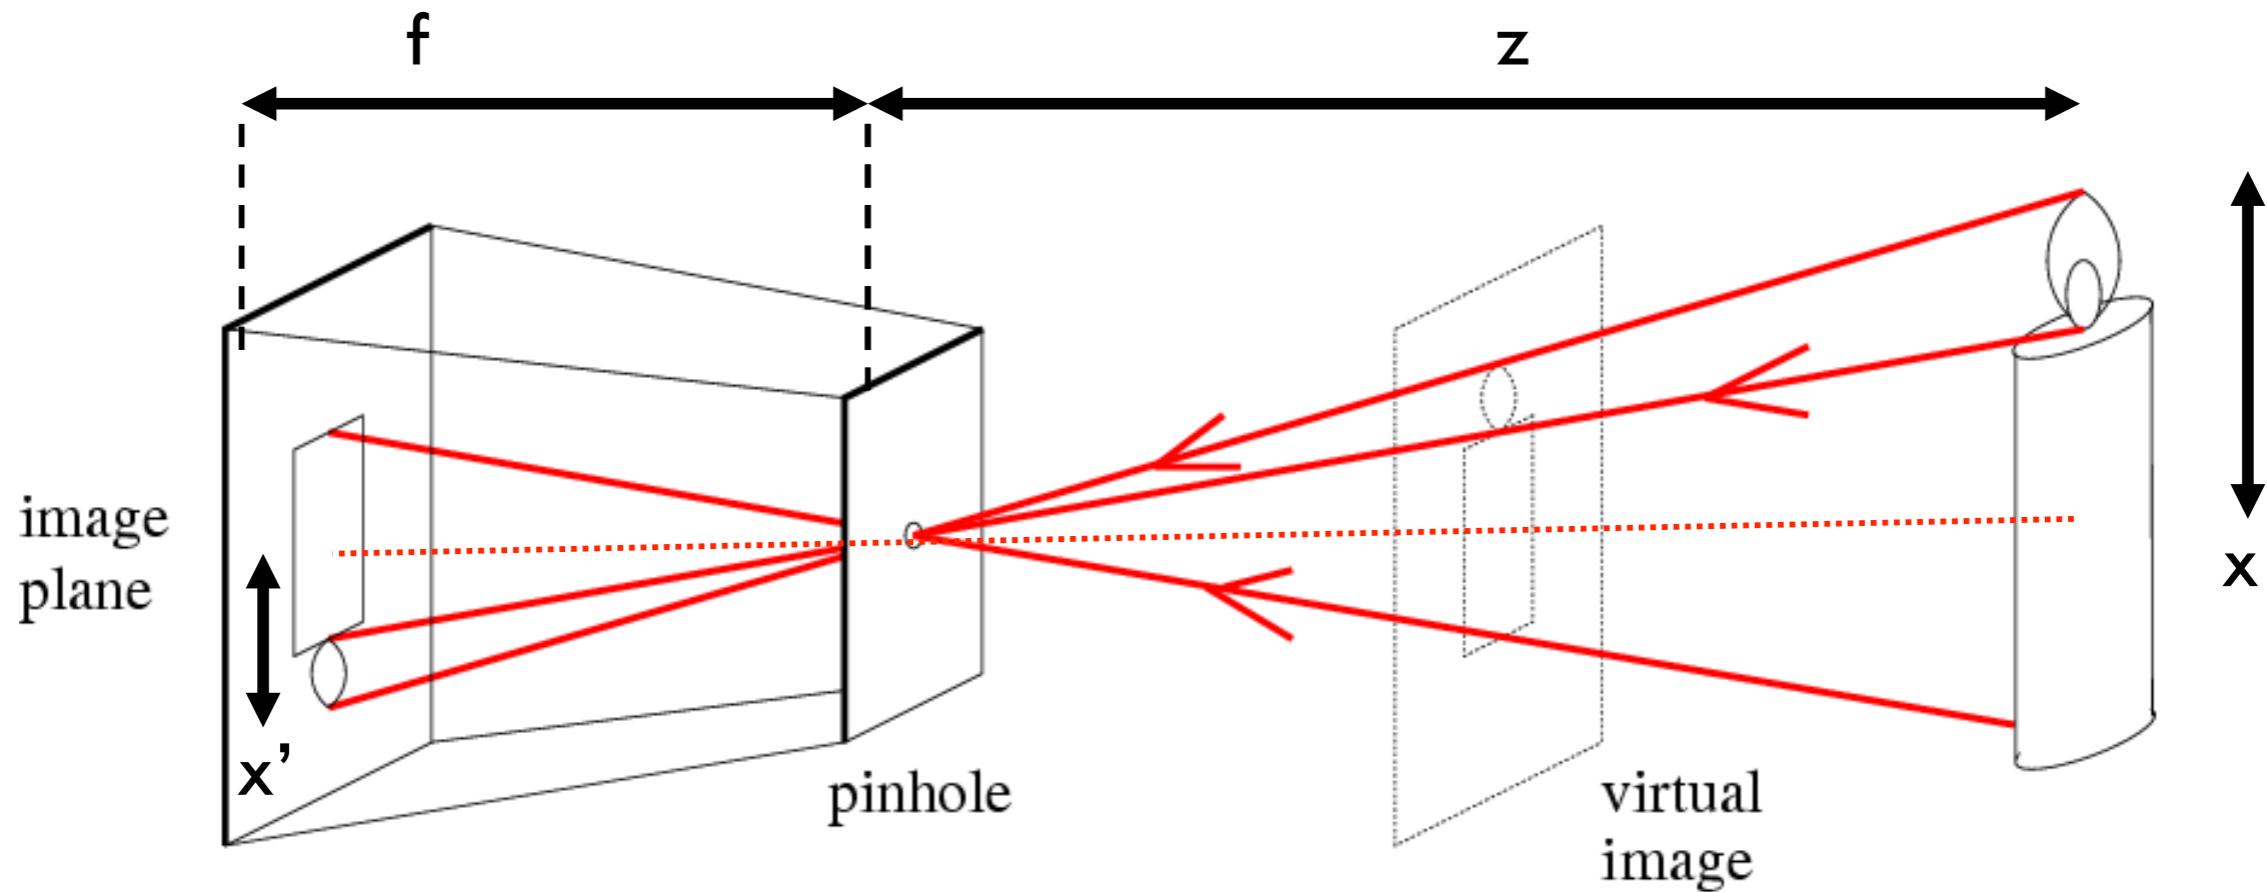
- A lens collects rays with a particular divergence and refocuses them to a point.
  - ▶ But points at the “wrong” distance won’t be refocused exactly.
- *Depth of field*: how much of the scene is in focus
- We’re going to ignore this today, however--- we’re going to assume a “pin-hole” model.

# Perspective Projection



- The pinhole creates two similar triangles
  - ▶ Allows us to determine  $x'$  in terms of  $x$

# Perspective Projection

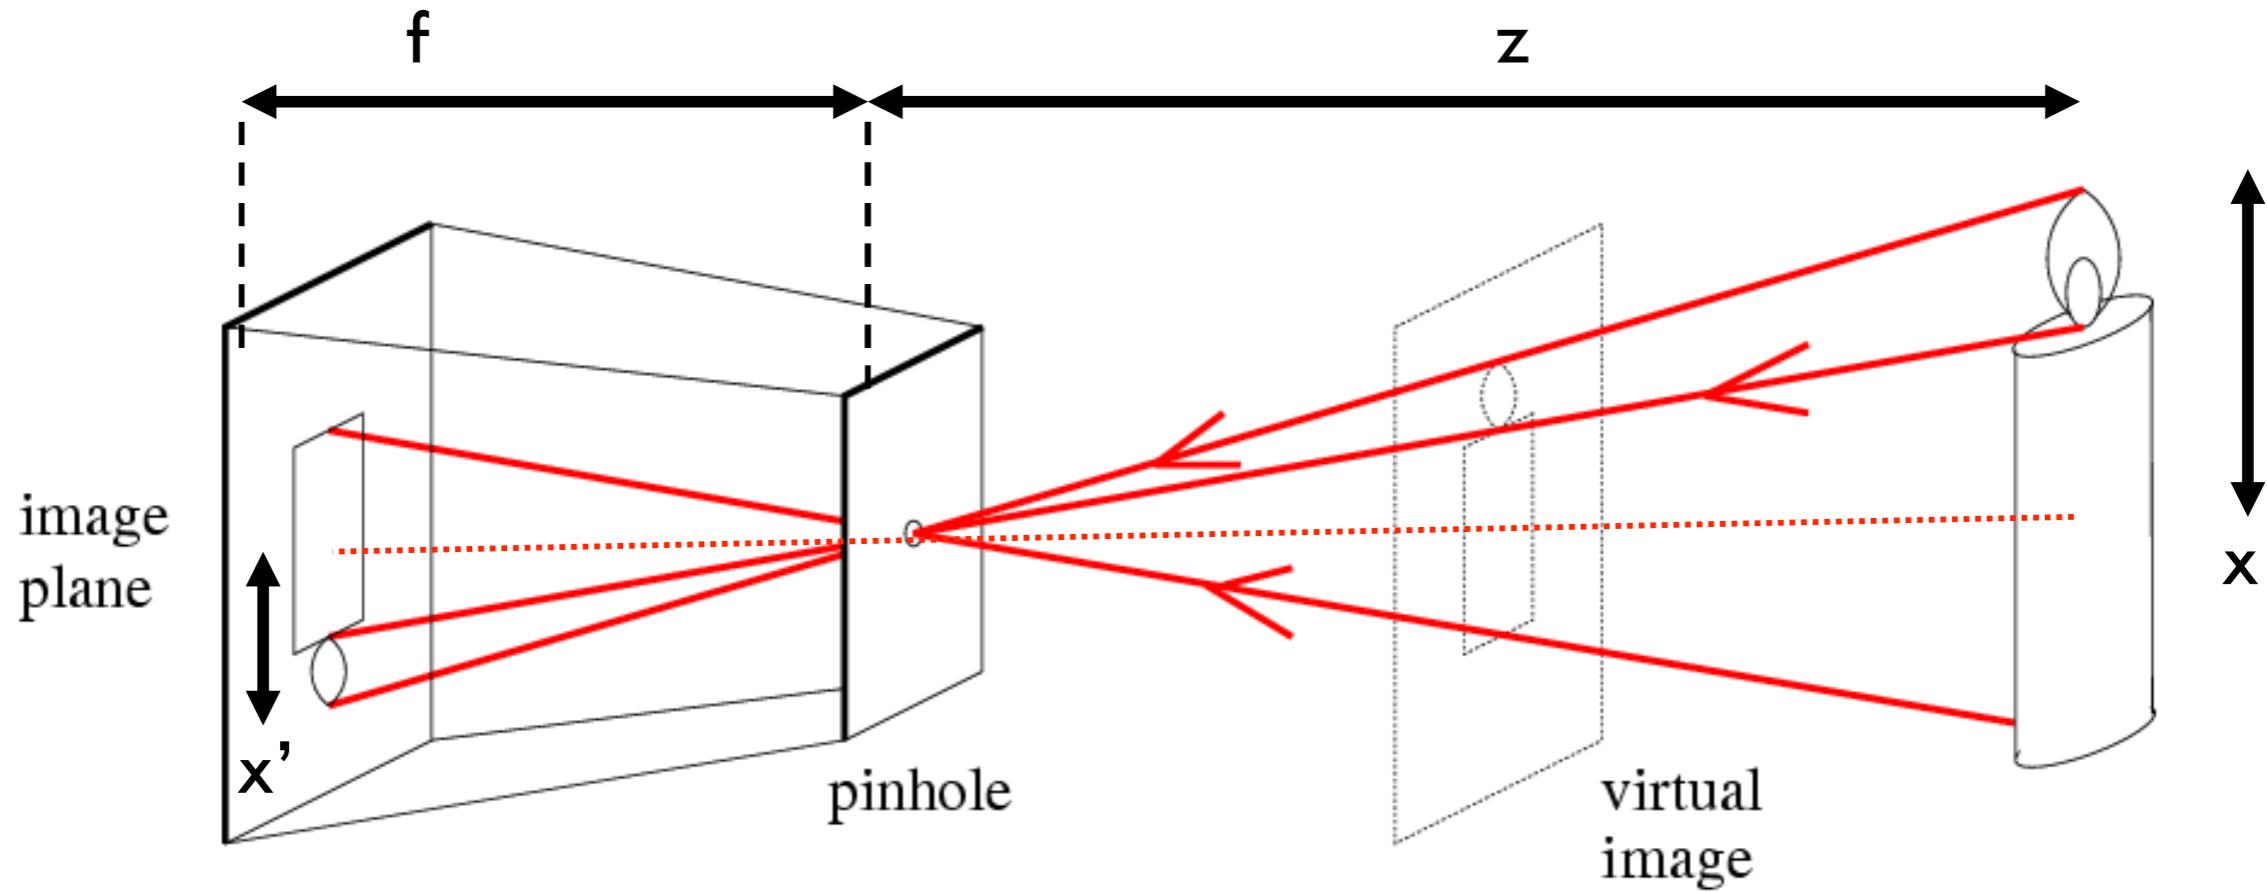


- The pinhole creates two similar triangles
  - ▶ Allows us to determine  $x'$  in terms of  $x$

$$x' = -xf/z$$

(why is it negative? we'll assume from here on out that the camera “unflips” the image.)

# Perspective Projection



- What are the pixel coordinates where the flame appears?
  - ▶  $x' = fx/z + c$
  - ▶ Measure  $f$  in “pixels” and add an offset (so that the “middle” pixel is in the middle of the image)

# The Perspective Matrix

- Suppose we write a point in the world (like the position of the candle flame) as a vector:

$$p = \begin{bmatrix} x \\ y \\ z \end{bmatrix}$$

the matrix M is a function of z! We want M to work for any point p.

We also are unable to handle cx, cy

- Can we write a matrix so that  $p' = Mp$ ?

$$p' = \begin{bmatrix} fx/z + c_x \\ fy/z + c_y \end{bmatrix} \stackrel{?}{=} \begin{bmatrix} f/z & 0 & 0 \\ 0 & f/z & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \end{bmatrix}$$

# The Perspective Matrix

- Suppose we write a point in the world (like the position of the candle flame) as a vector:

$$p = \begin{bmatrix} x \\ y \\ z \end{bmatrix}$$

the matrix M is a function of z! We want M to work for any point p.

We also are unable to handle cx, cy

- Can we write a matrix so that  $p' = Mp$ ?

$$p' = \begin{bmatrix} fx/z + c_x \\ fy/z + c_y \end{bmatrix} \stackrel{?}{=} \begin{bmatrix} f/z & 0 & 0 \\ 0 & f/z & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \end{bmatrix}$$

What's wrong with this?

# Homogenous Coordinates

- We'll introduce a new convention, *homogenous coordinates*.
- We write points just the way we did before, but add an extra row:
  - ▶ The extra row is a *scale factor* for the whole vector.

there's a whole related set of geometry, but all we need to know in this lecture is that there's a funny scale factor.

# Homogenous Coordinates

- We'll introduce a new convention, *homogenous coordinates*.
- We write points just the way we did before, but add an extra row:
  - ▶ The extra row is a *scale factor* for the whole vector.

$$p = \begin{bmatrix} x \\ y \\ z \end{bmatrix}$$

becomes

$$p = \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

there's a whole related set of geometry, but all we need to know in this lecture is that there's a funny scale factor.

# Homogenous Coordinates

- We'll introduce a new convention, *homogenous coordinates*.
- We write points just the way we did before, but add an extra row:
  - ▶ The extra row is a *scale factor* for the whole vector.

$$p = \begin{bmatrix} x \\ y \\ z \end{bmatrix}$$

becomes

$$p = \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

What point does this correspond to?

$$\begin{bmatrix} 10 \\ 20 \\ 15 \\ 5 \end{bmatrix}$$

there's a whole related set of geometry, but all we need to know in this lecture is that there's a funny scale factor.

# Do homogeneous coordinates help?

- Eureka!

$$p' = \begin{bmatrix} fx/z + c_x \\ fy/z + c_y \\ 1 \end{bmatrix} = \begin{bmatrix} fx + c_x z \\ fy + c_y z \\ z \end{bmatrix} = \begin{bmatrix} f & 0 & c_x & 0 \\ 0 & f & c_y & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

2d  
homogeneous  
coordinates

2d  
homogeneous  
coordinates

perspective  
transform

3d  
homogeneous  
coordinates

# Moving the camera

- Basic idea:
  - ▶ Moving the camera is exactly the same thing as moving the world in the opposite way.
- But how do we represent motion?
  - ▶ With a matrix!

# Translation

- Suppose I want to shift all objects by  $T_x, T_y, T_z$ :

$$\begin{bmatrix} x + T_x \\ y + T_y \\ z + T_z \\ 1 \end{bmatrix} = \begin{bmatrix} ? \\ \square \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

# Translation

- Suppose I want to shift all objects by  $T_x, T_y, T_z$ :

$$\begin{bmatrix} x + T_x \\ y + T_y \\ z + T_z \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & T_x \\ 0 & 1 & 0 & T_y \\ 0 & 0 & 1 & T_z \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

# Rotation

- How about a rotation of 90 degrees around the Z axis?

$$\begin{bmatrix} -y \\ x \\ z \\ 1 \end{bmatrix} = \begin{bmatrix} ? \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

# Rotation

- How about a rotation of 90 degrees around the Z axis?

$$\begin{bmatrix} -y \\ x \\ z \\ 1 \end{bmatrix} = \begin{bmatrix} 0 & -1 & 0 & 0 \\ 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

# Rotation Matrices: Intuition

- What do those R's mean?
  - They correspond to the directions in the direction!

You can easily verify  
that the  $[1 \ 0 \ 0 \ 1]$  vector  
is transformed to  $[R_{00} \ R_{10} \ R_{20} \ 1]$

$$\begin{bmatrix} x' \\ y' \\ z' \\ 1 \end{bmatrix} = \begin{bmatrix} \hat{x}' & \hat{y}' & \hat{z}' \\ R_{00} & R_{01} & R_{02} \\ R_{10} & R_{11} & R_{12} \\ R_{20} & R_{21} & R_{22} \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

# Rotation Matrices: Intuition

- What do those R's mean?
  - They correspond to the directions in the direction!

You can easily verify  
that the  $[1 \ 0 \ 0 \ 1]$  vector  
is transformed to  $[R_{00} \ R_{10} \ R_{20} \ 1]$

$$\begin{bmatrix} x' \\ y' \\ z' \\ 1 \end{bmatrix} = \begin{bmatrix} \hat{x}' & \hat{y}' & \hat{z}' \\ R_{00} & R_{01} & R_{02} \\ R_{10} & R_{11} & R_{12} \\ R_{20} & R_{21} & R_{22} \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

## Important Properties:

1. Each column is a unit vector
2. Each column is perpendicular to the others

# Rigid-Body Transformations

- The product of two rigid-body transformations is **always** another rigid-body transformation!
- So no matter how the object has been translated or rotated, we can describe its position with a single 4x4 matrix, which has the structure:

$$\begin{bmatrix} x' \\ y' \\ z' \\ 1 \end{bmatrix} = \begin{bmatrix} R_{00} & R_{01} & R_{02} & T_x \\ R_{10} & R_{11} & R_{12} & T_y \\ R_{20} & R_{21} & R_{22} & T_z \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

# Putting it all together

$$\begin{bmatrix} x' \\ y' \\ s \end{bmatrix} = \begin{bmatrix} f & 0 & c_x & 0 \\ 0 & f & c_y & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} R_{00} & R_{01} & R_{02} & T_x \\ R_{10} & R_{11} & R_{12} & T_y \\ R_{20} & R_{21} & R_{22} & T_z \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

2d  
homogeneous  
pixel (camera)  
coordinates

Focal length and  
focal center of camera

Rigidly move every object in  
the world to simulate the  
camera's true position

3d  
homogenous  
(world)  
coordinates

# Putting it all together

$$\begin{bmatrix} x' \\ y' \\ s \end{bmatrix} = \begin{bmatrix} f & 0 & c_x & 0 \\ 0 & f & c_y & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} R_{00} & R_{01} & R_{02} & T_x \\ R_{10} & R_{11} & R_{12} & T_y \\ R_{20} & R_{21} & R_{22} & T_z \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

2d  
homogeneous  
pixel (camera)  
coordinates

Focal length and  
focal center of camera

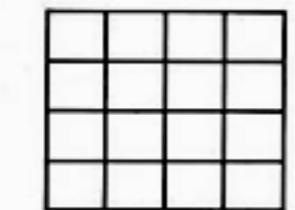
Rigidly move every object in  
the world to simulate the  
camera's true position

3d  
homogenous  
(world)  
coordinates

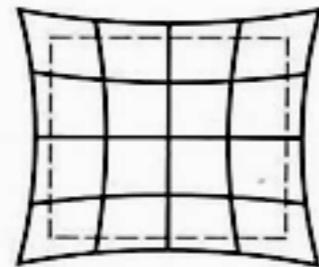
“Intrinsics”      “Extrinsics”

# Lens distortions

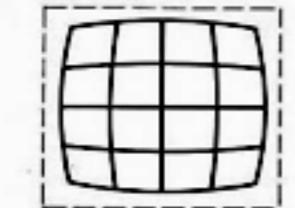
- Unfortunately, real (imperfect) lenses further complicate life.



Undistorted



Pin cushion

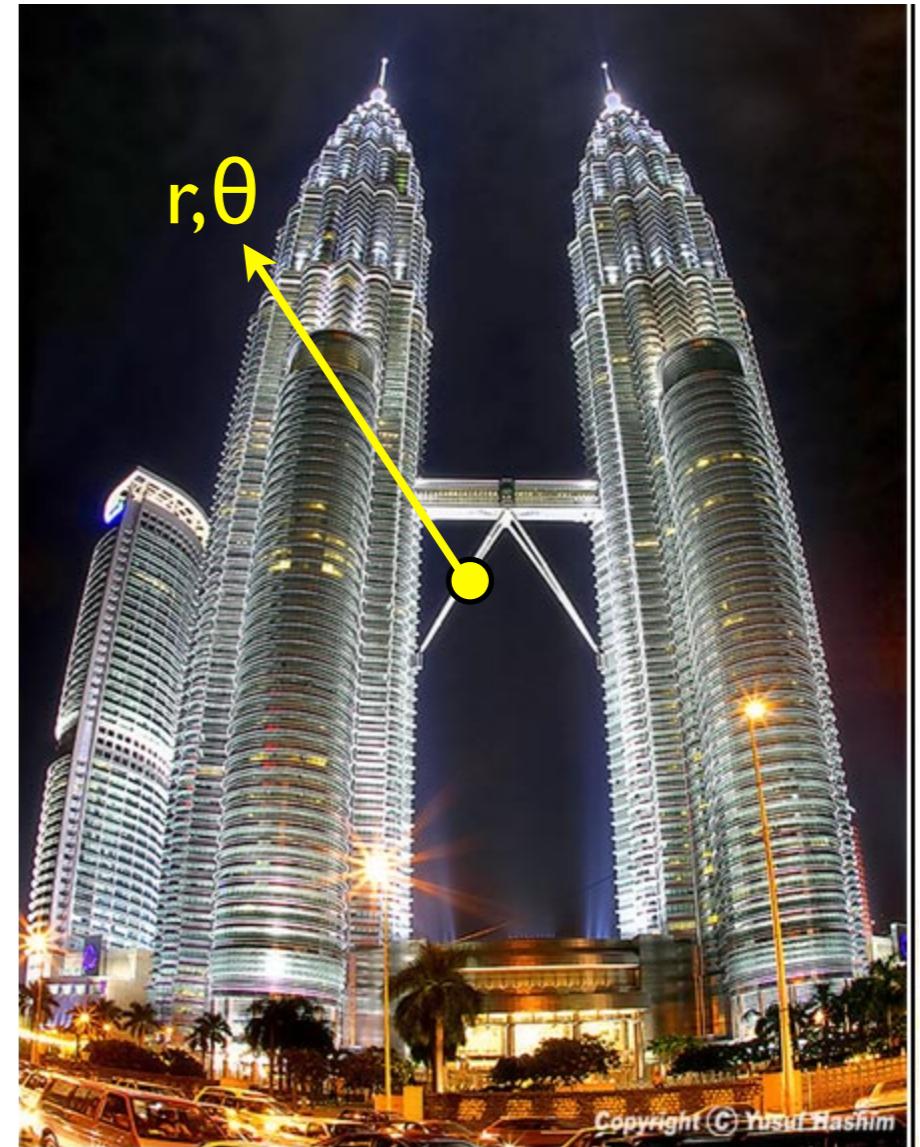


Barrel (common)



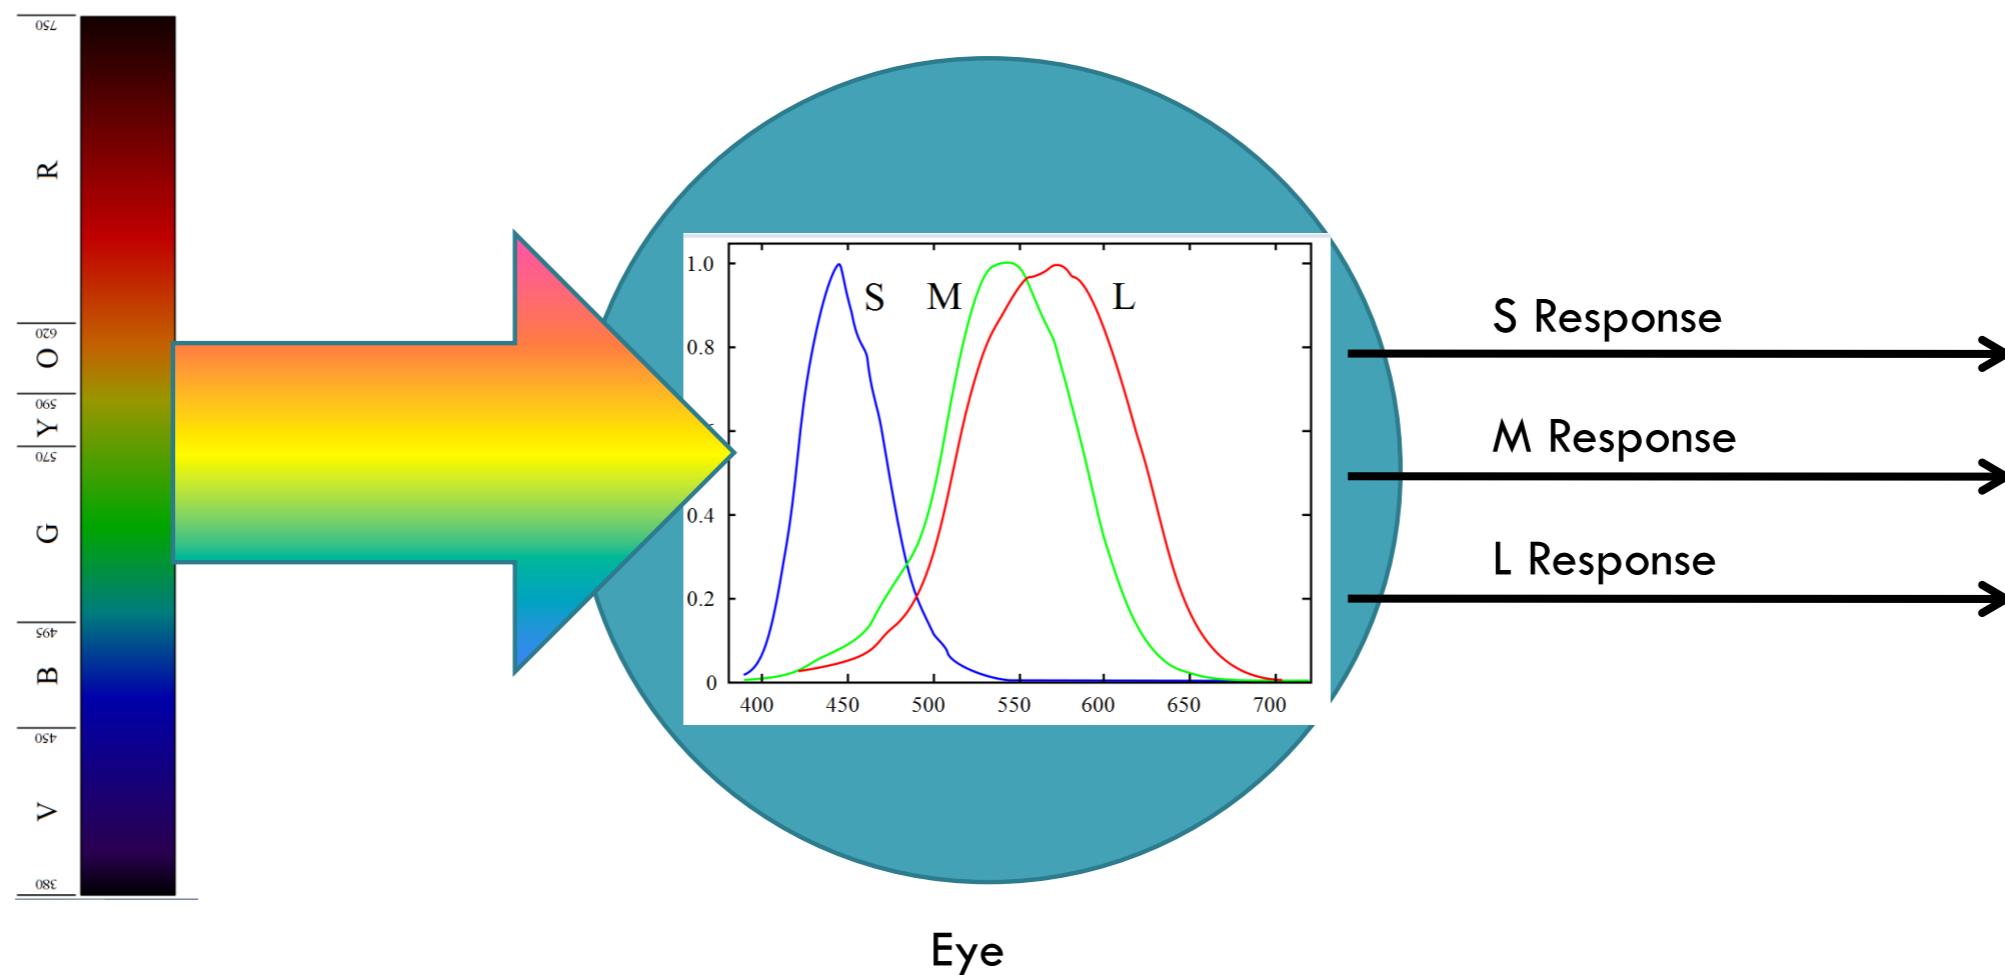
# Correcting for lens distortion

- Radial Distortion
  1. Compute the pixel coordinates assuming the lens is undistorted
  2. Convert to polar form
  3. Compute  $r' = f(r)$
  4. Convert  $r'$  and  $\theta$  back to Cartesian coordinates.
- Function  $f()$  is typically nasty polynomial functions.
  - ▶ We find the parameters by using non-linear optimization algorithms

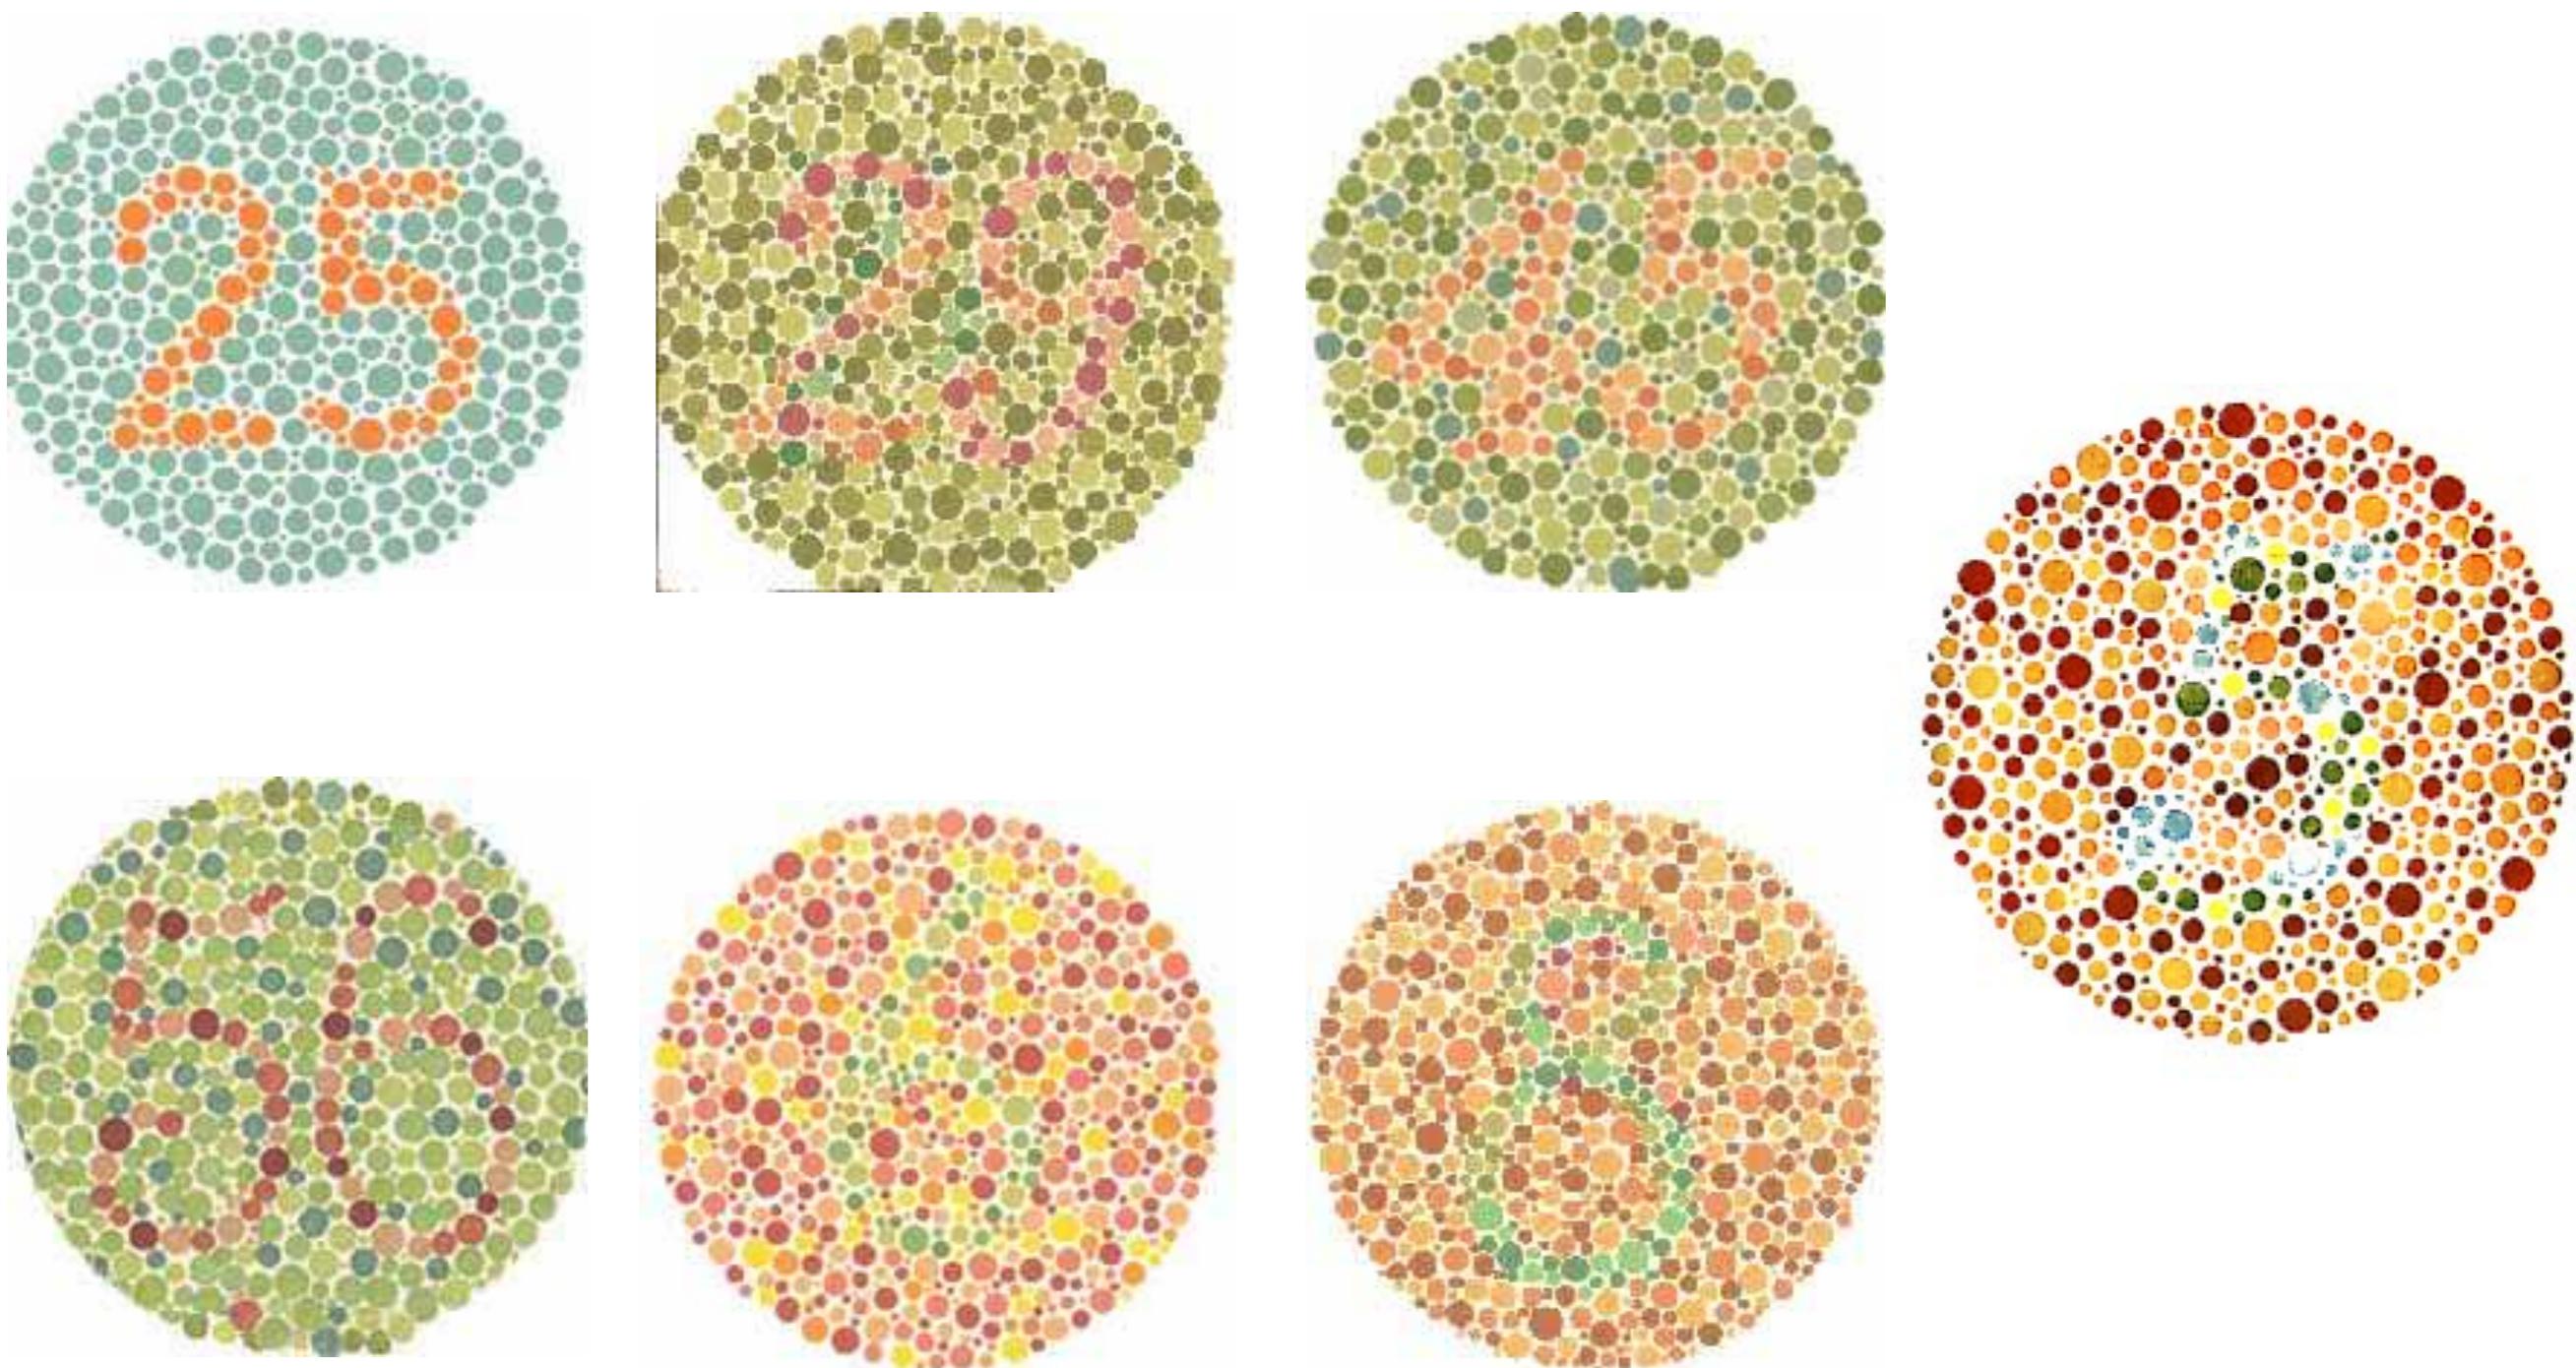


# Color Cameras

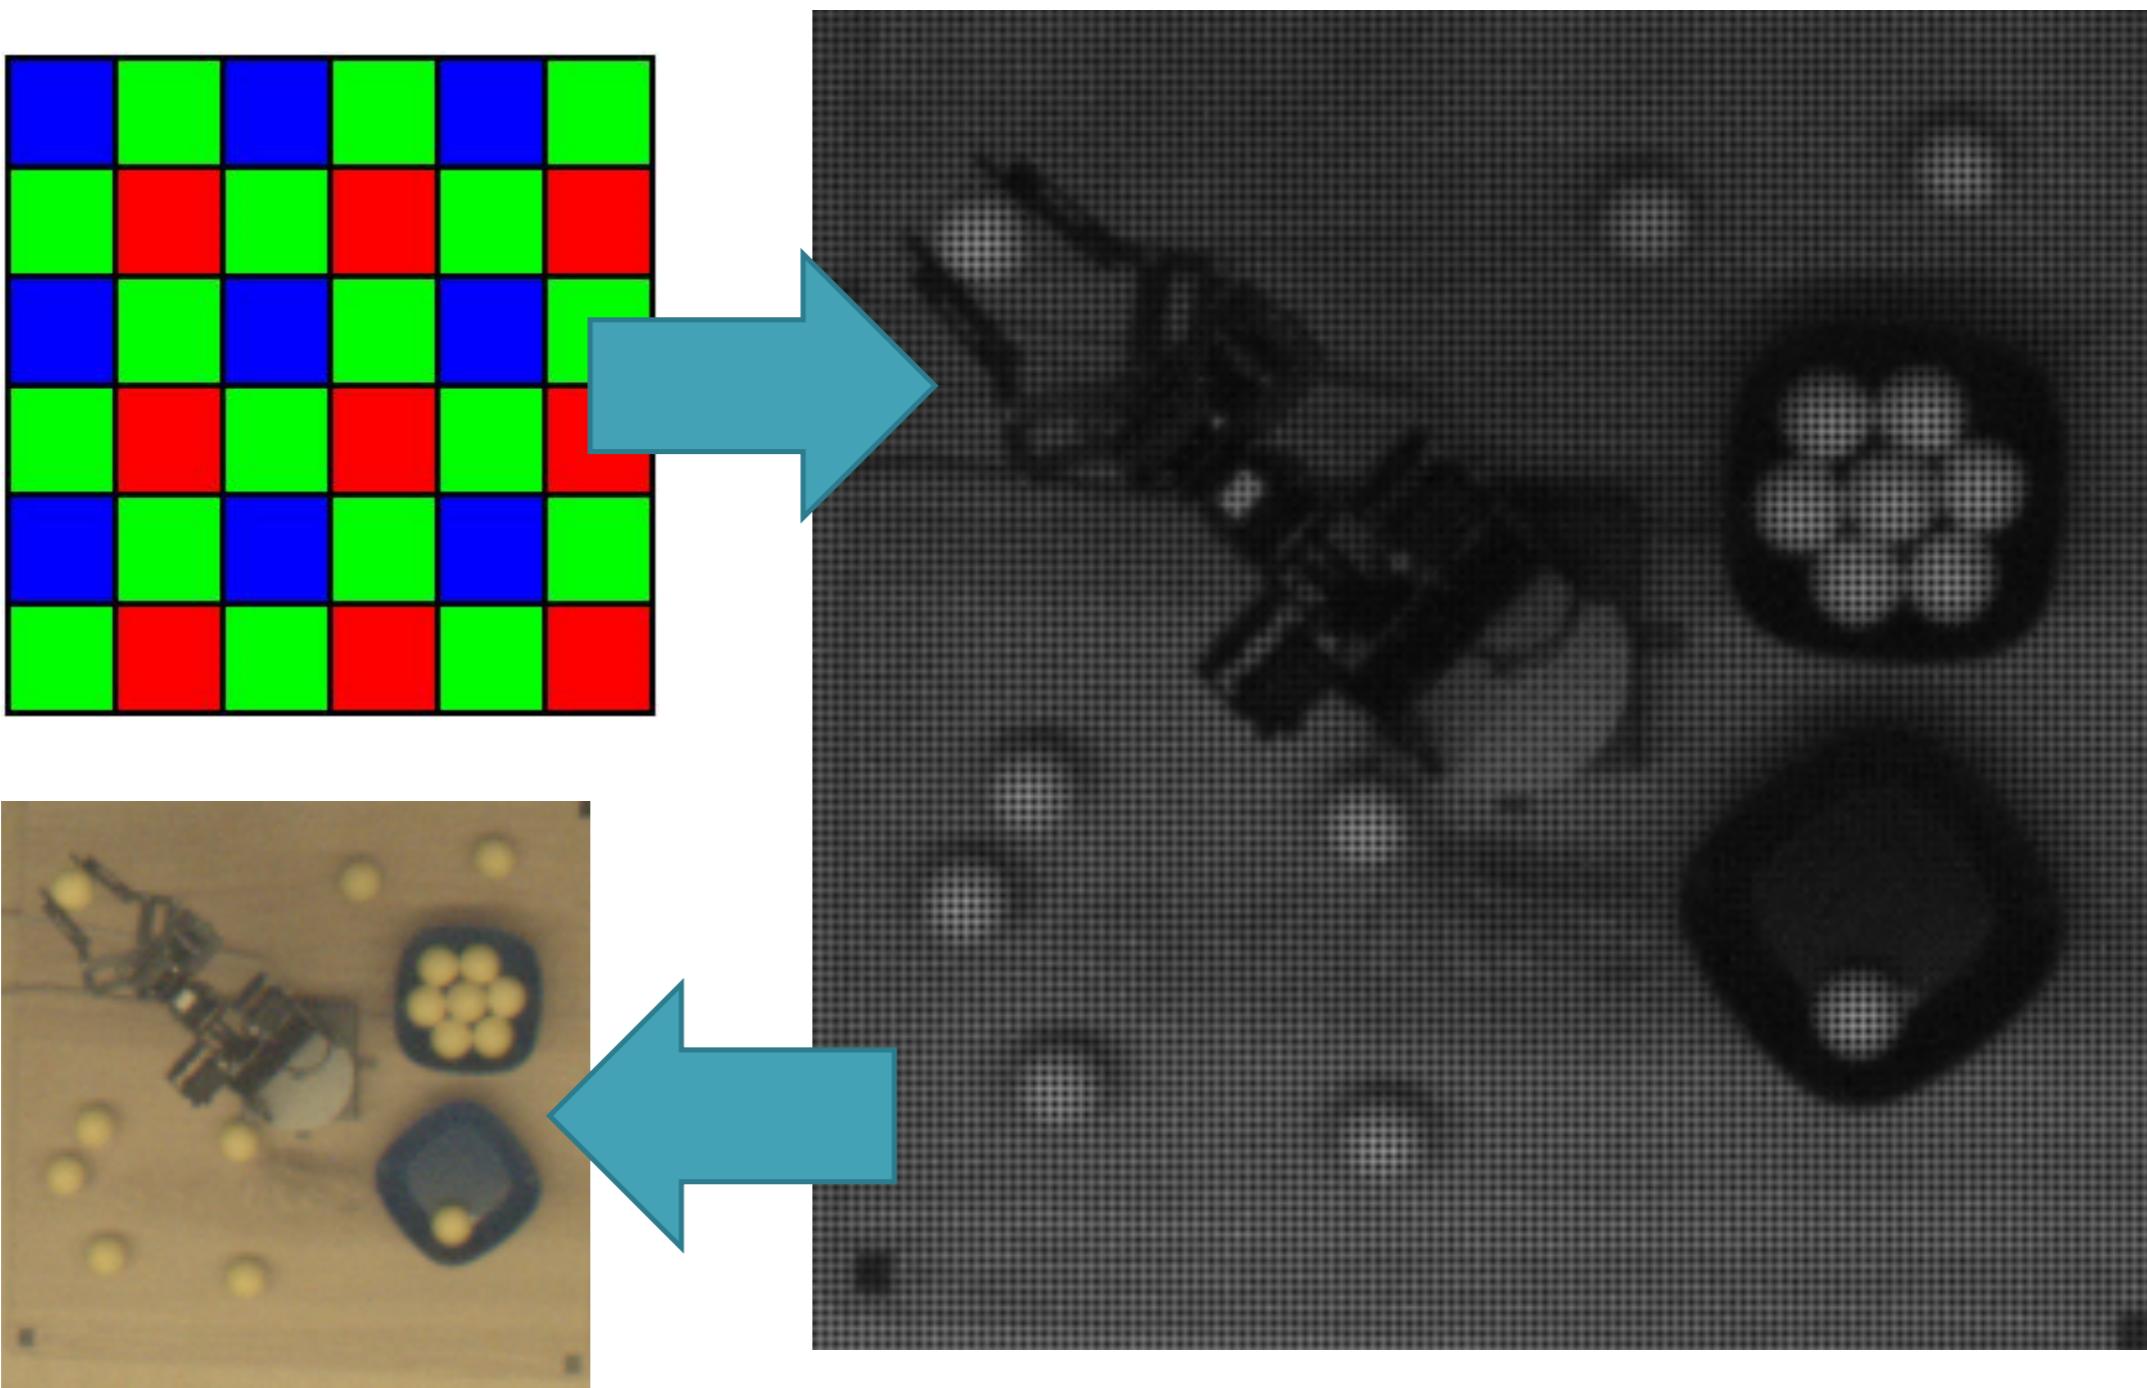
- Incoming light is described in terms of a *power spectral density*
- “Color” isn’t a physical property of light
  - ▶ It’s made up by our eyes and brain!
  - ▶ Different types of incoming light can have the same “color”



# Just for fun...



# Bayer Patterns



# Bayer Patterns

- Why does this matter?
  - ▶ At each pixel, two color channels are interpolated based on nearby pixels
- Thus, a color camera is more blurry than a monochrome camera.
  - ▶ e.g., Monochrome cameras give slightly better results for AprilTags

